

# EECS 370 - Lecture 13

## Pipelining and Data Hazards



# Announcements

- **Project 2a due today at 11:55pm**
- **Homework 2 due Monday 2/23 @11:55pm**
  - Come to office hours with questions
  - Also there are special office hours in 1690 BBB on Monday from 6-9pm for last minute questions.

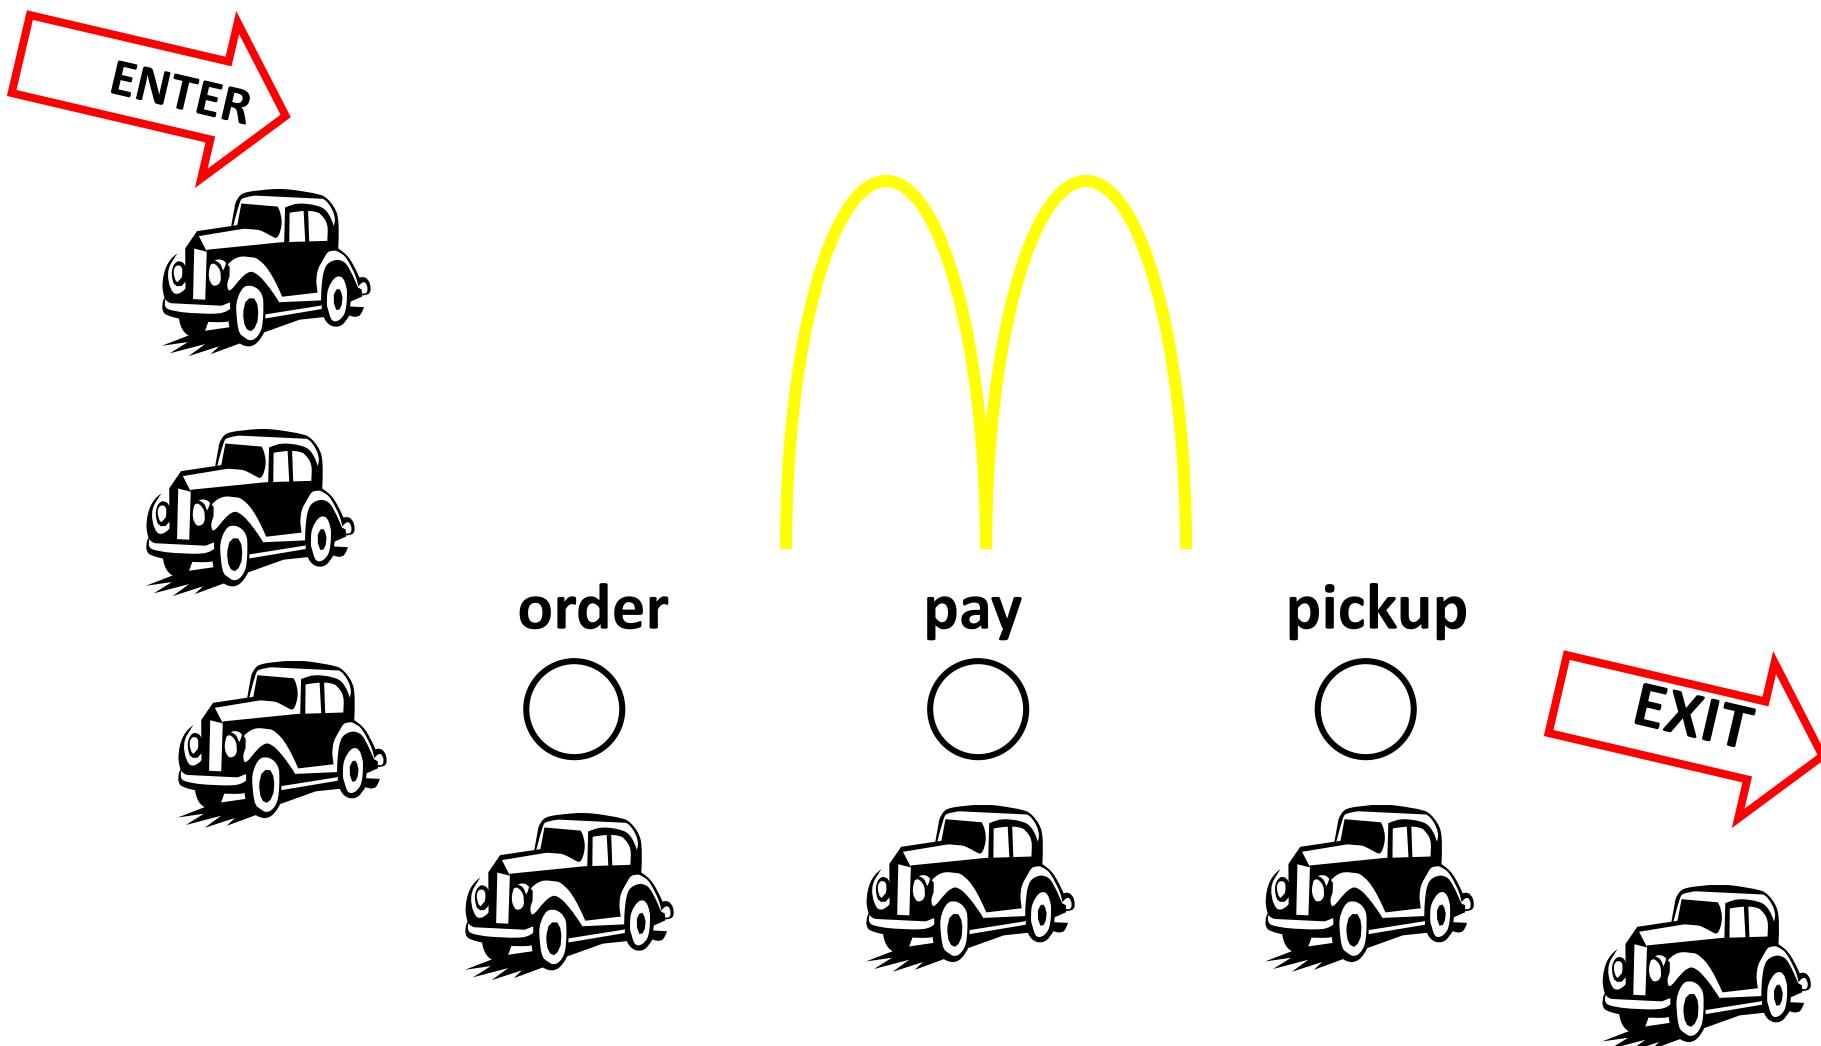


# Pipelining

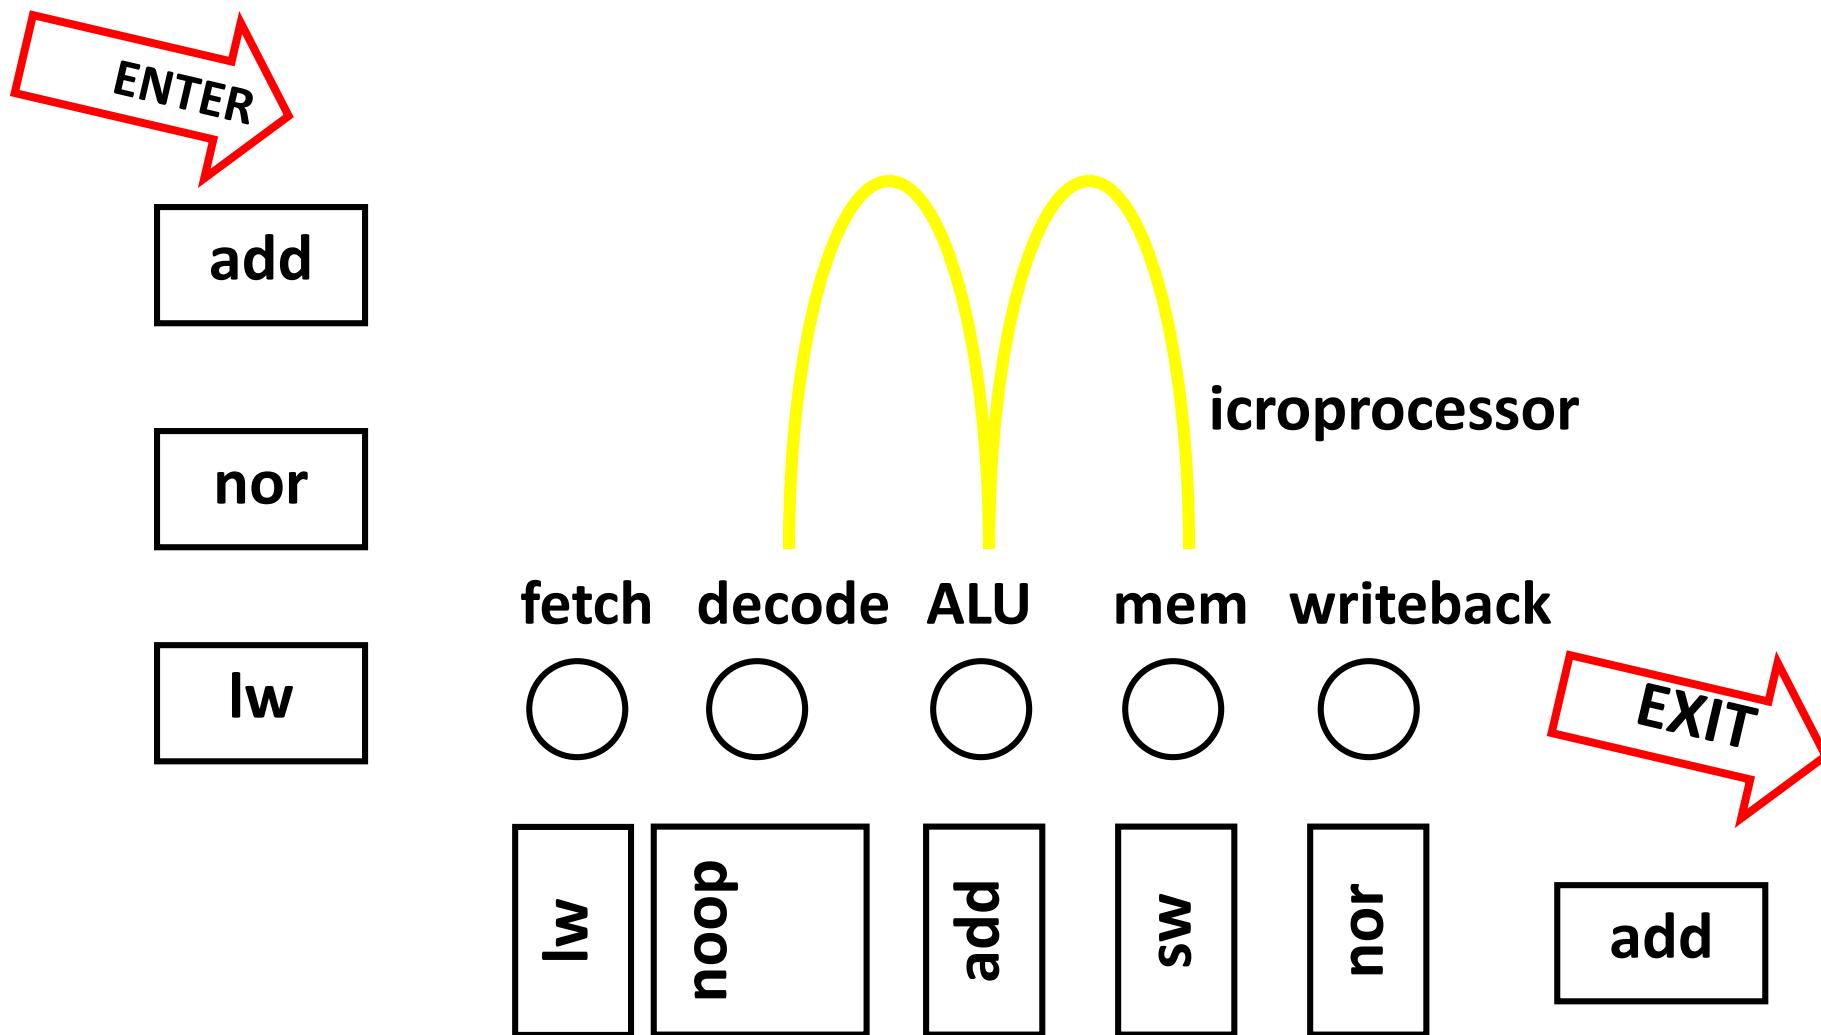
- Want to execute an instruction?
  - Build a processor (multi-cycle)
  - Find instructions
  - Line up instructions (1, 2, 3, ...)
  - Overlap execution
    - Cycle #1: Fetch 1
    - Cycle #2: Decode 1      Fetch 2
    - Cycle #3: ALU 1            Decode 2            Fetch 3
    - .....
  - This is called pipelining instruction execution.
  - Used extensively for the first time on IBM 360 (1960s).
  - CPI approaches 1.



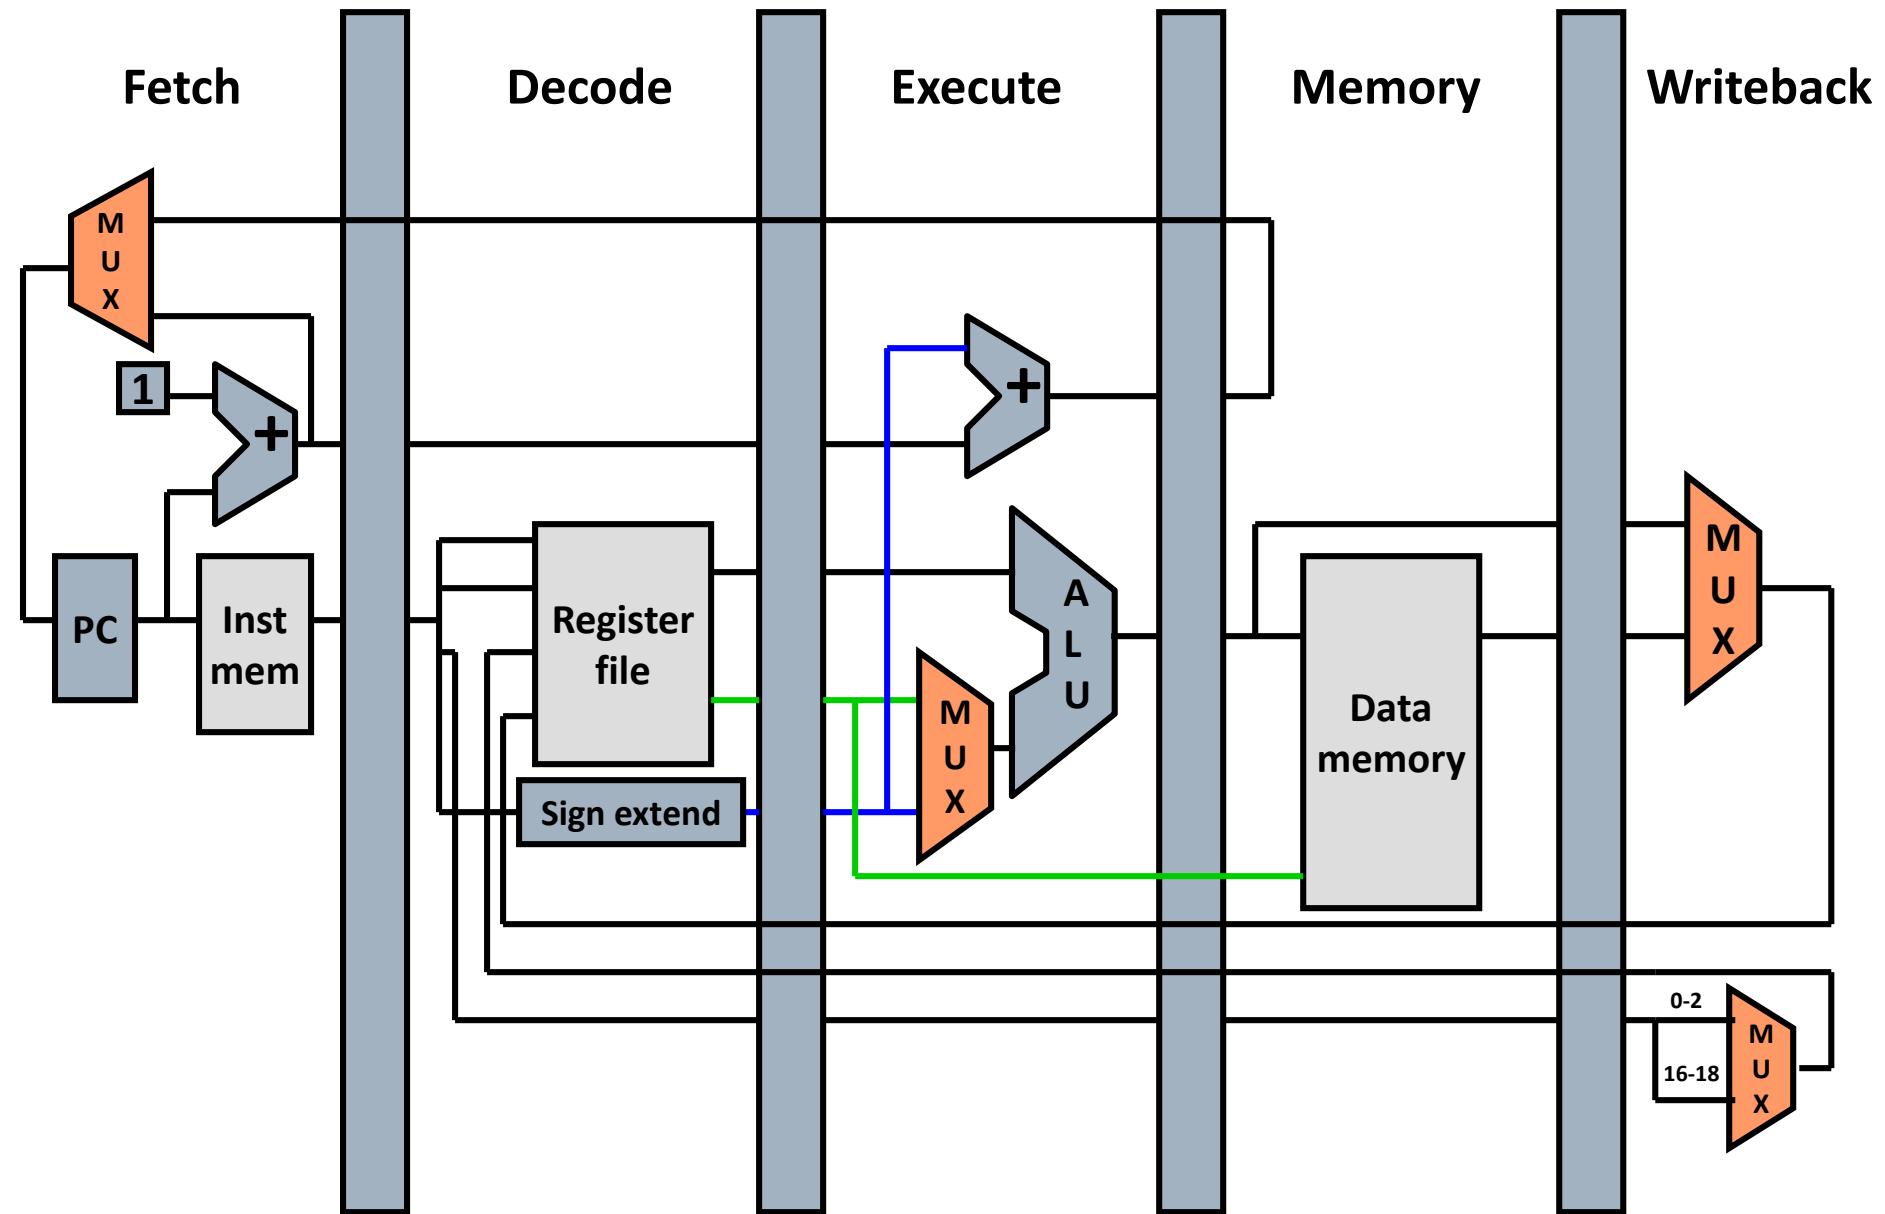
# Pipelining



# Pipelining



# Our new pipelined datapath

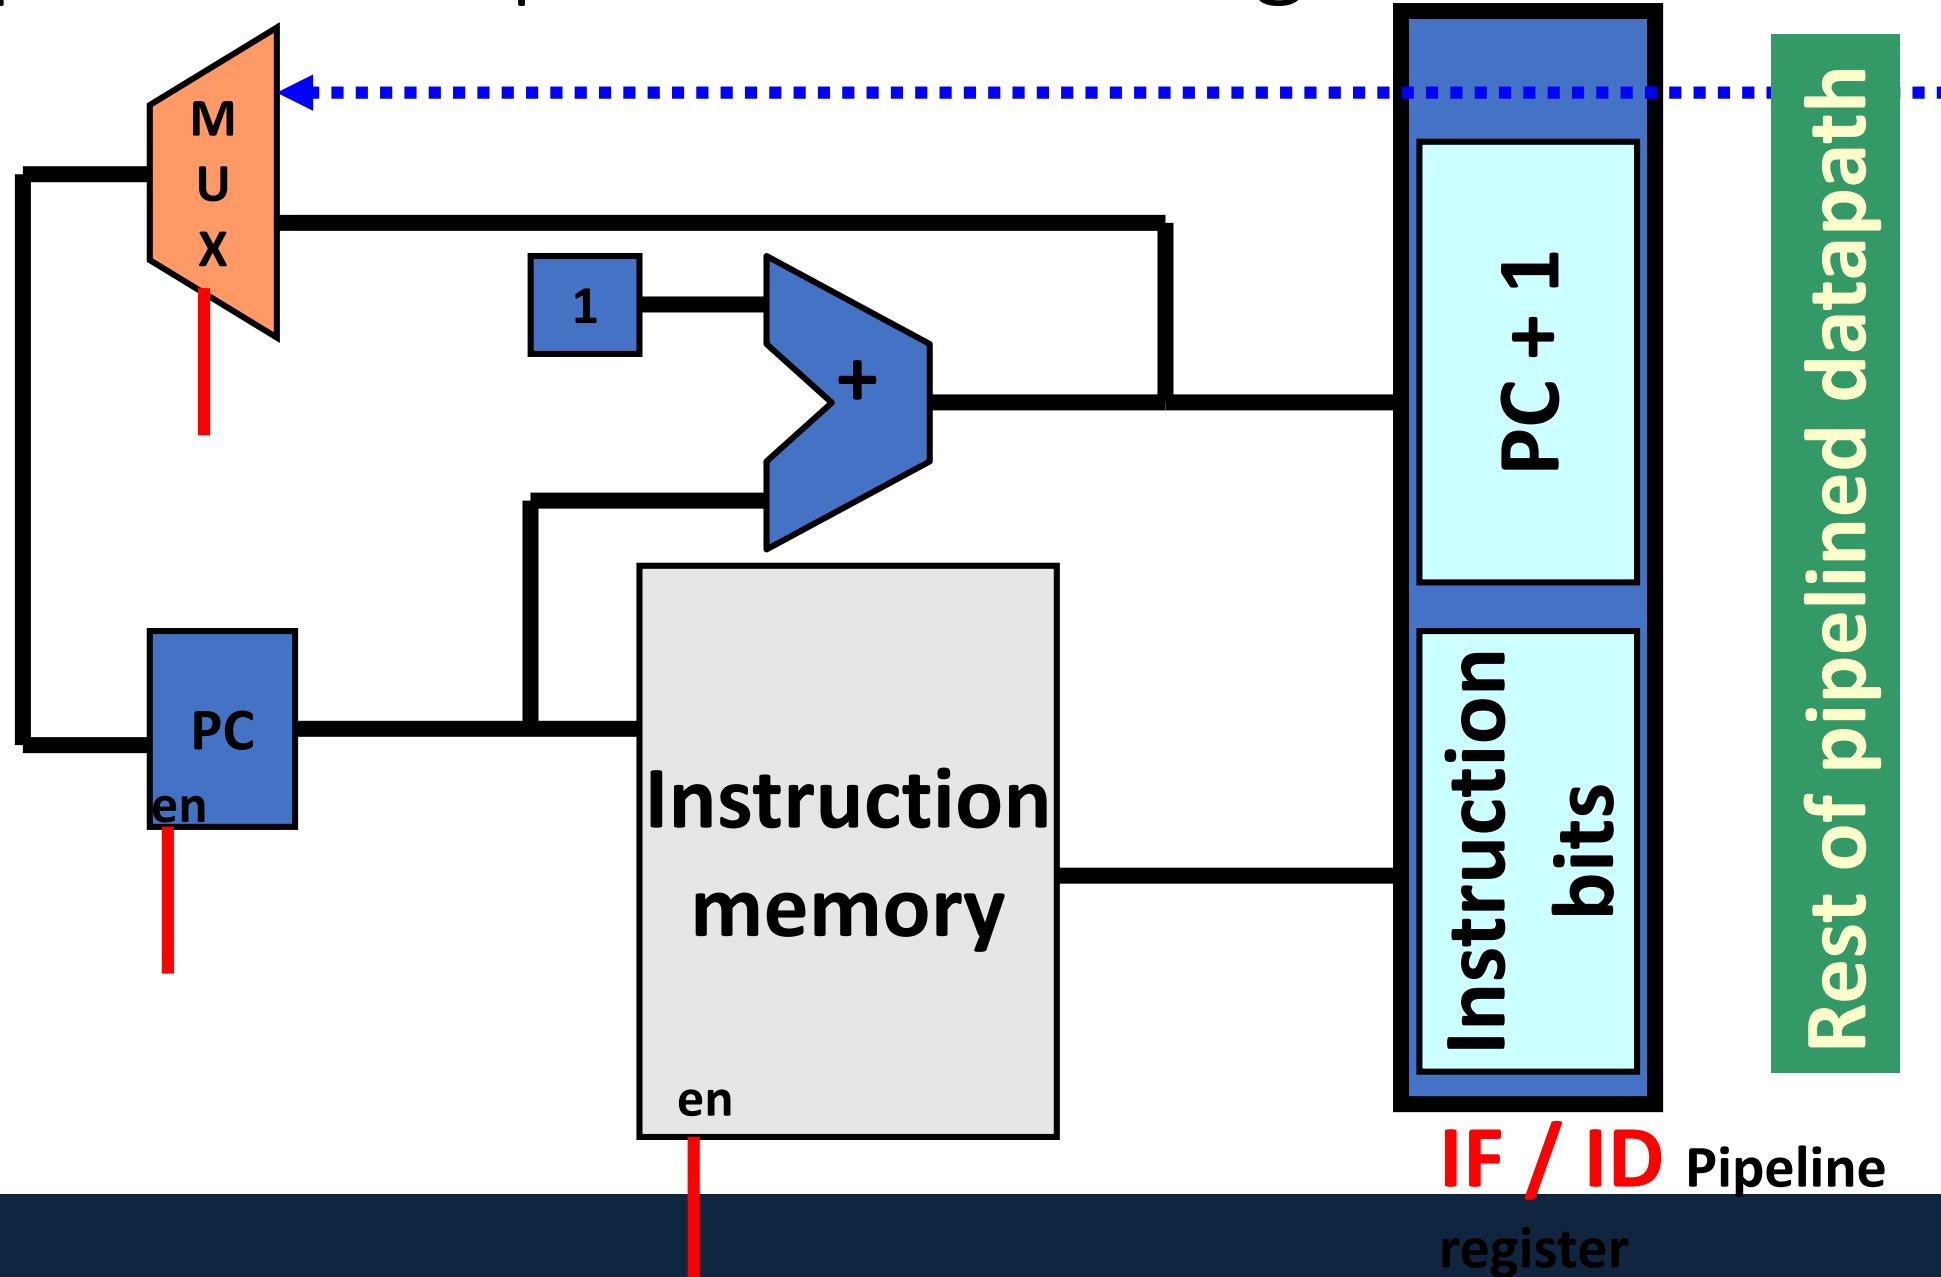


# Stage 1: Fetch

- Design a datapath that can fetch an instruction from memory every cycle.
  - Use PC to index memory to read instruction
  - Increment the PC (assume no branches for now)
- Write everything needed to complete execution to the **pipeline register (IF/ID)**
  - The next **stage** will read this pipeline register



# Pipeline datapath – Fetch stage

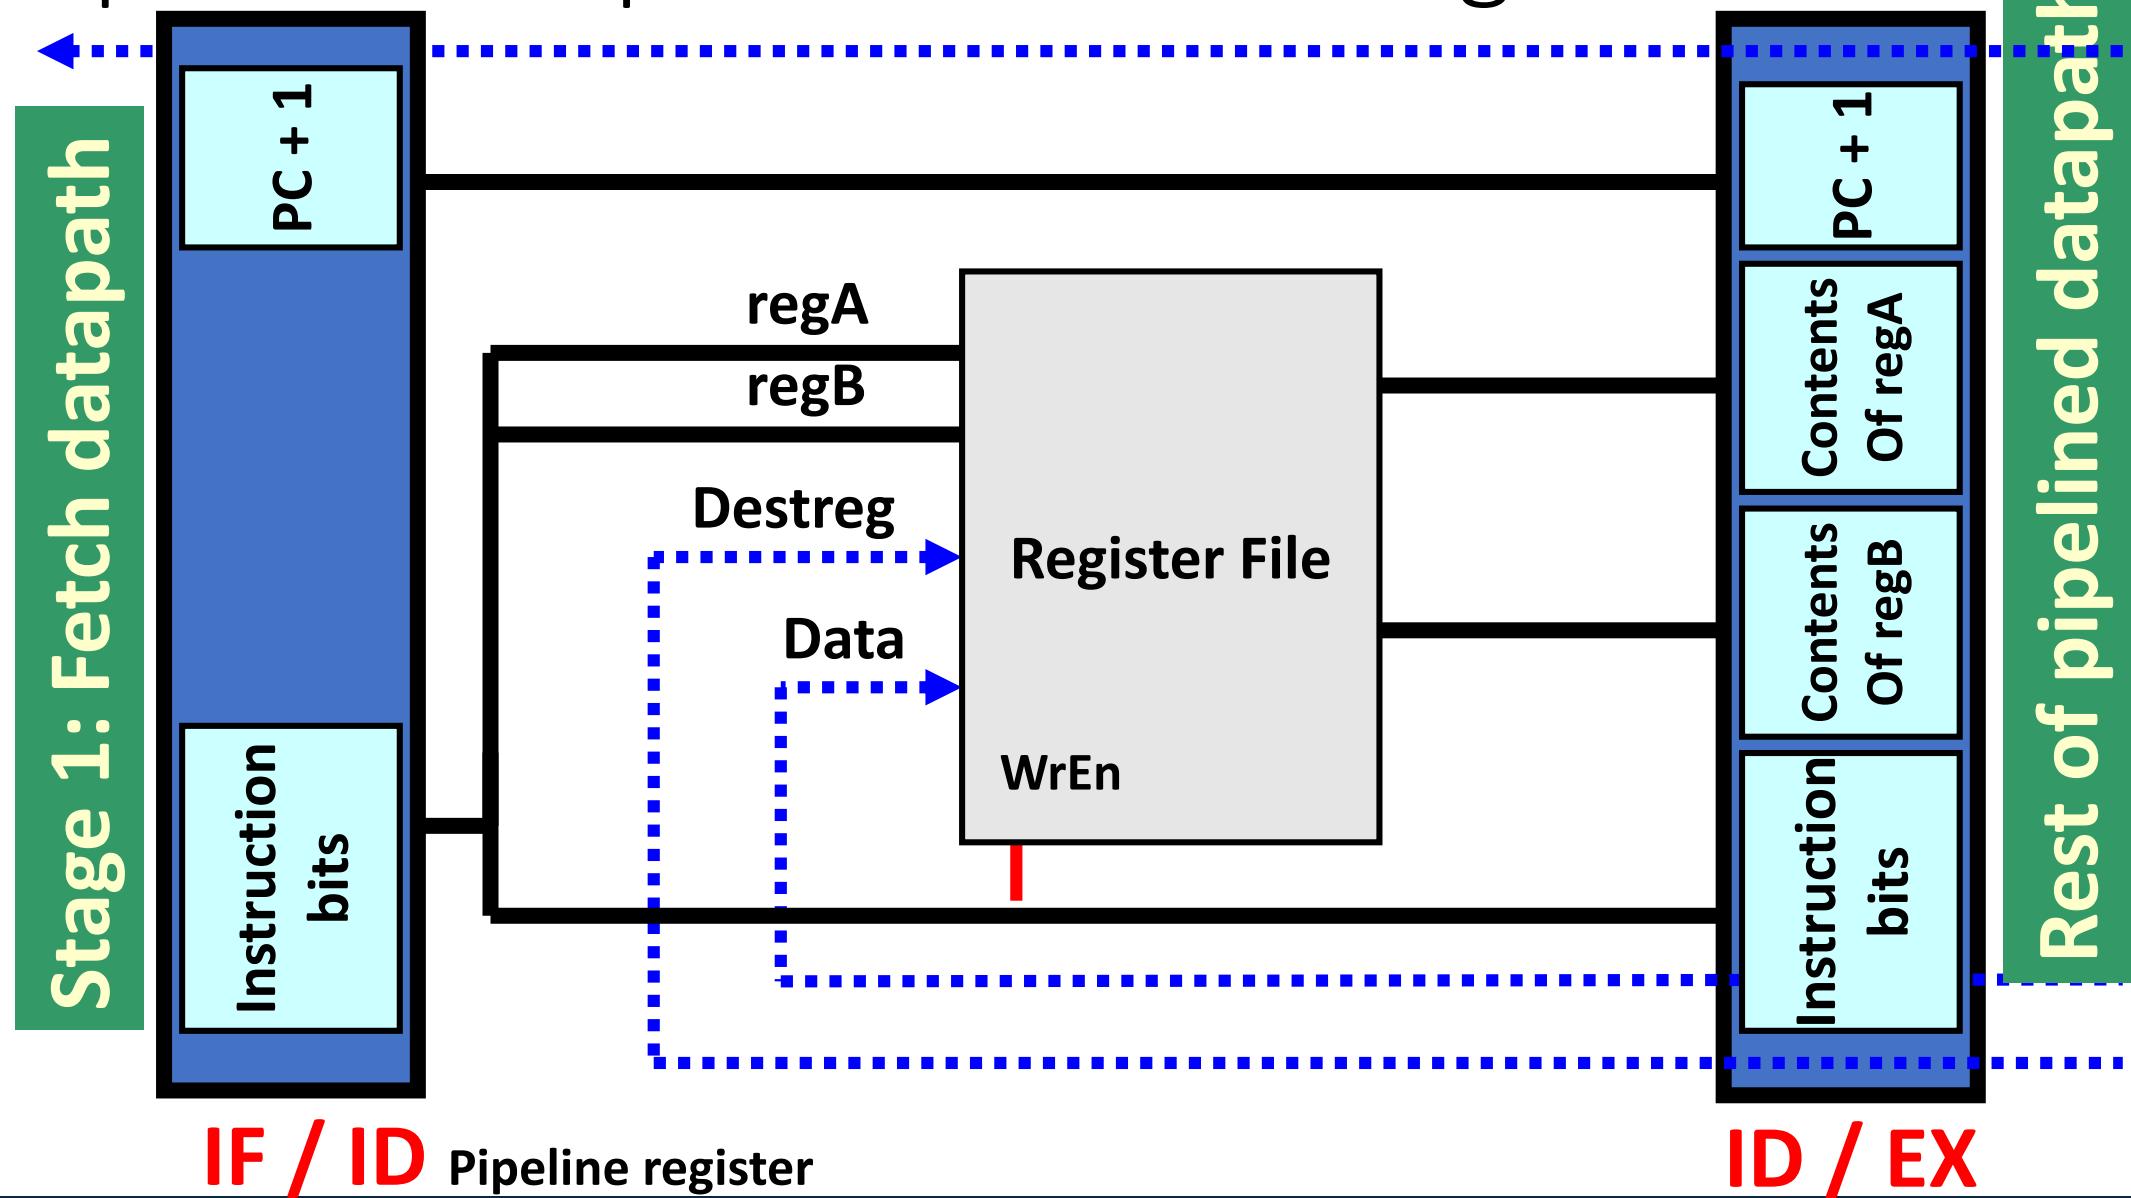


## Stage 2: Decode

- Design a datapath that reads the IF/ID pipeline register, decodes instruction and reads register file (specified by regA and regB of instruction bits).
  - Decode is easy, just pass on the opcode and let later stages figure out their own control signals for the instruction.
- Write everything needed to complete execution to the **pipeline register (ID/EX)**
  - Pass on the offset field and both destination register specifiers (or simply pass on the whole instruction!).
  - Including PC+1 even though decode didn't use it.



# Pipeline datapath – Decode stage

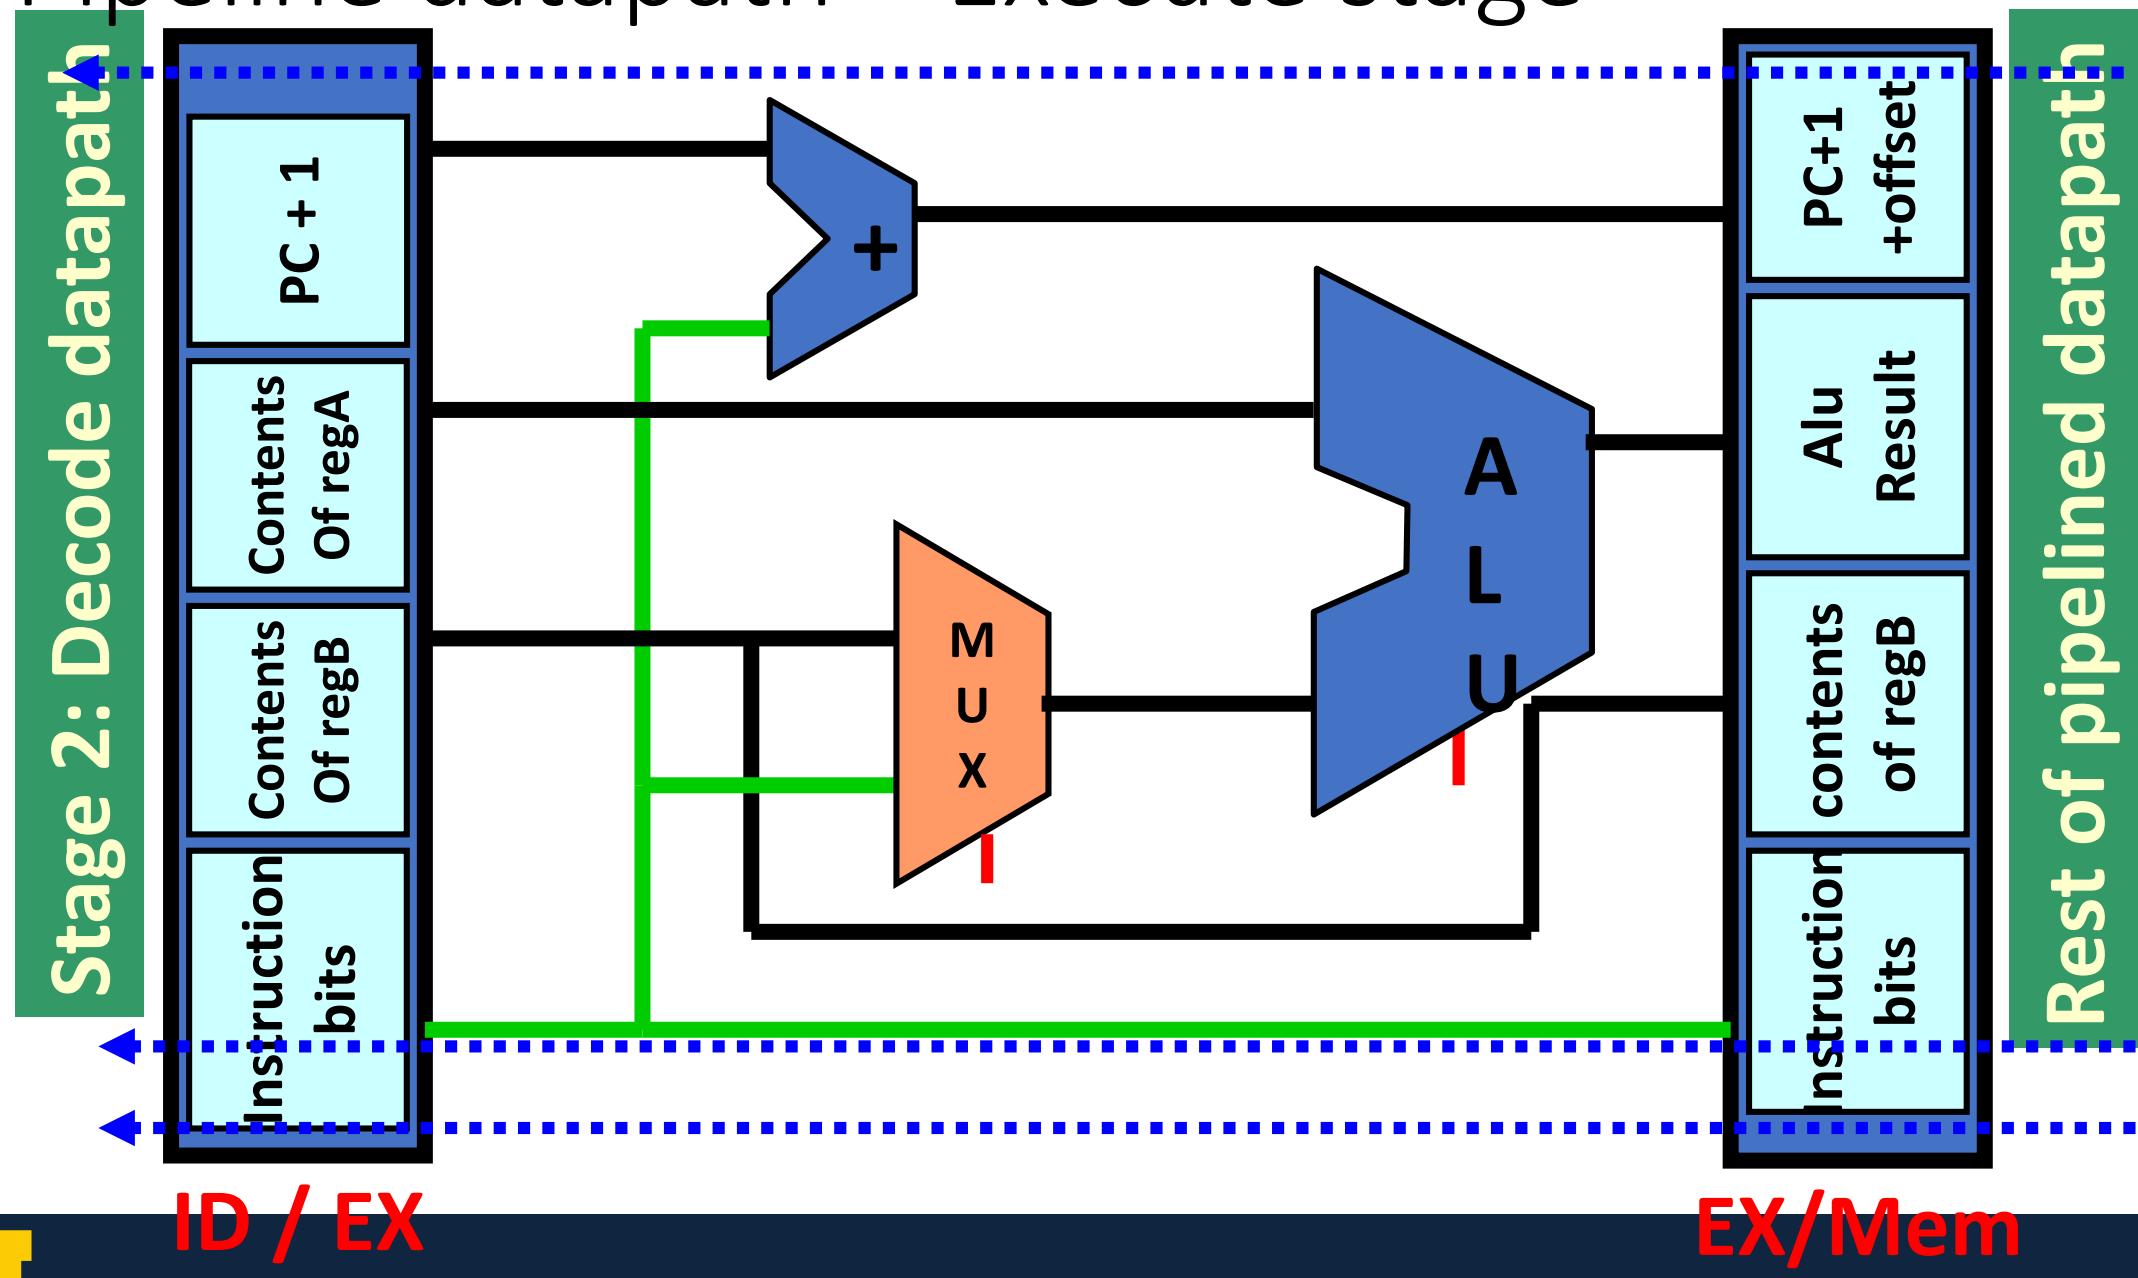


## Stage 3: Execute

- Design a datapath that performs the proper ALU operation for the instruction specified and the values present in the ID/EX pipeline register.
  - The inputs are the contents of regA and either the contents of regB or the offset field on the instruction.
  - Also, calculate PC+1+offset in case this is a branch.
- Write everything needed to complete execution to the **pipeline register (EX/Mem)**
  - ALU result, contents of regB and PC+1+offset
  - Instruction bits for opcode and destReg specifiers
  - Result from comparison of regA and regB contents



# Pipeline datapath – Execute stage

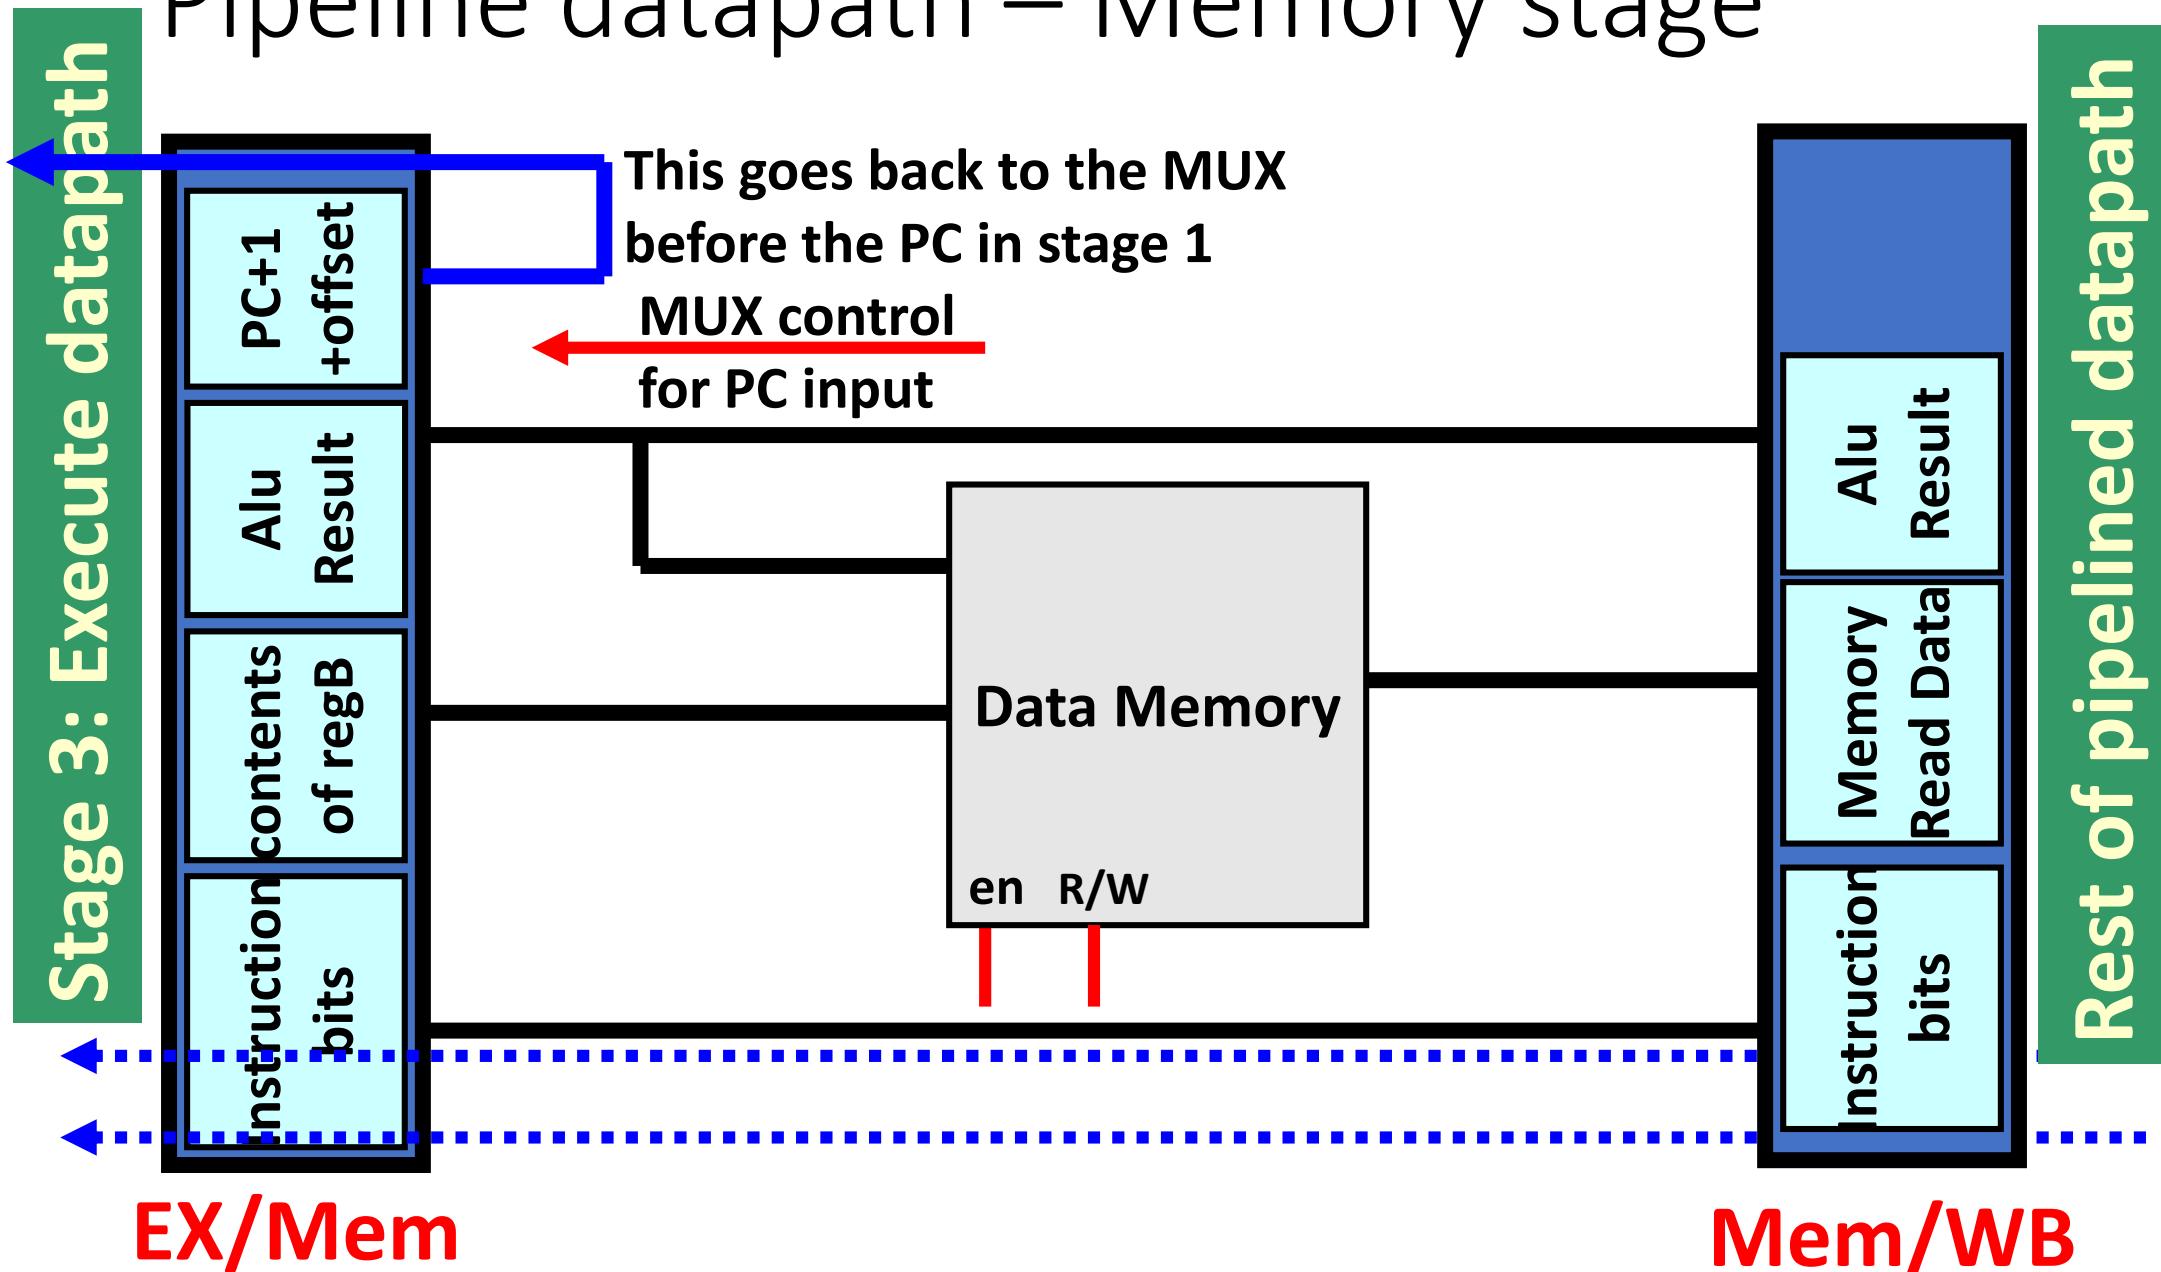


# Stage 4: Memory Operation

- Design a datapath that performs the proper memory operation for the instruction specified and the values present in the EX/Mem pipeline register.
  - ALU result contains address for **ld** and **st** instructions.
  - Opcode bits control memory R/W and enable signals.
- Write everything needed to complete execution to the **pipeline register (Mem/WB)**
  - ALU result and MemData
  - Instruction bits for opcode and destReg specifiers



# Pipeline datapath – Memory stage

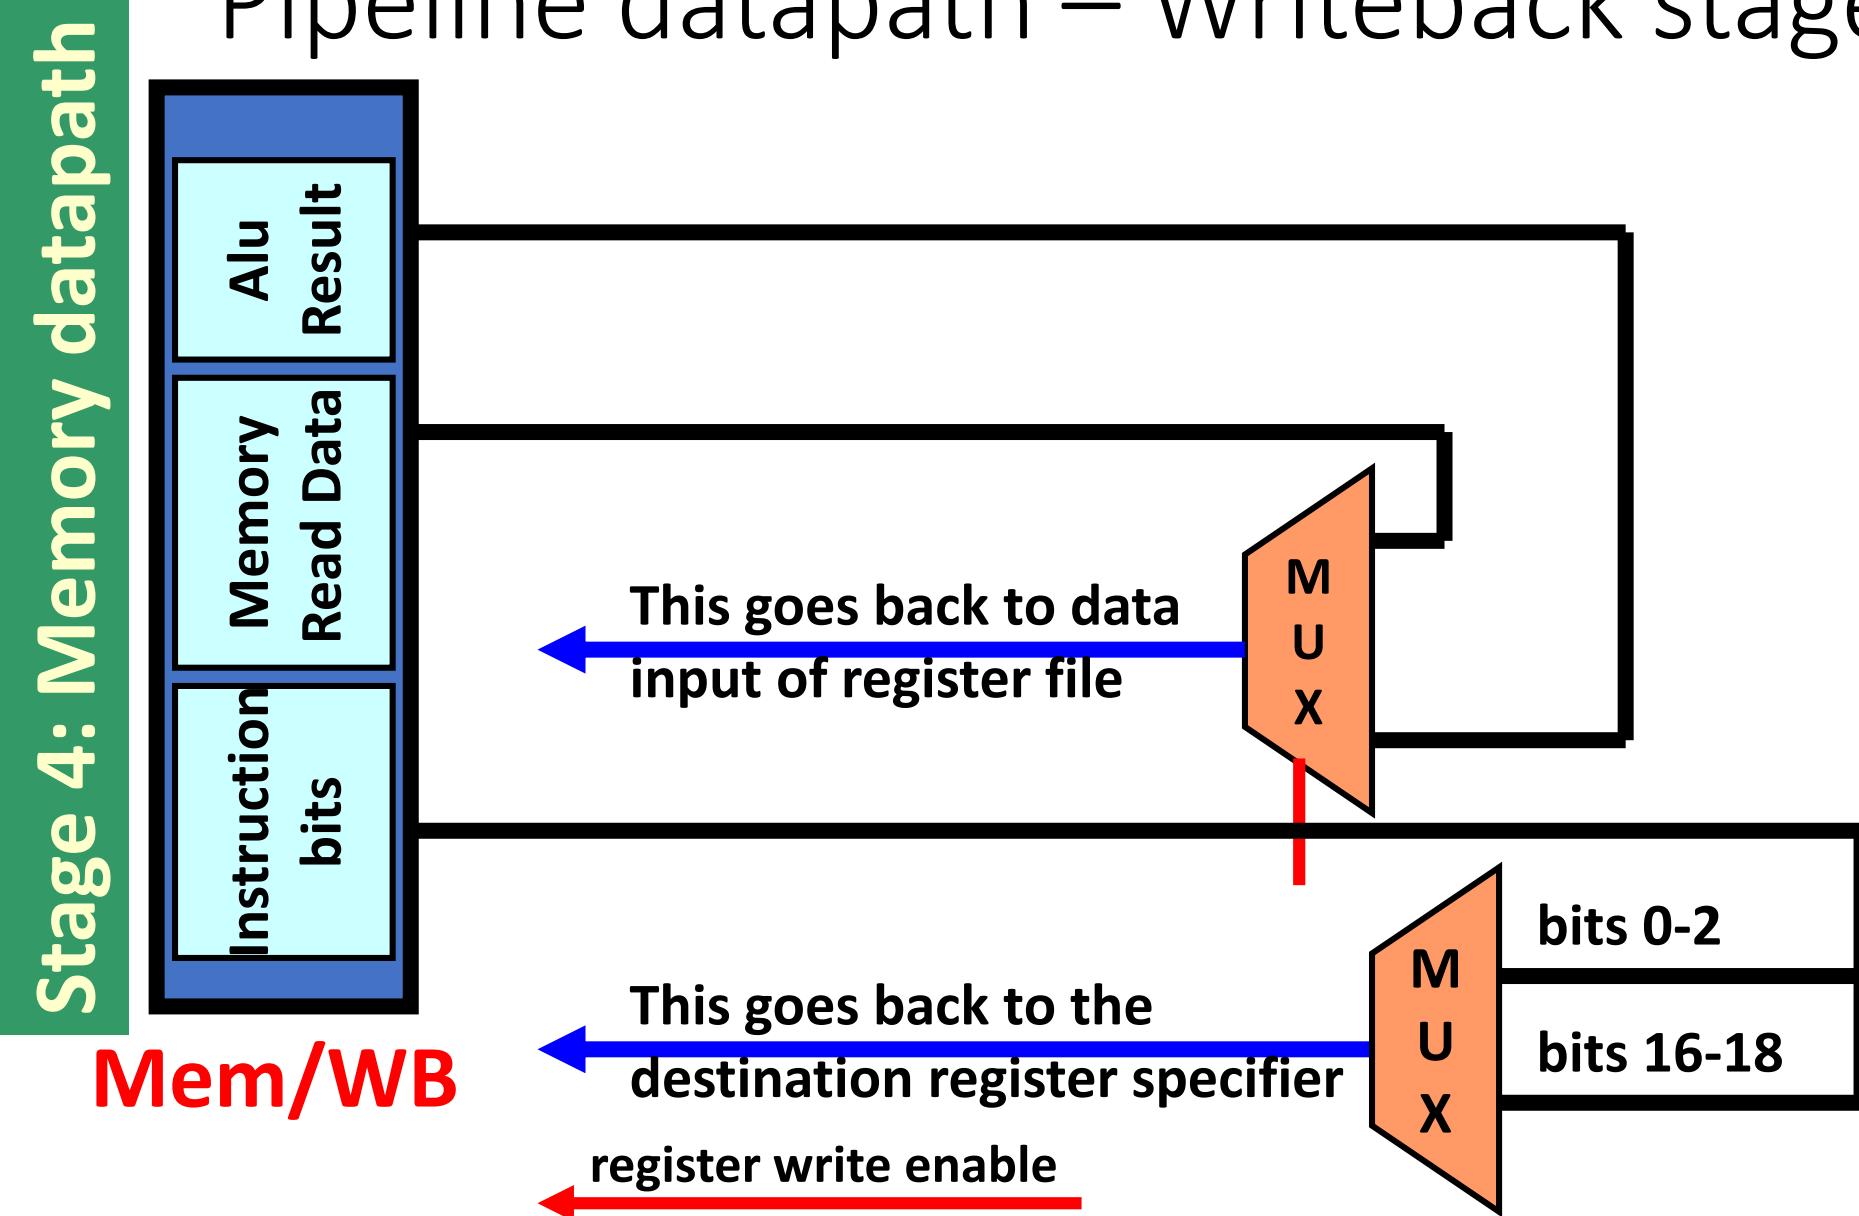


## Stage 5: Write back

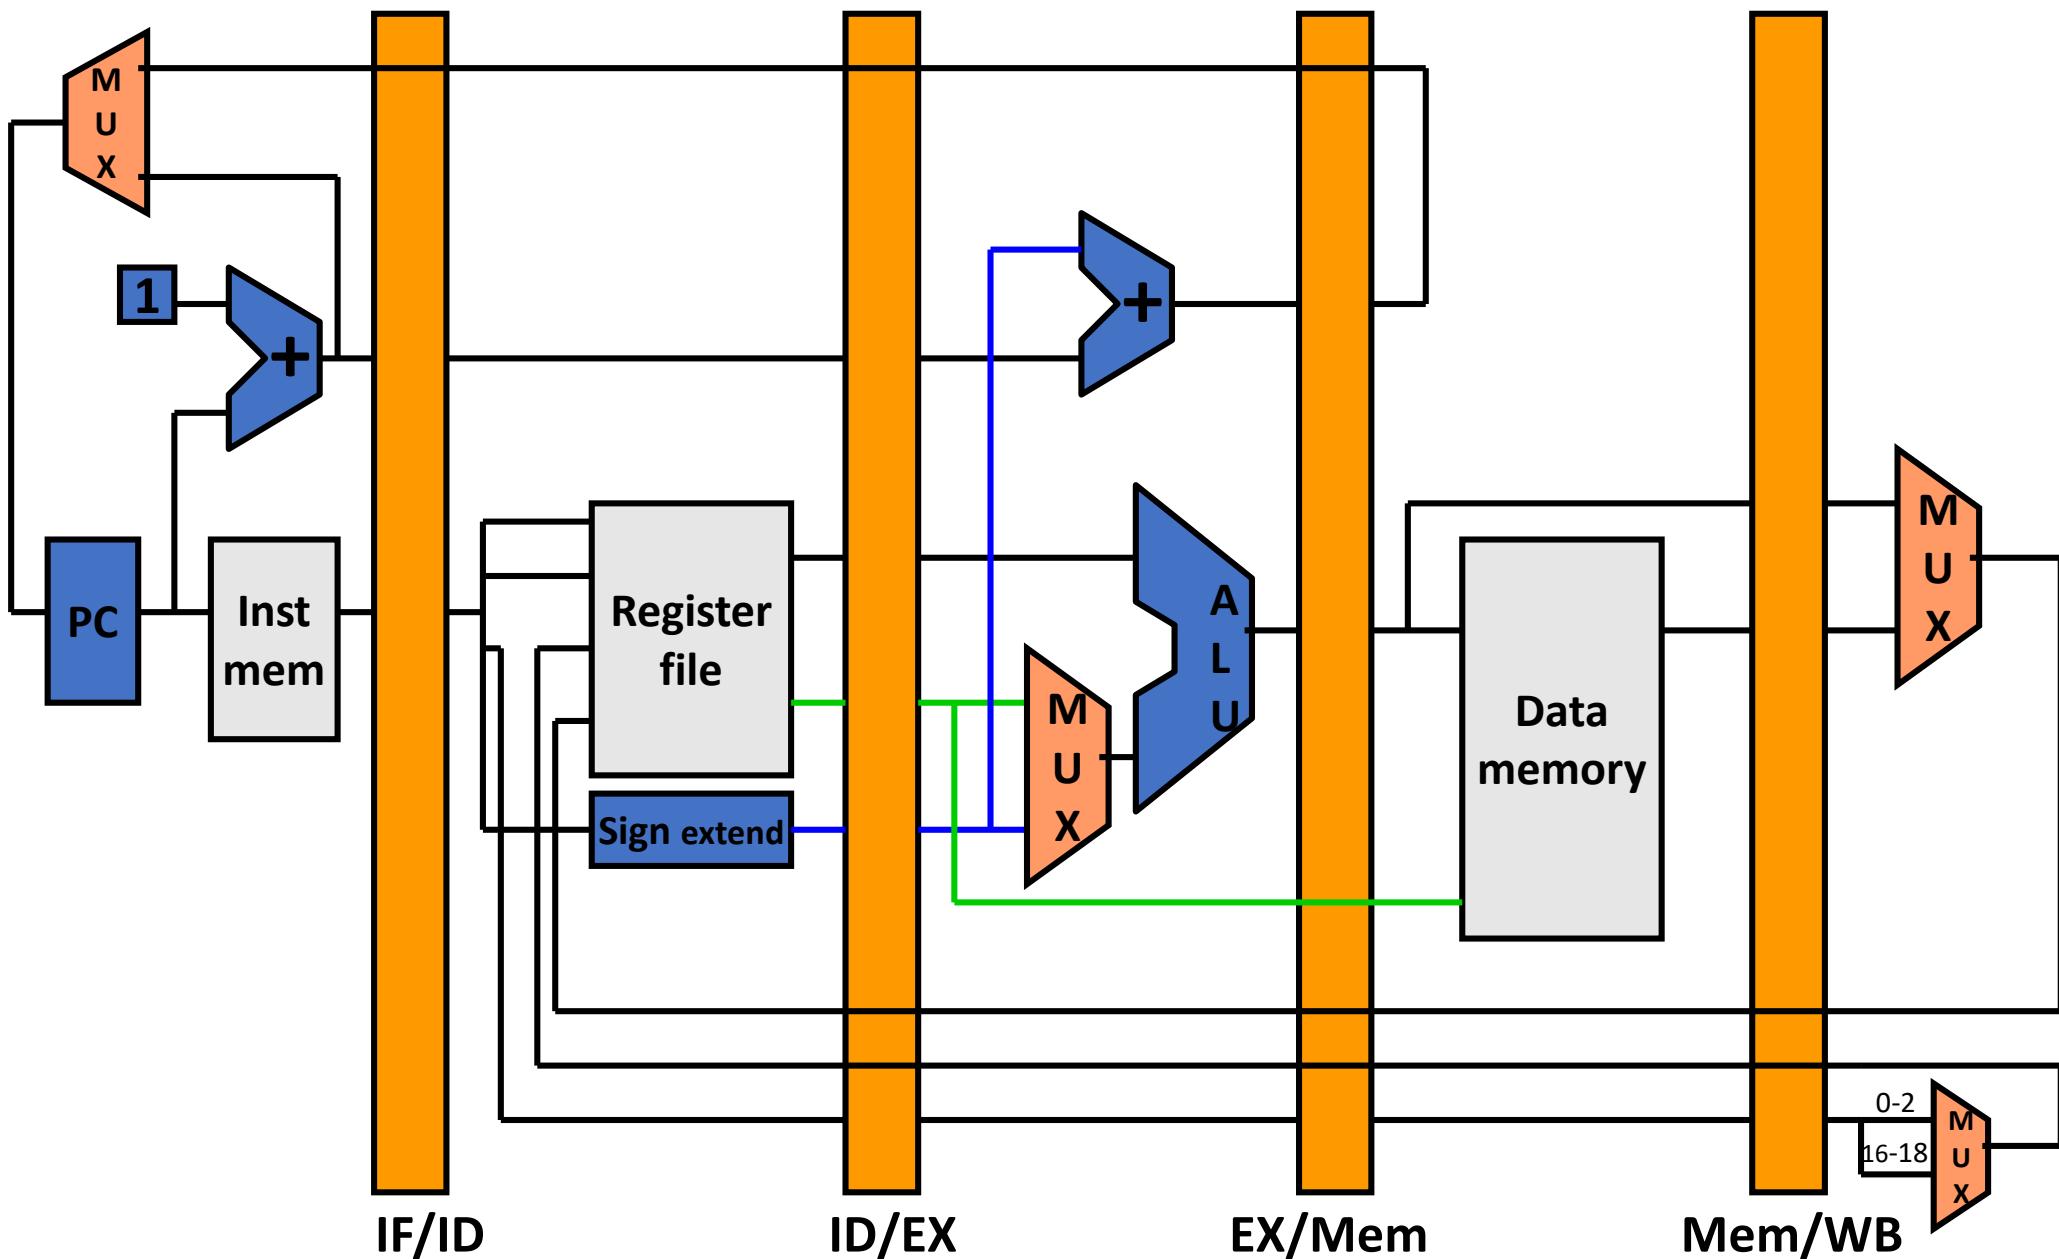
- Design a datapath that completes the execution of this instruction, writing to the register file if required.
  - Write MemData to destReg for Id instruction
  - Write ALU result to destReg for add or nor instructions.
  - Opcode bits also control register write enable signal.



# Pipeline datapath – Writeback stage



# Putting all together



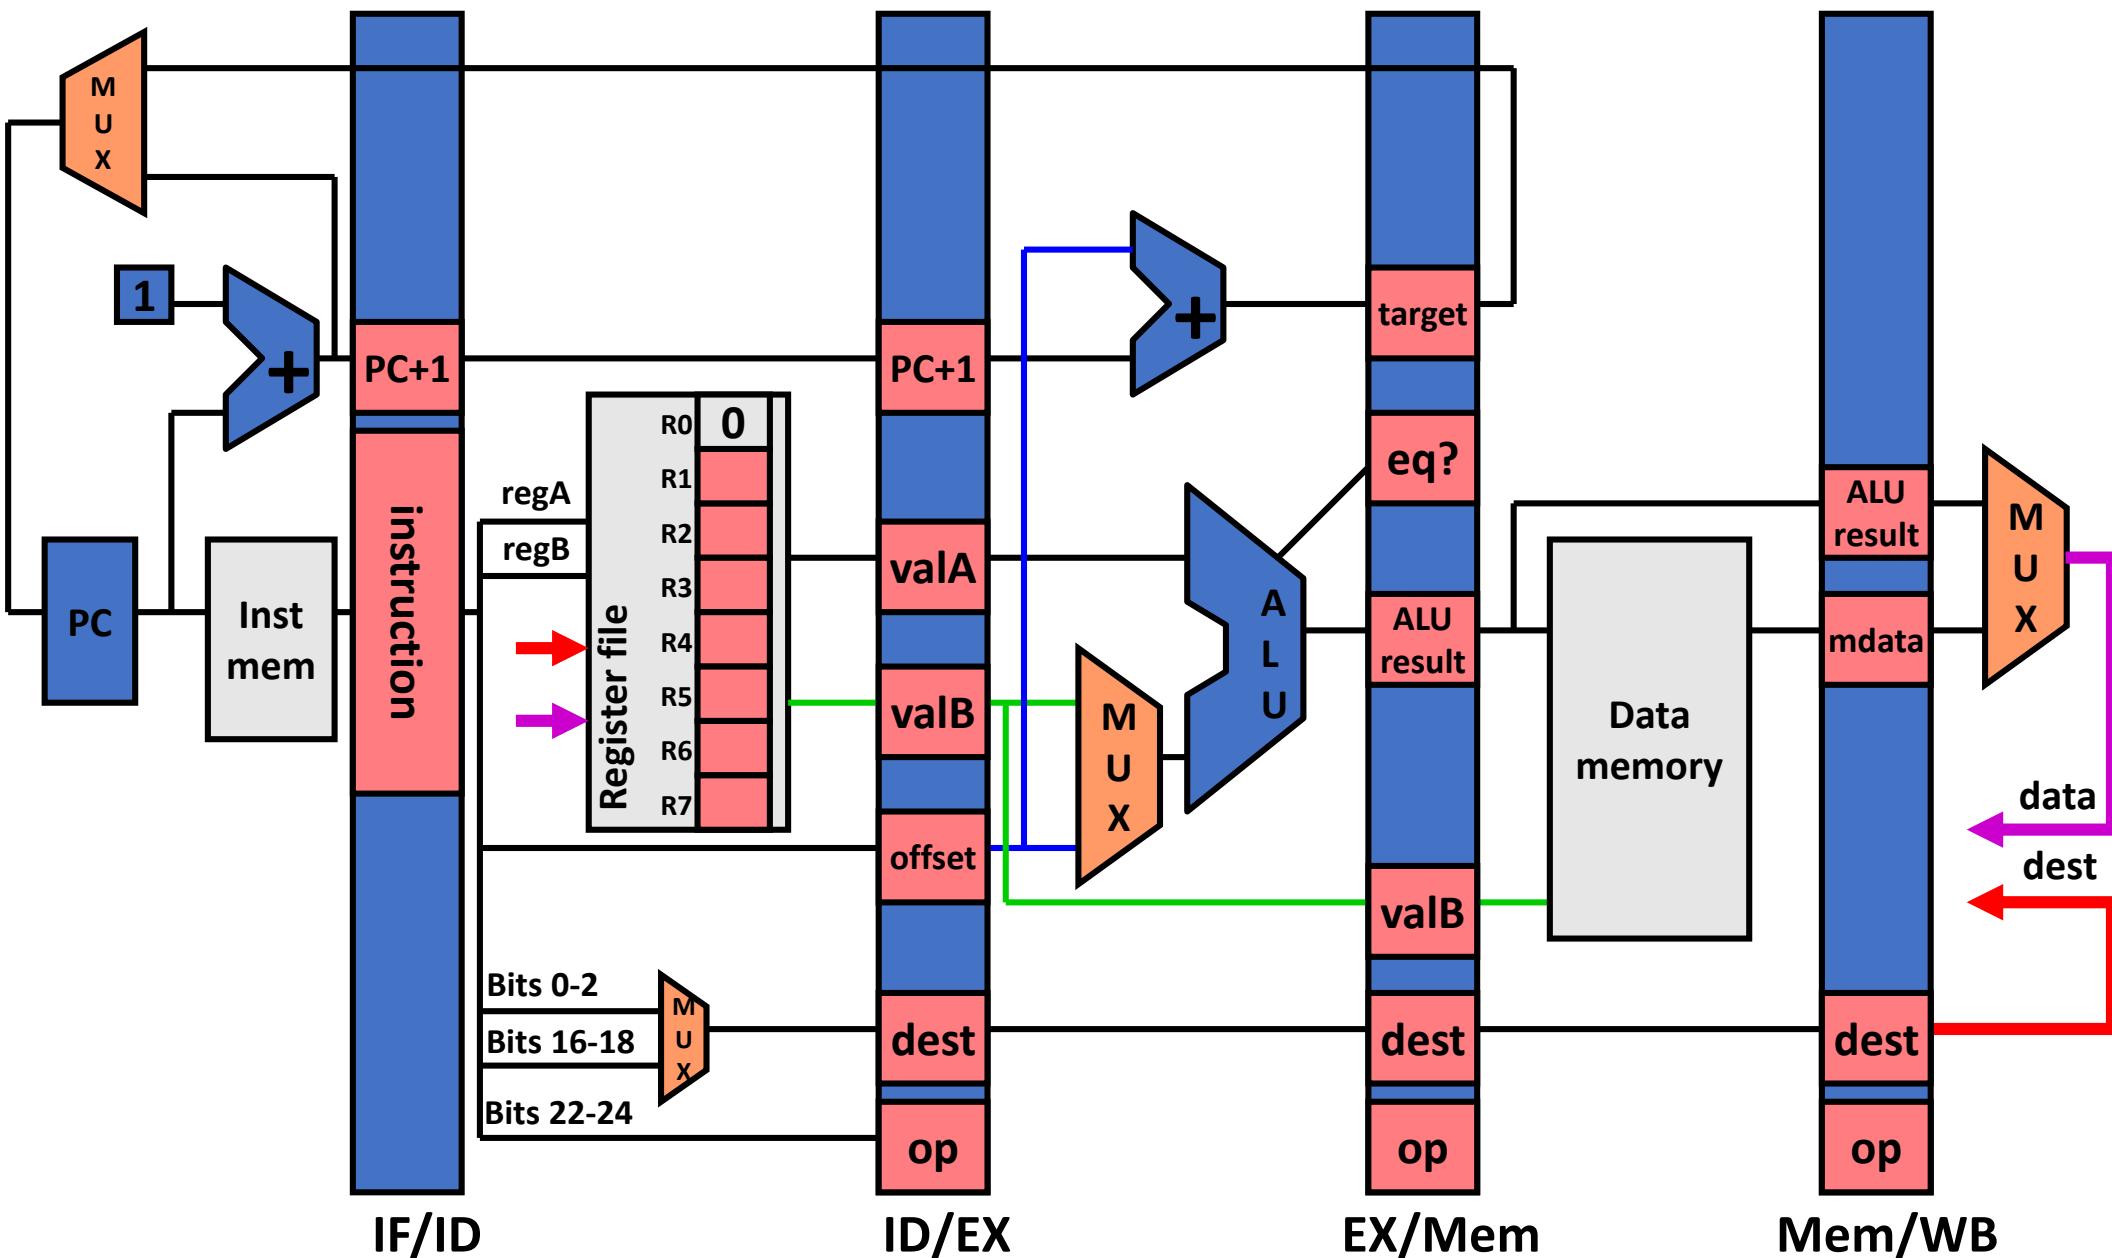
# Sample Code (Simple)

Let's run the following code on pipelined LC2K:

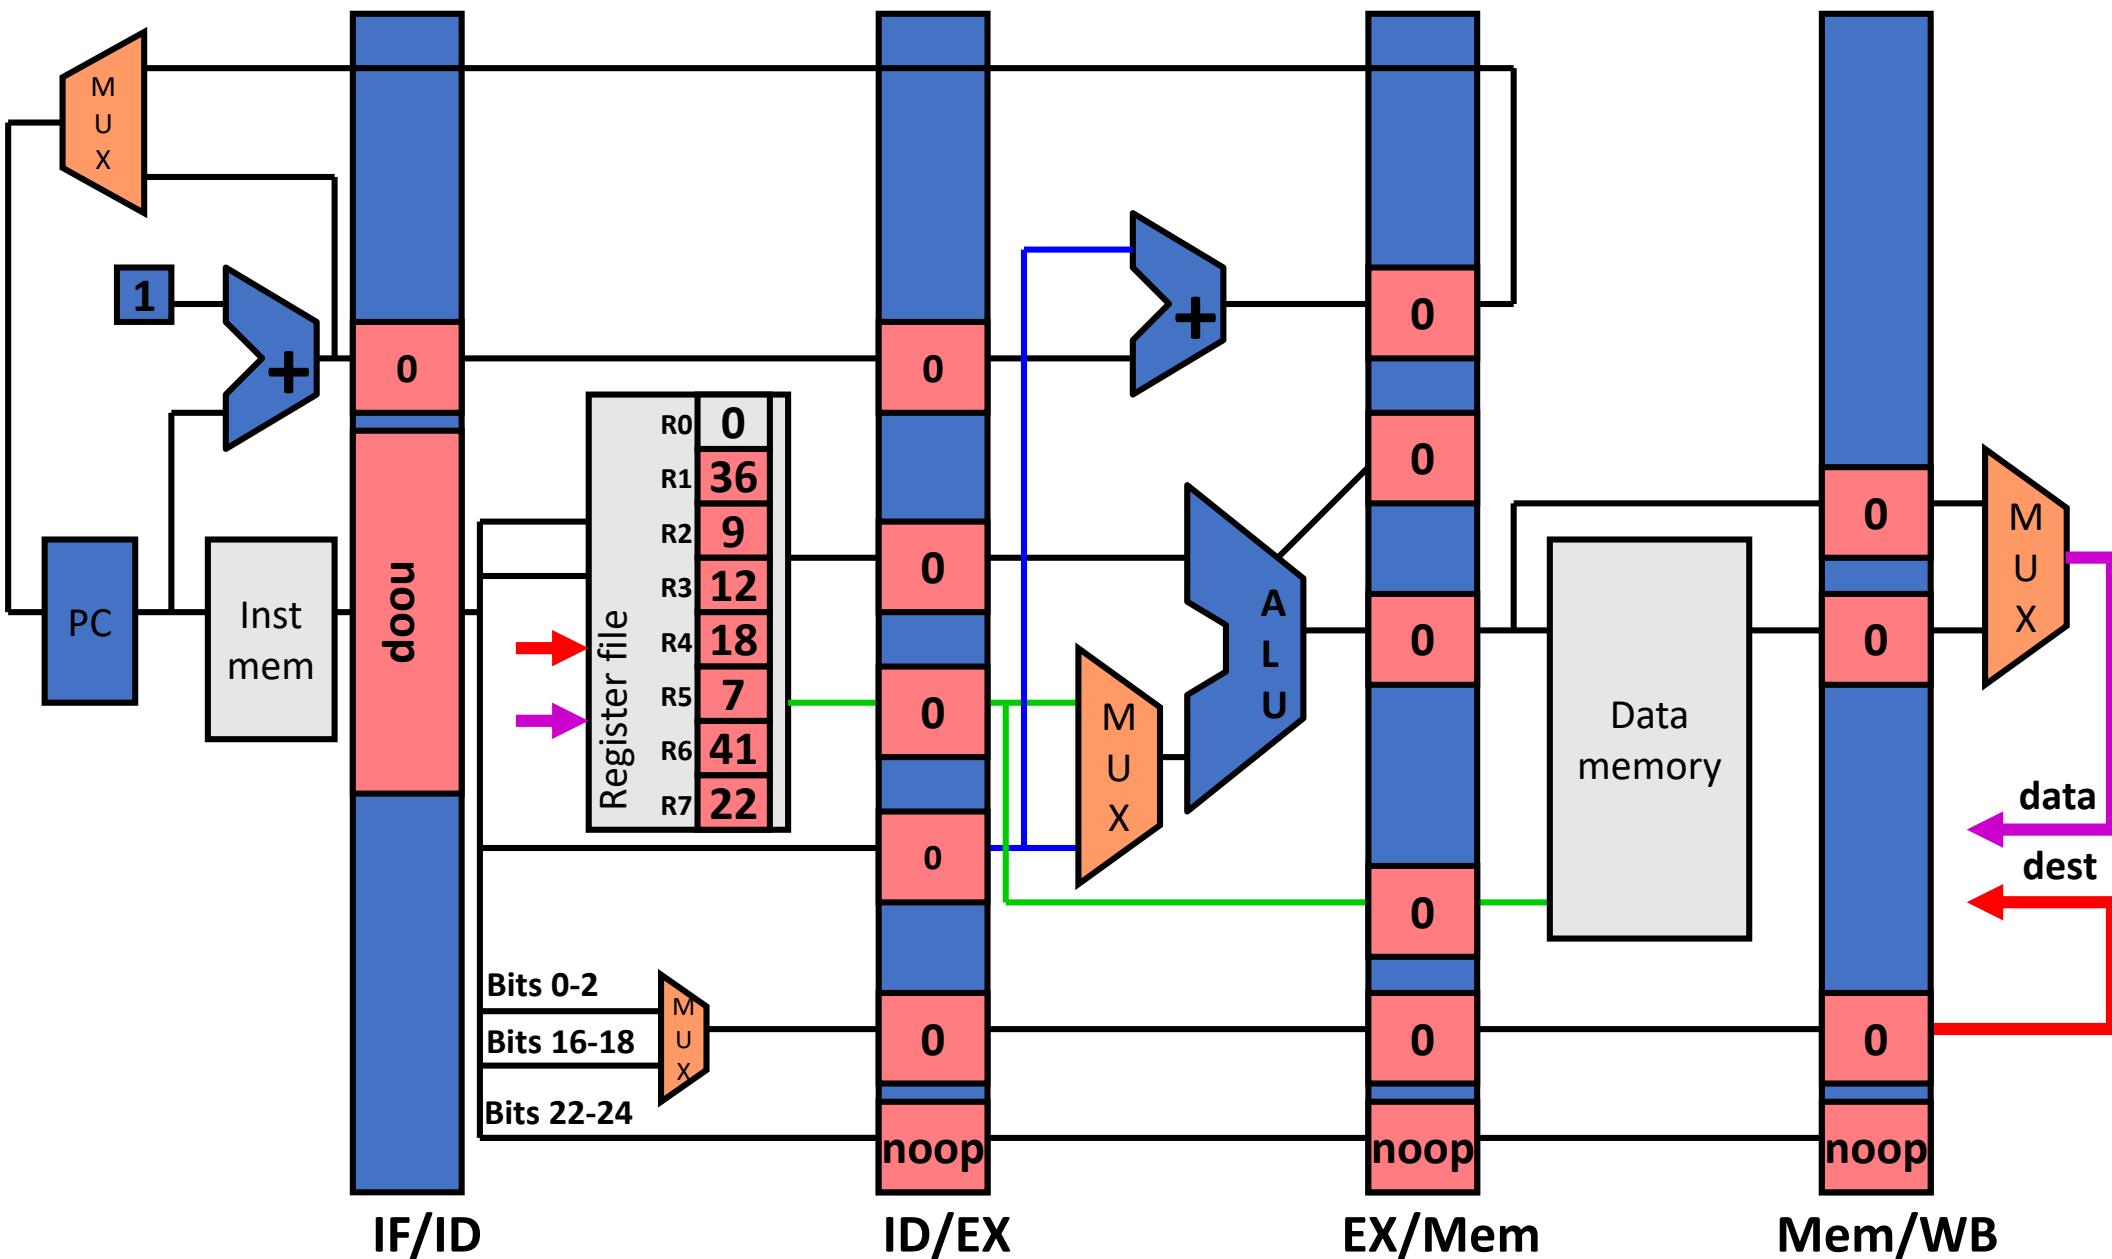
- add 1 2 3 ; reg 3 = reg 1 + reg 2
- nor 4 5 6 ; reg 6 = reg 4 nor reg 5
- lw 2 4 20 ; reg 4 = Mem[reg2+20]
- add 2 5 5 ; reg 5 = reg 2 + reg 5
- sw 3 7 10 ; Mem[reg3+10] =reg 7



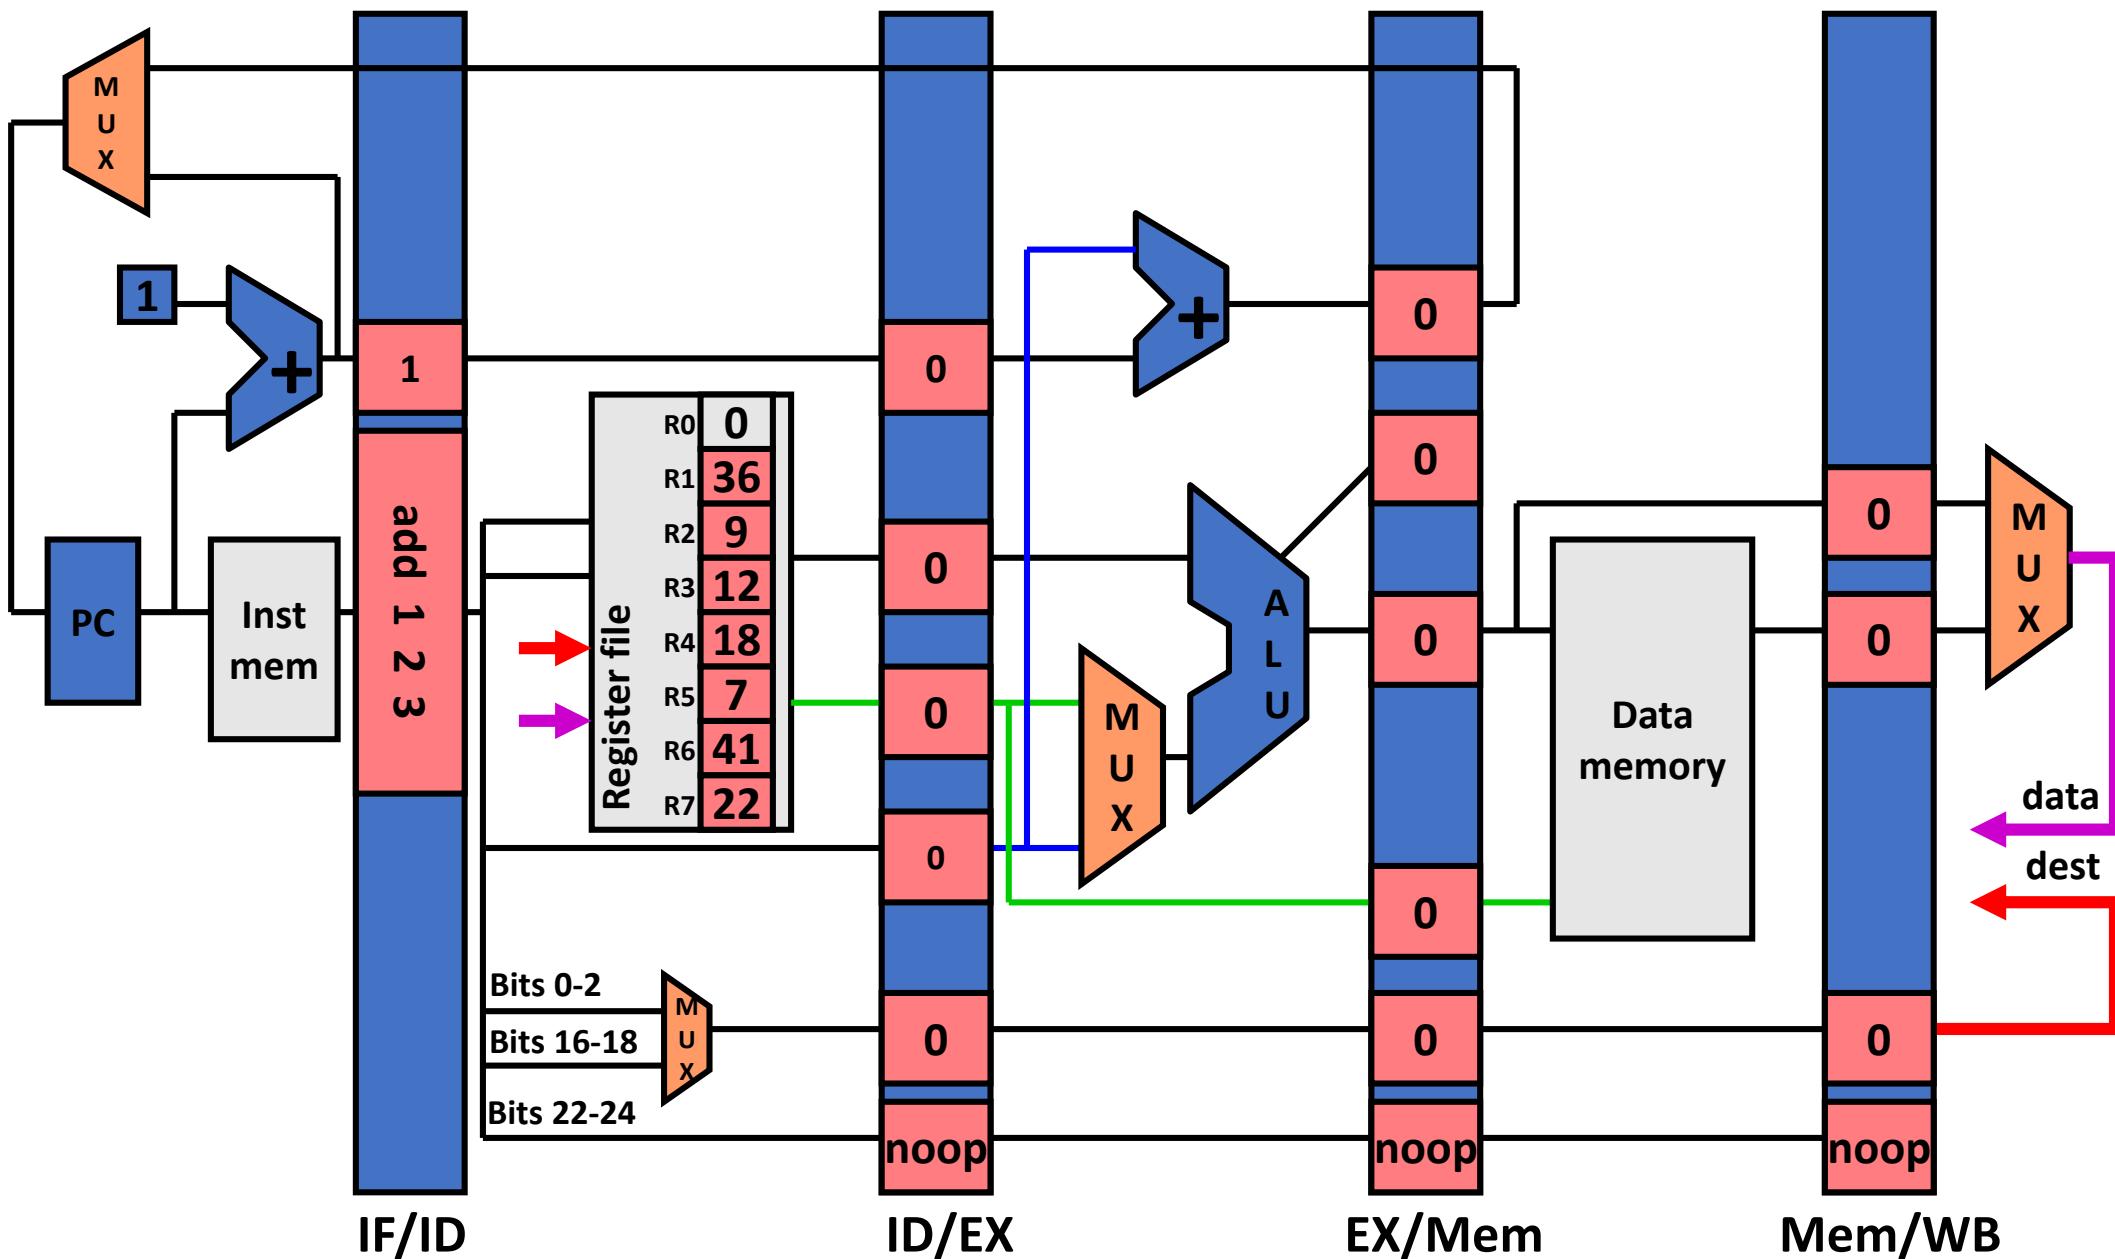
# Pipeline datapath



# Time 0 - Initial state

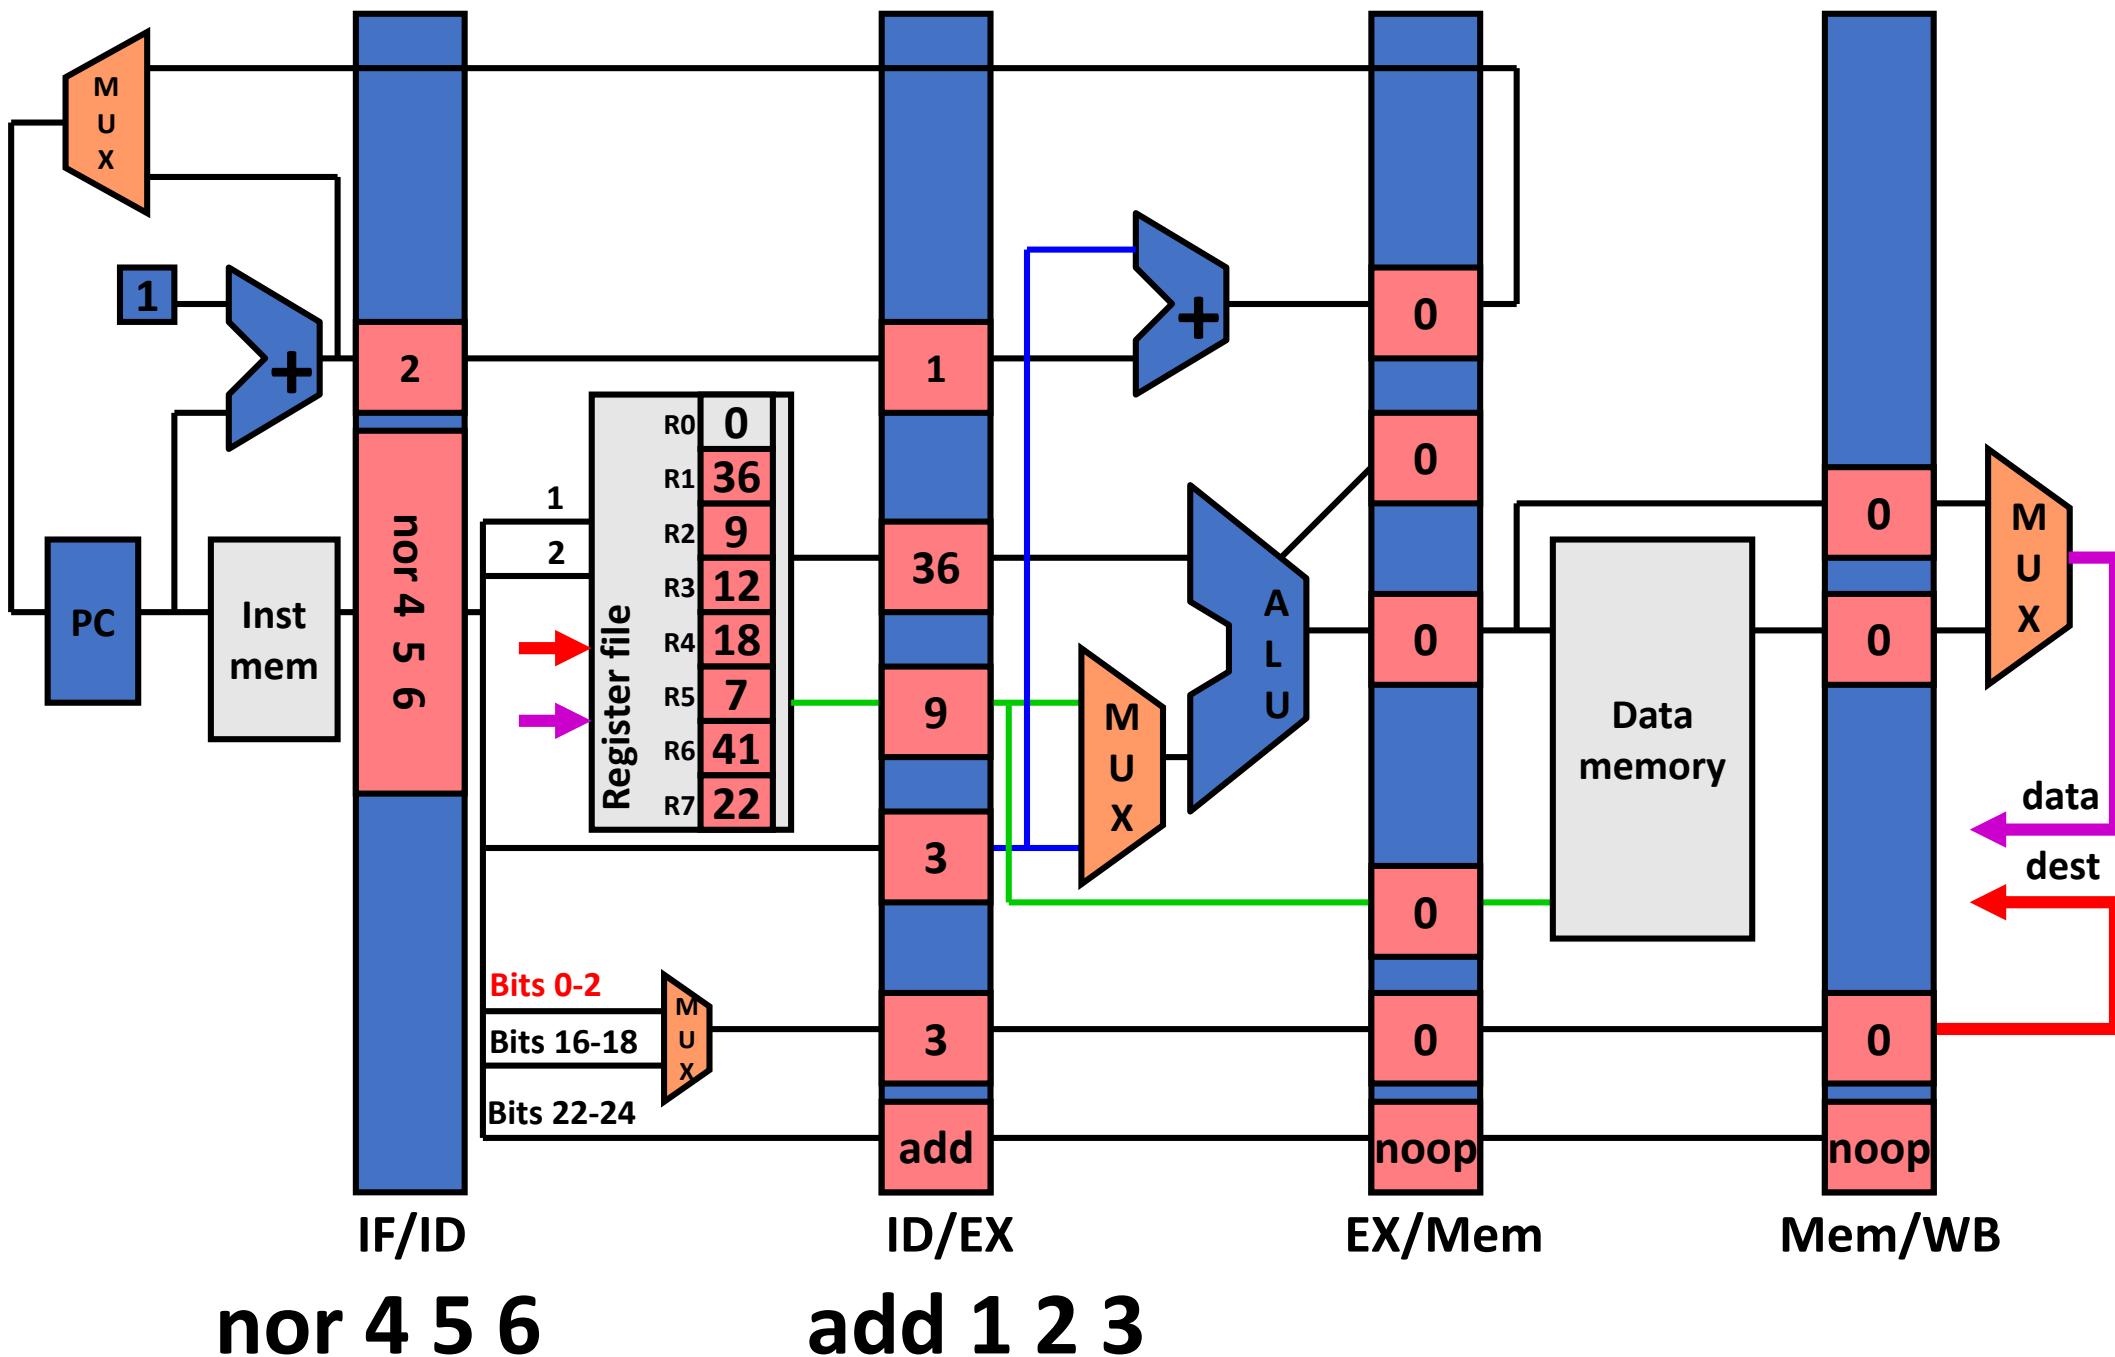


# Time 1 - Fetch: add 1 2 3

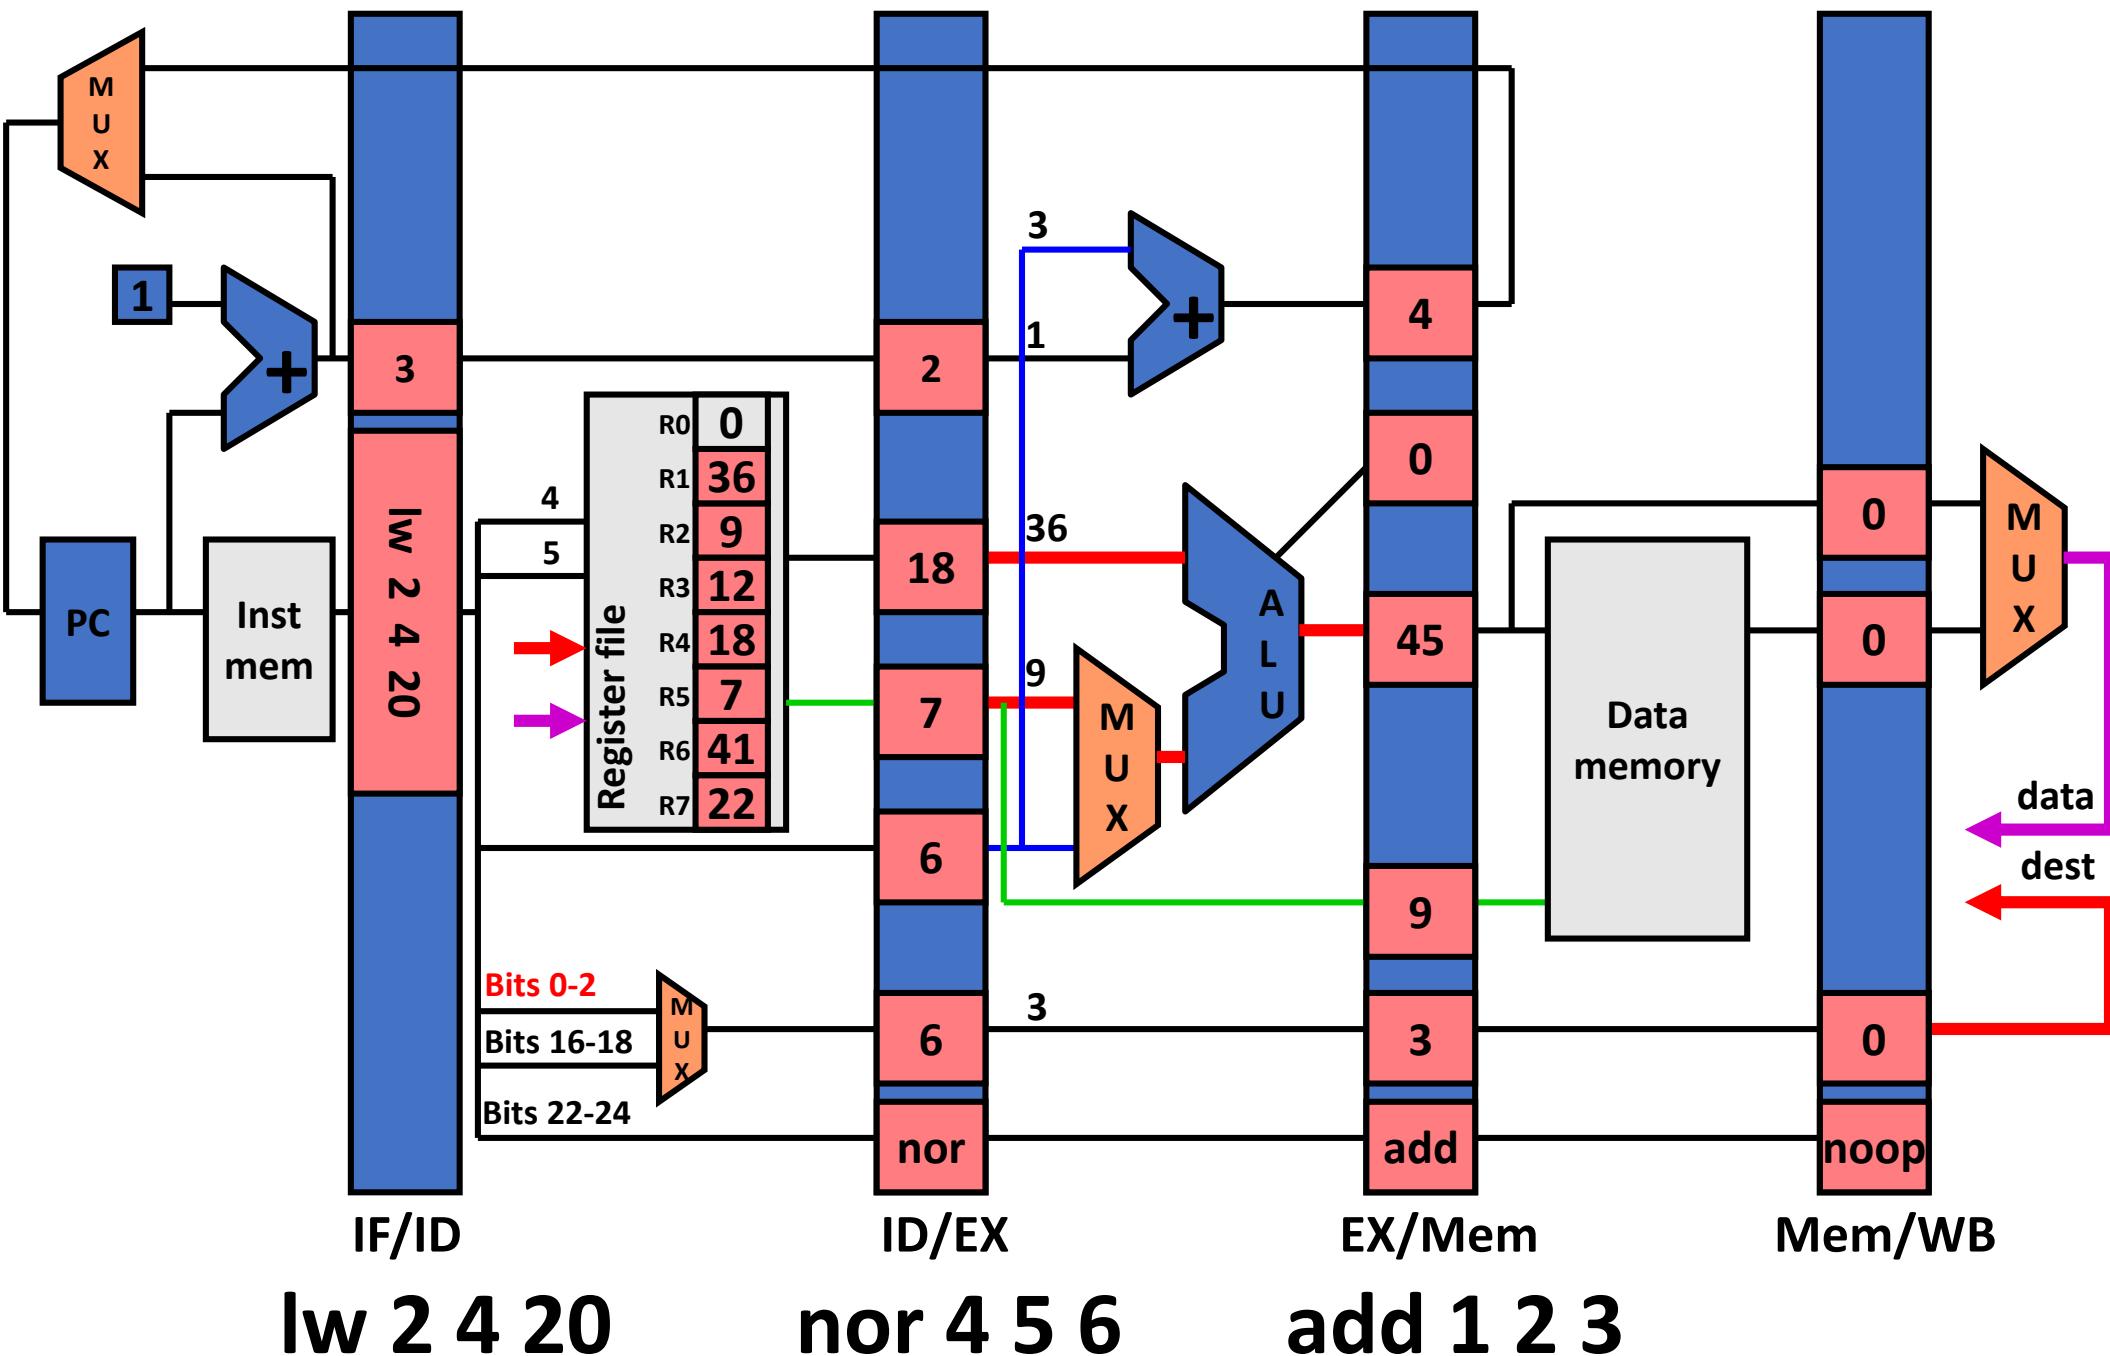


**add 1 2 3**

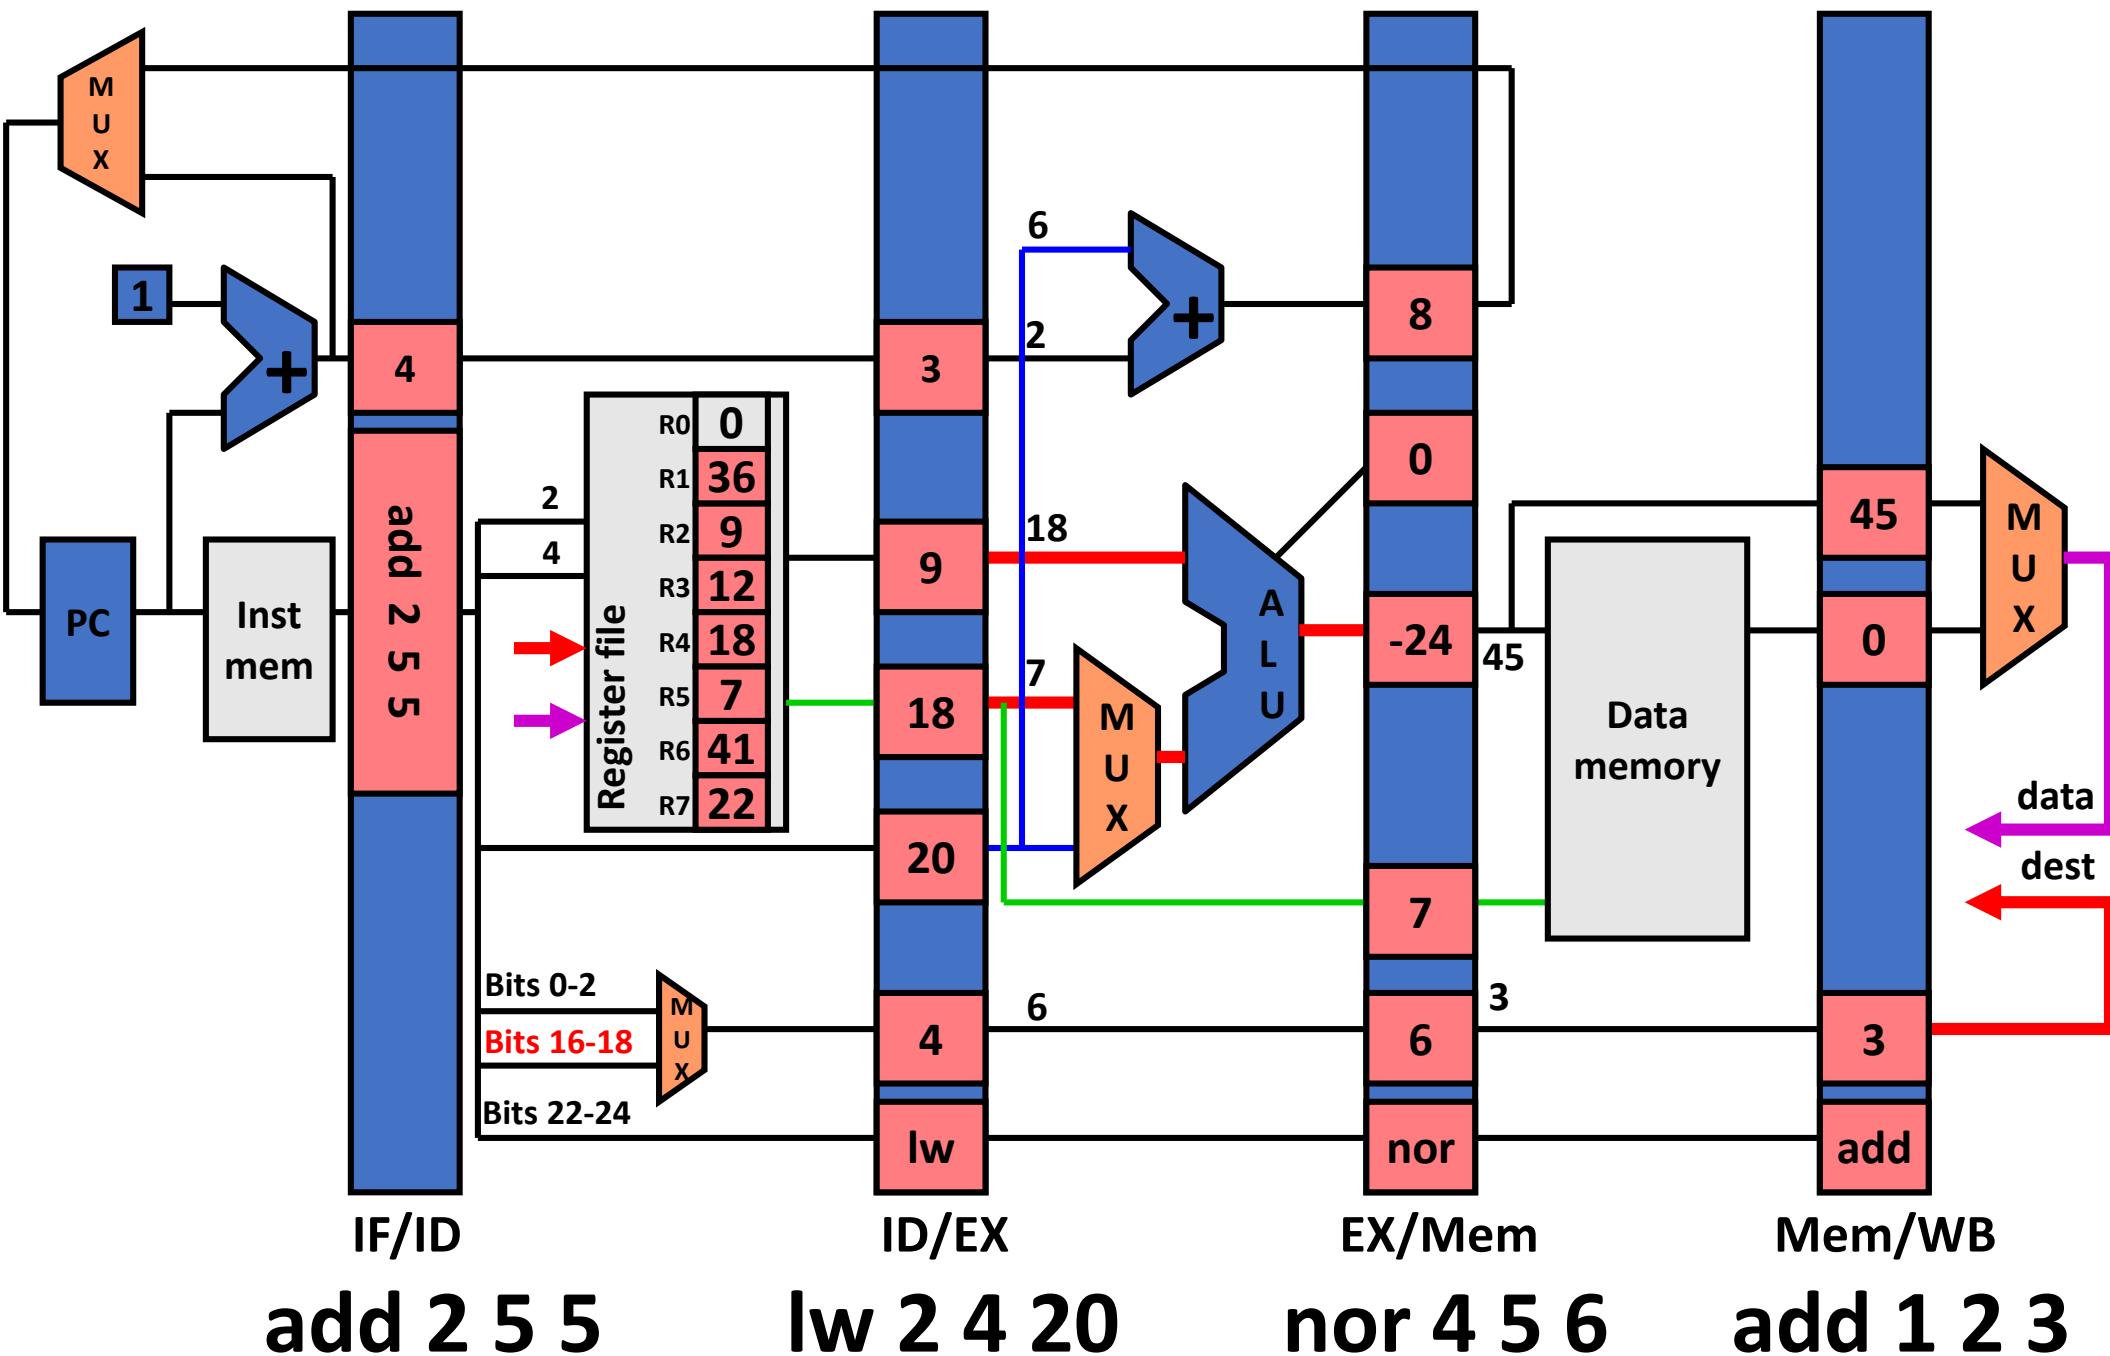
# Time 2 - Fetch: nor 4 5 6



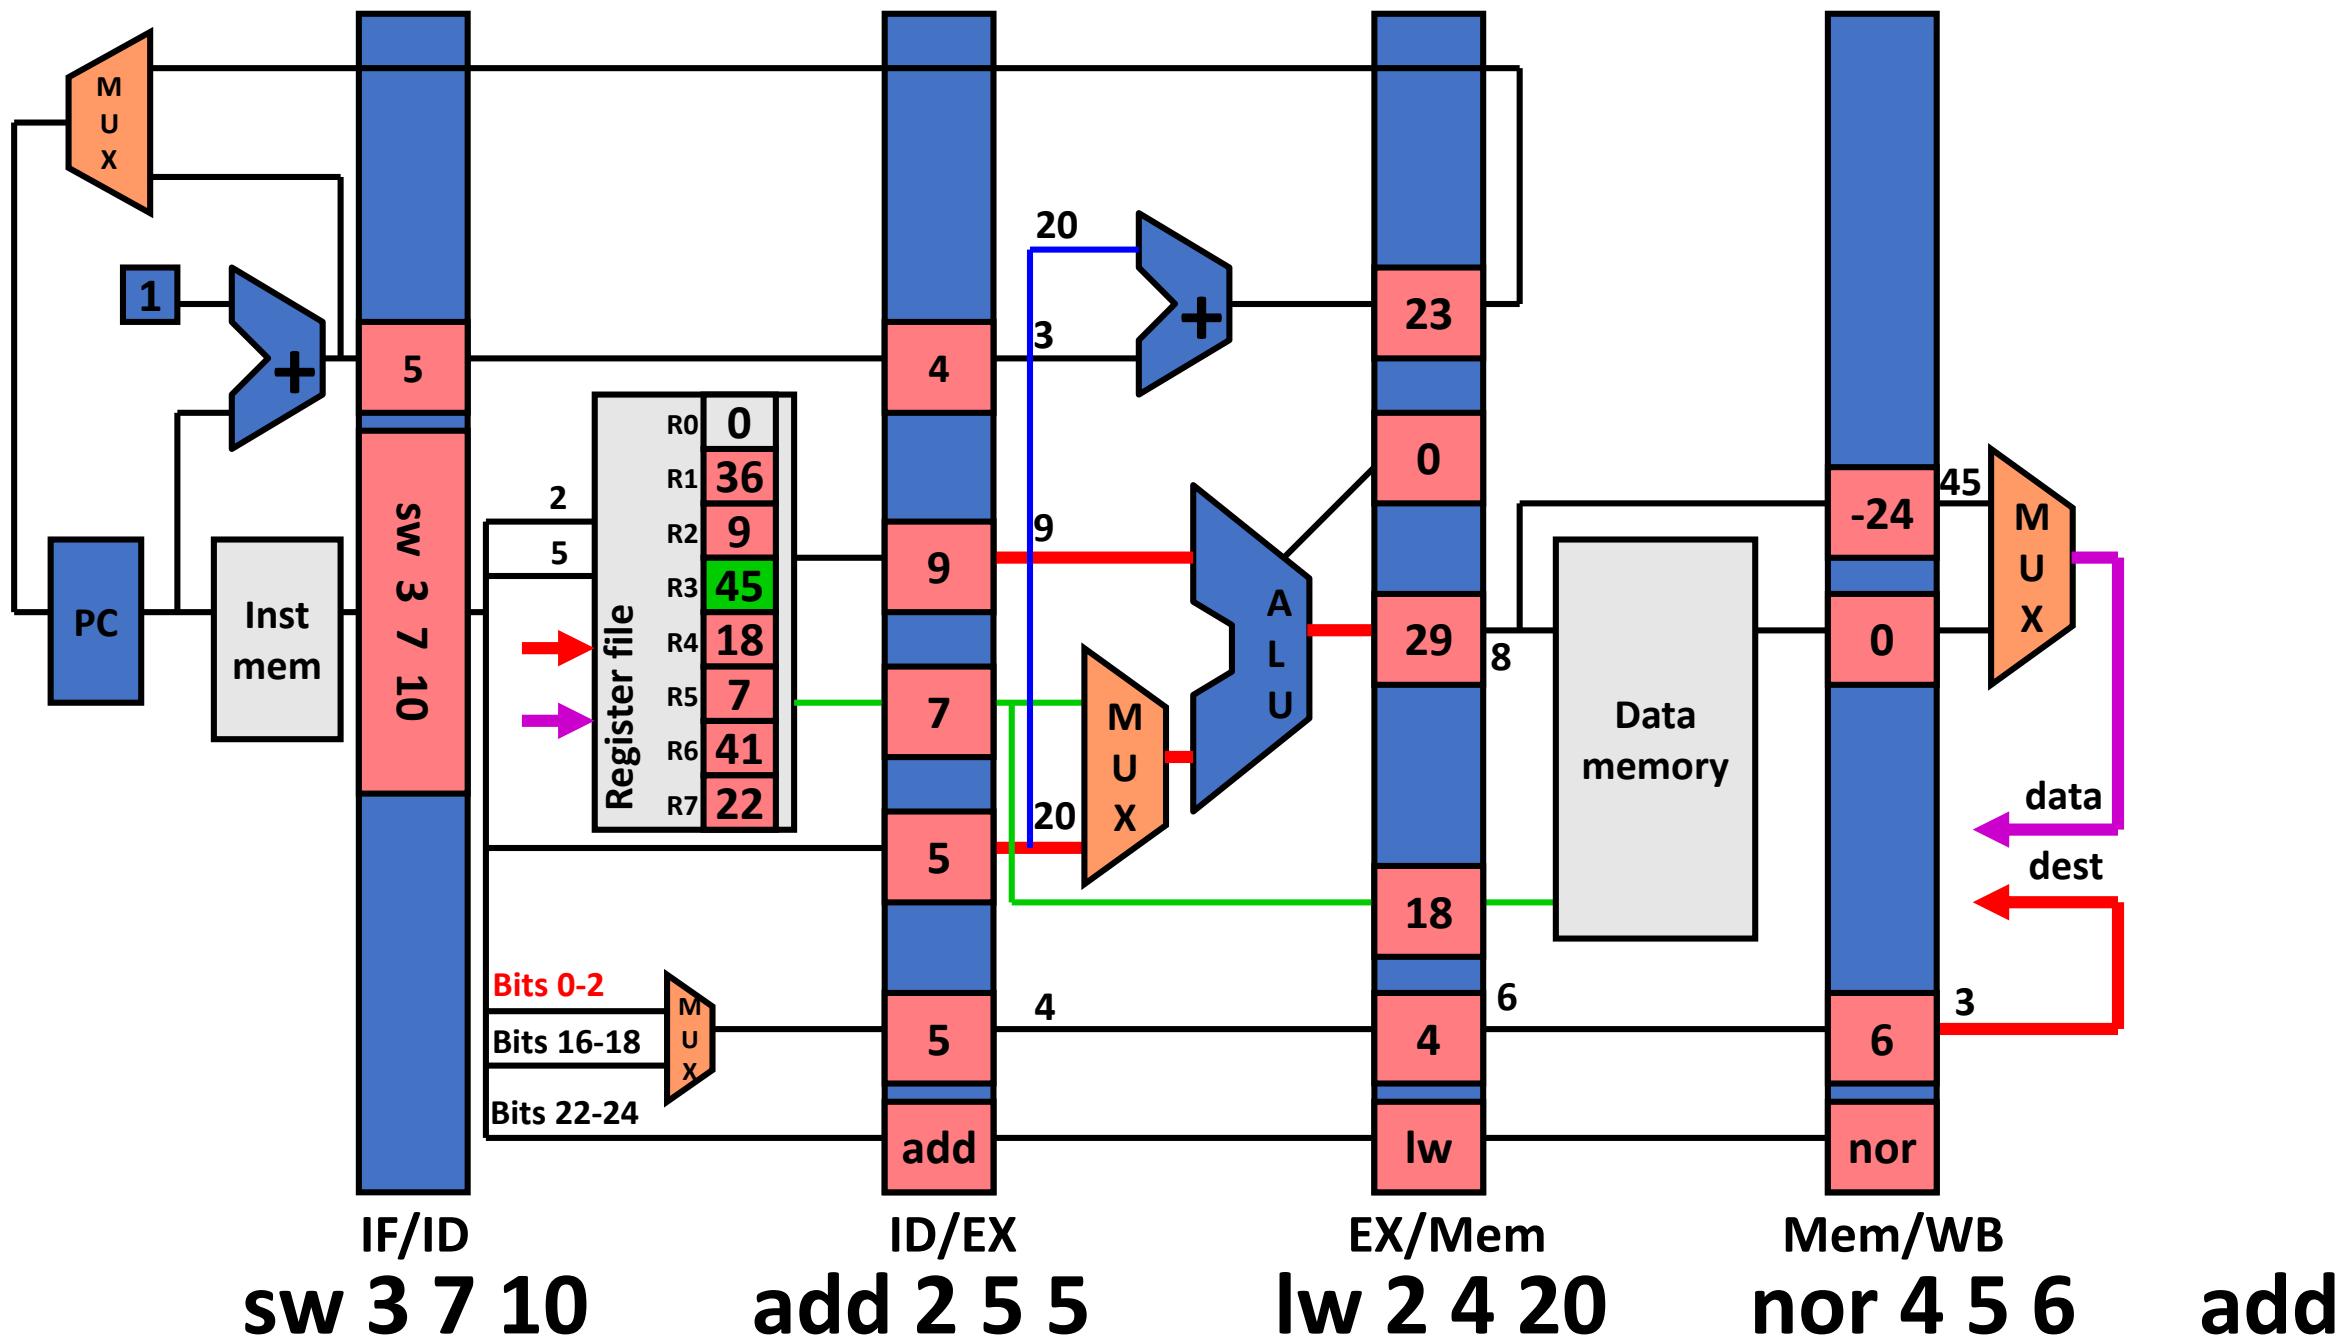
# Time 3 - Fetch: lw 2 4 20



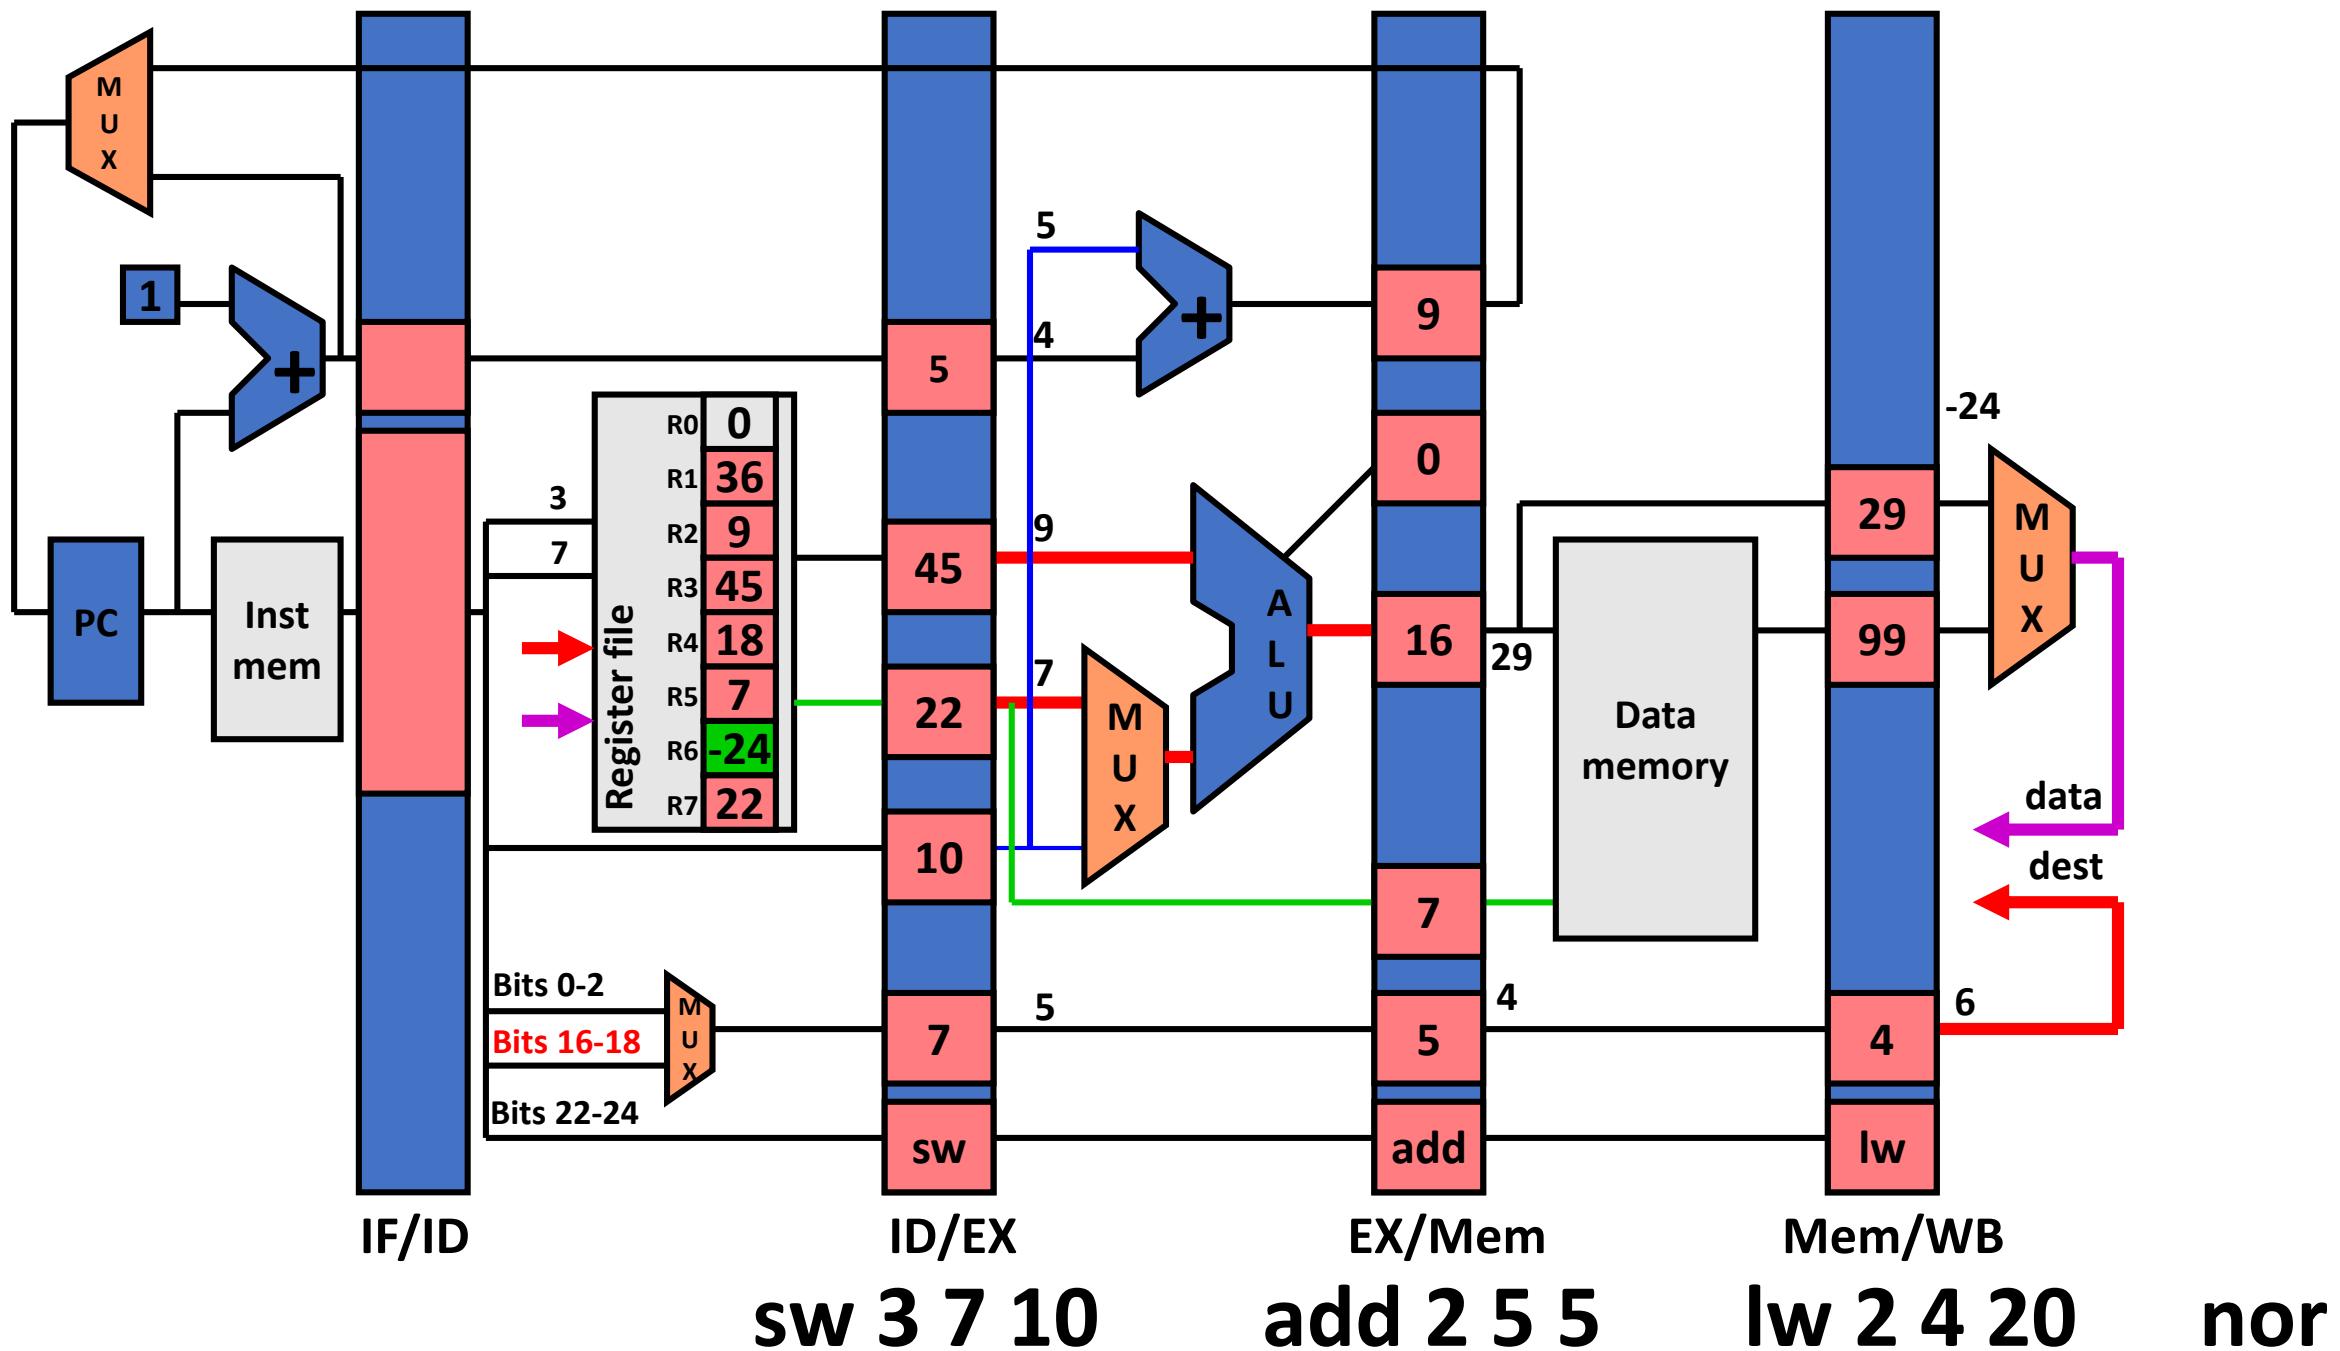
# Time 4 - Fetch: add 2 5 5



# Time 5 - Fetch: sw 3 7 10

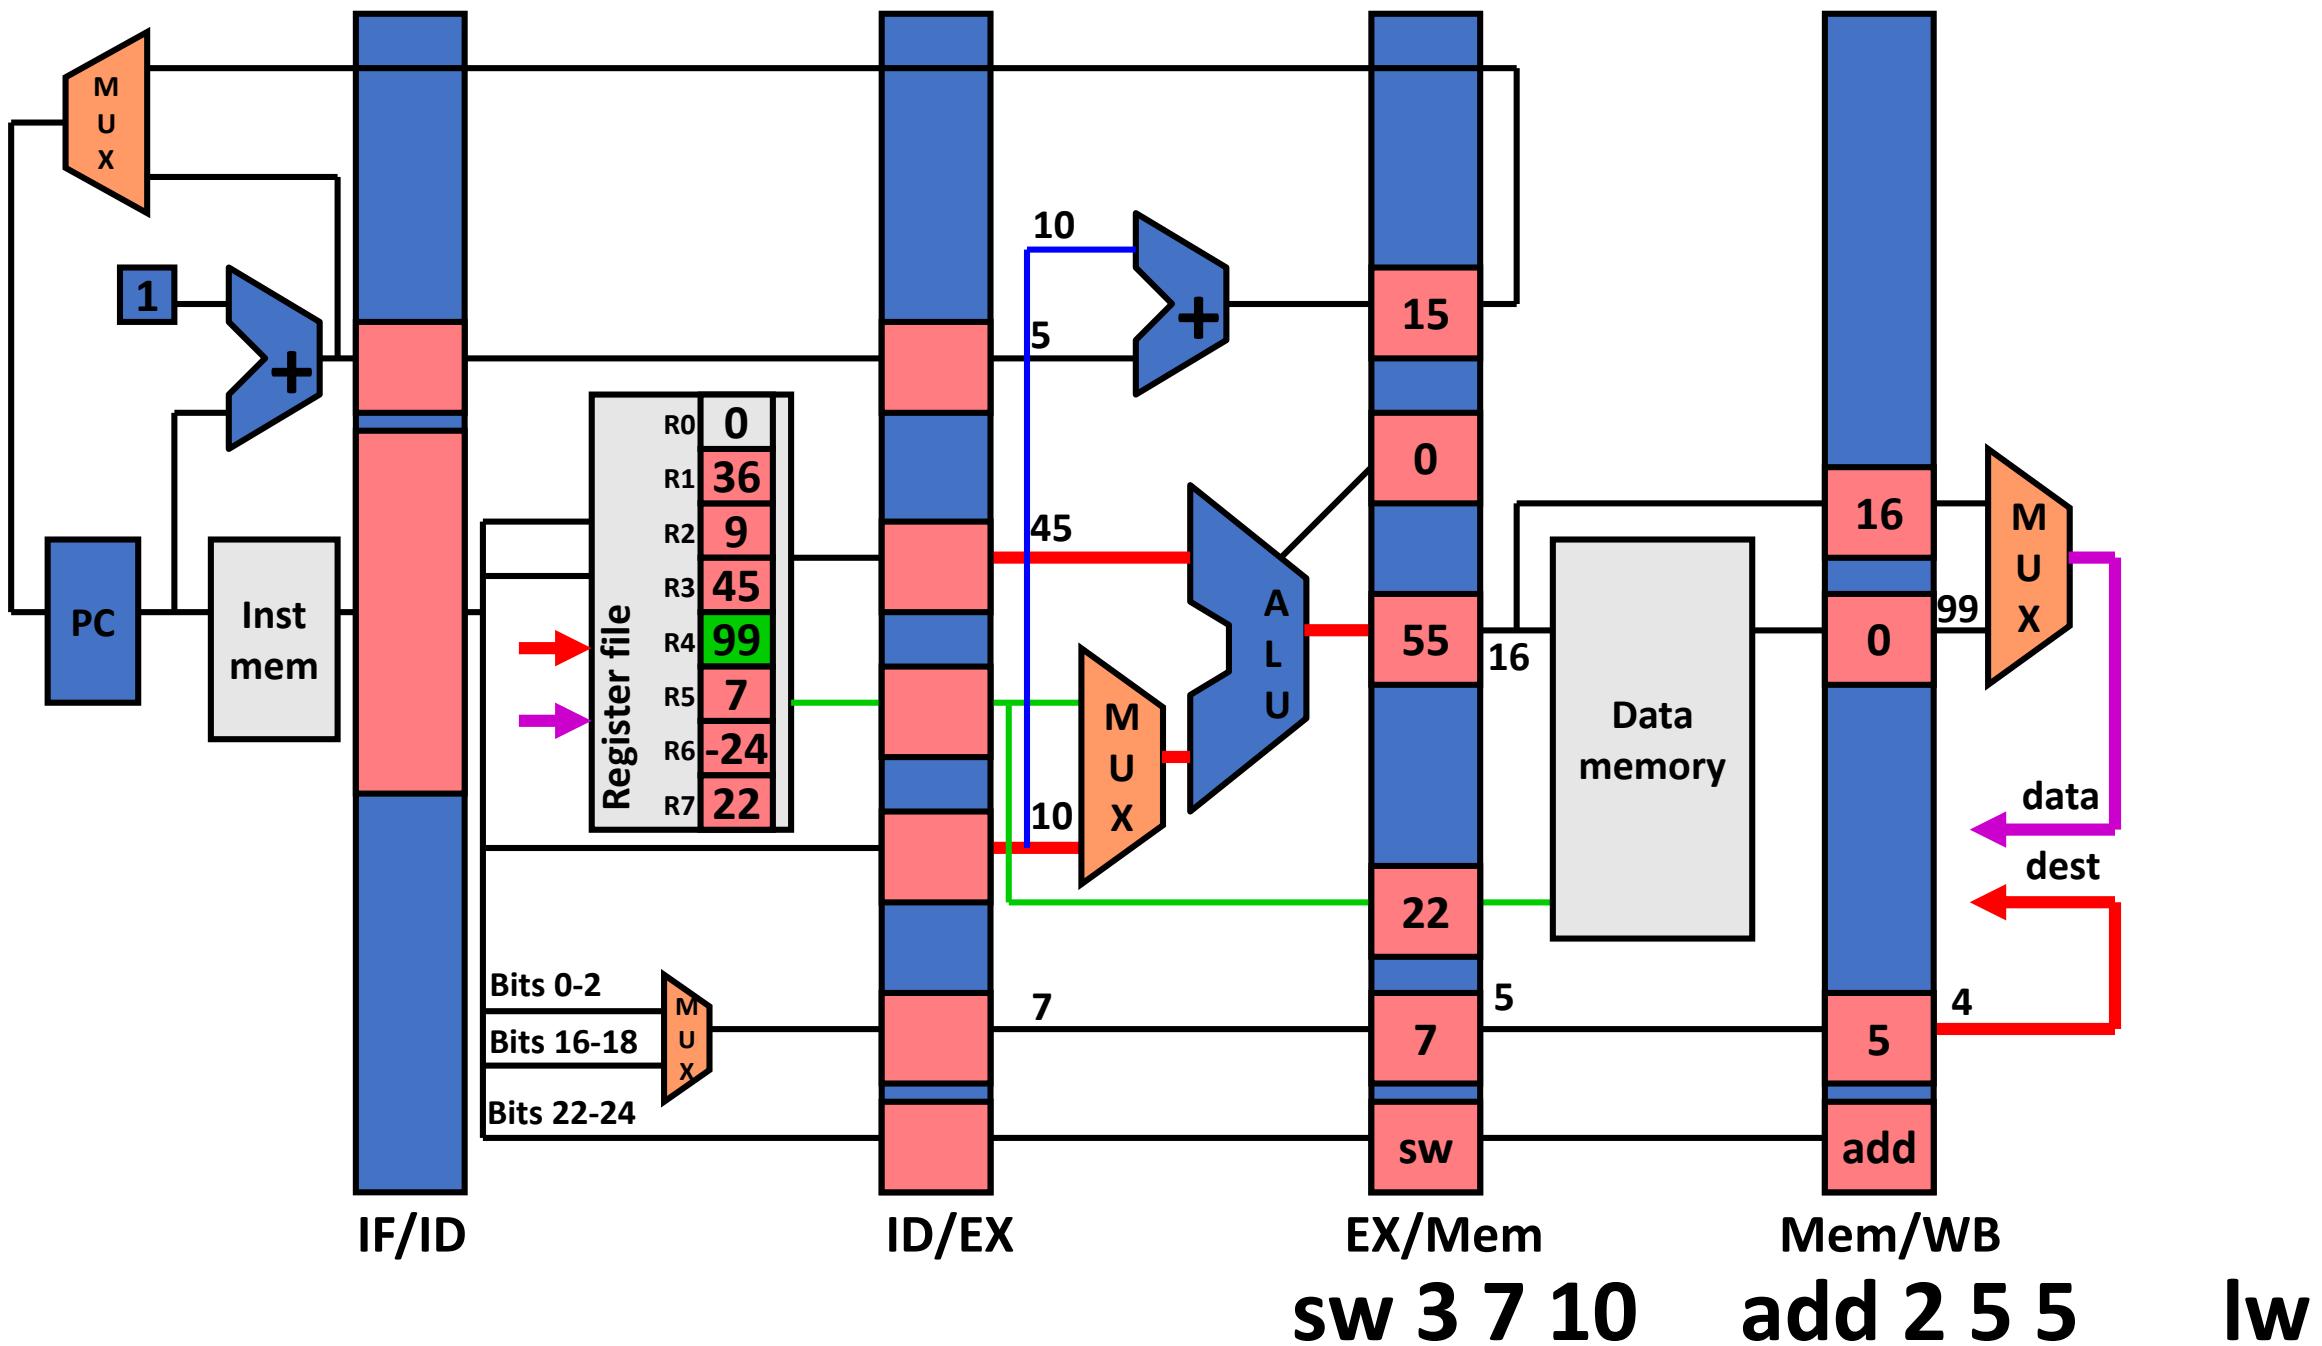


# Time 6 – no more instructions

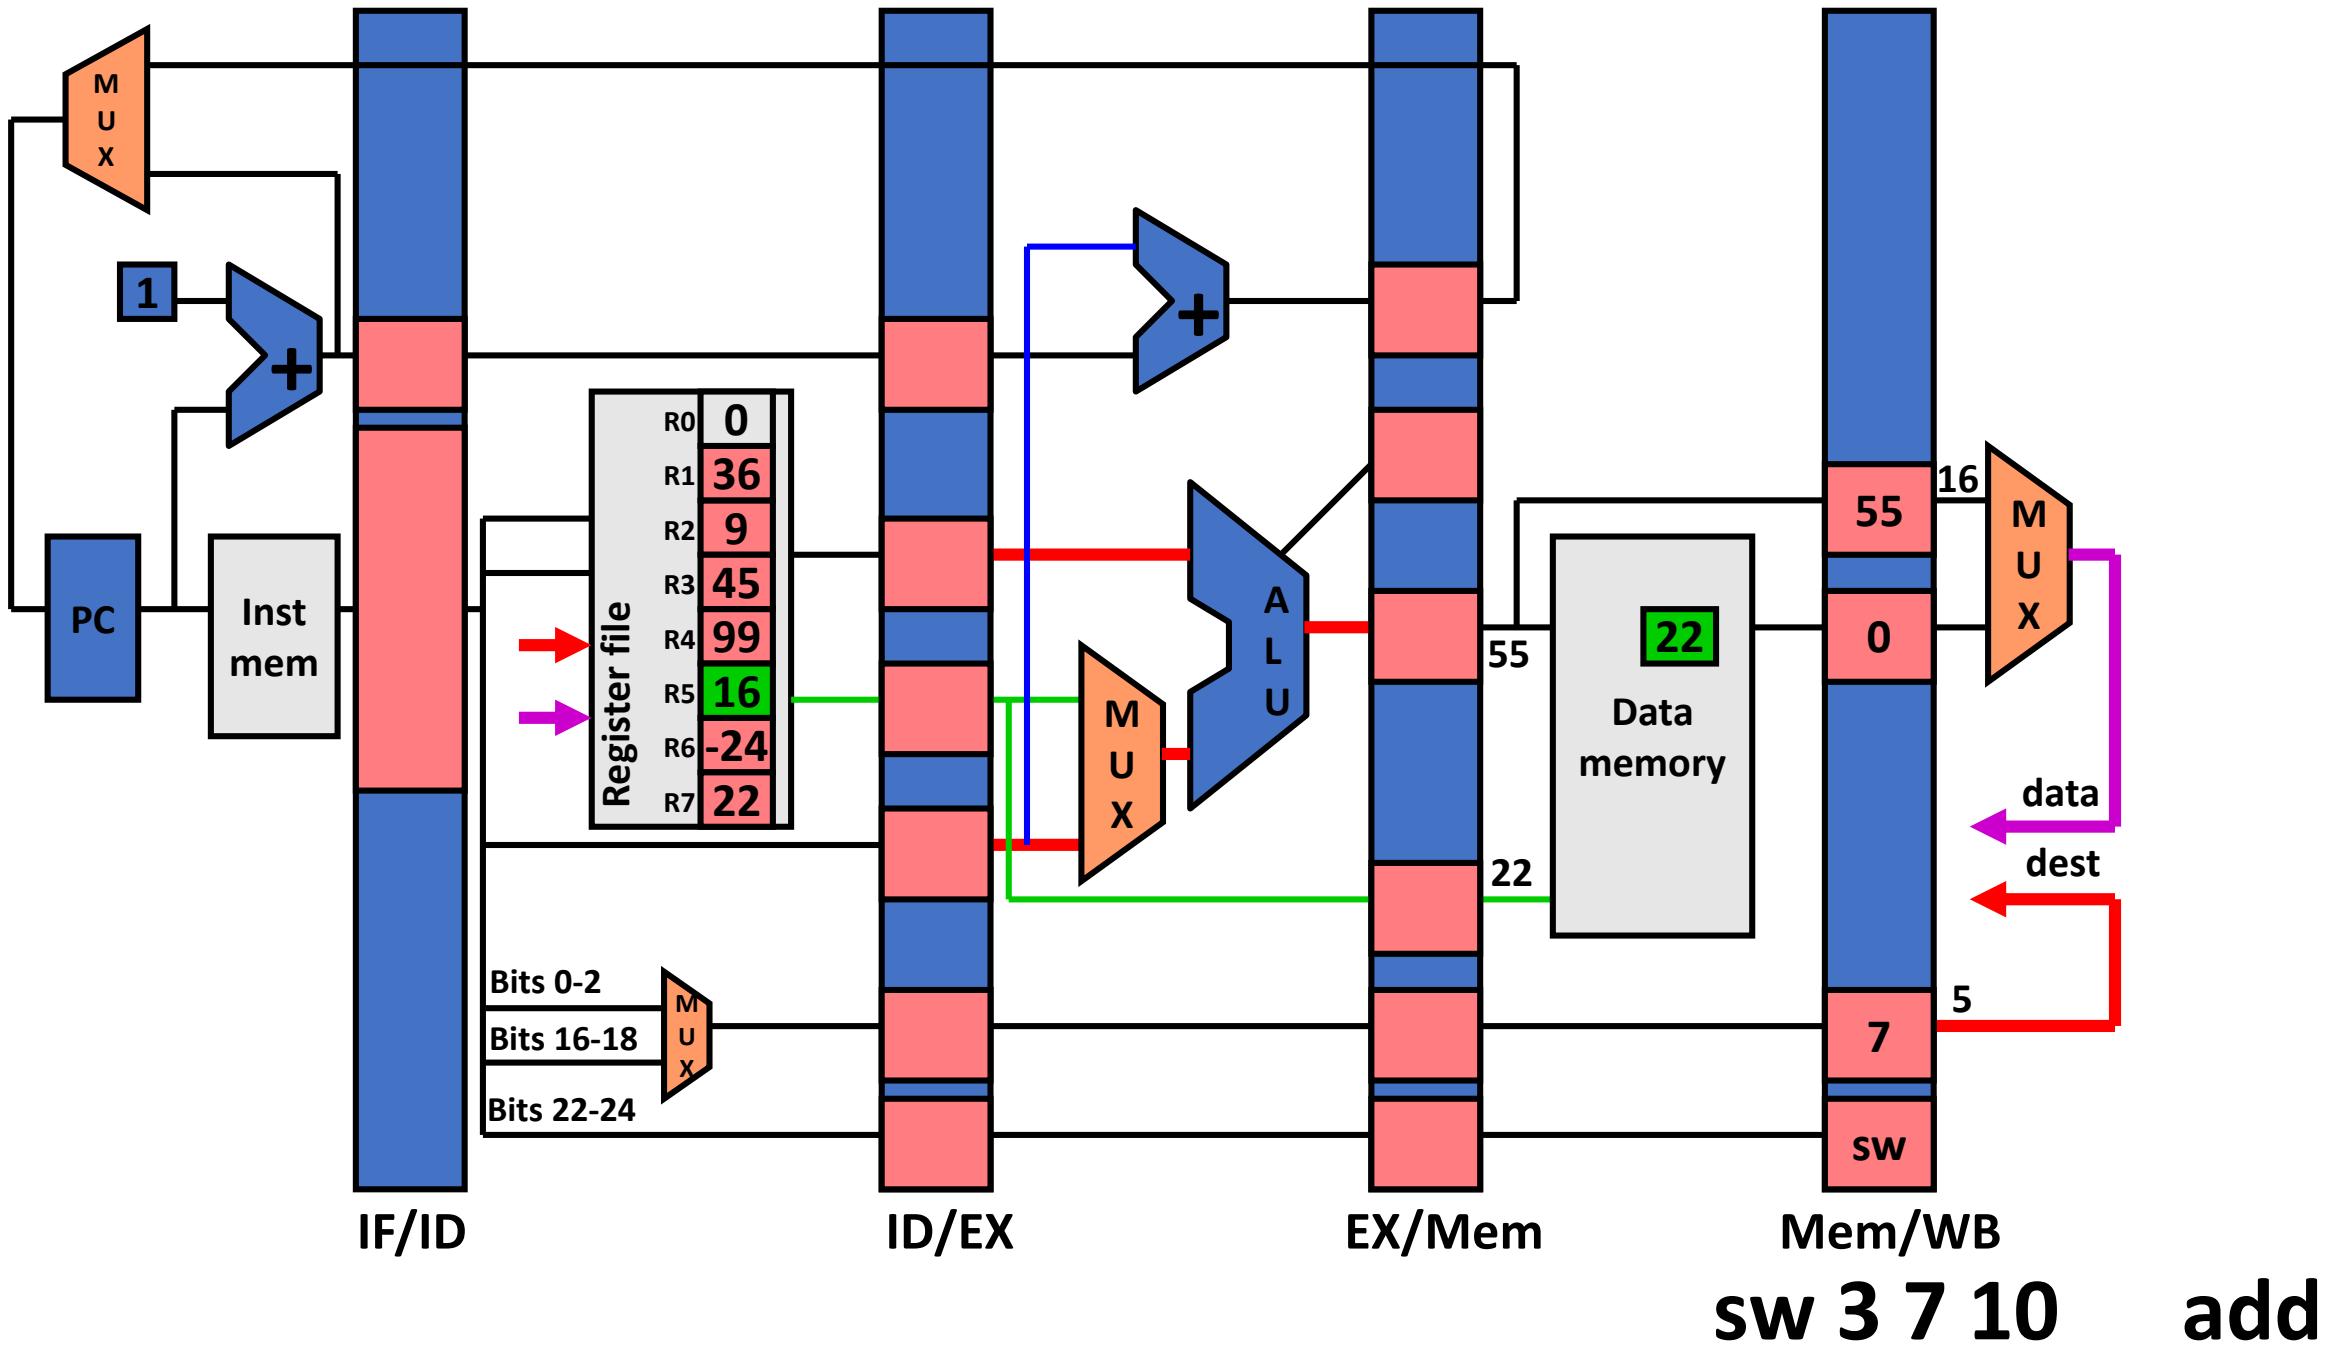


# Time 7 – no more instructions

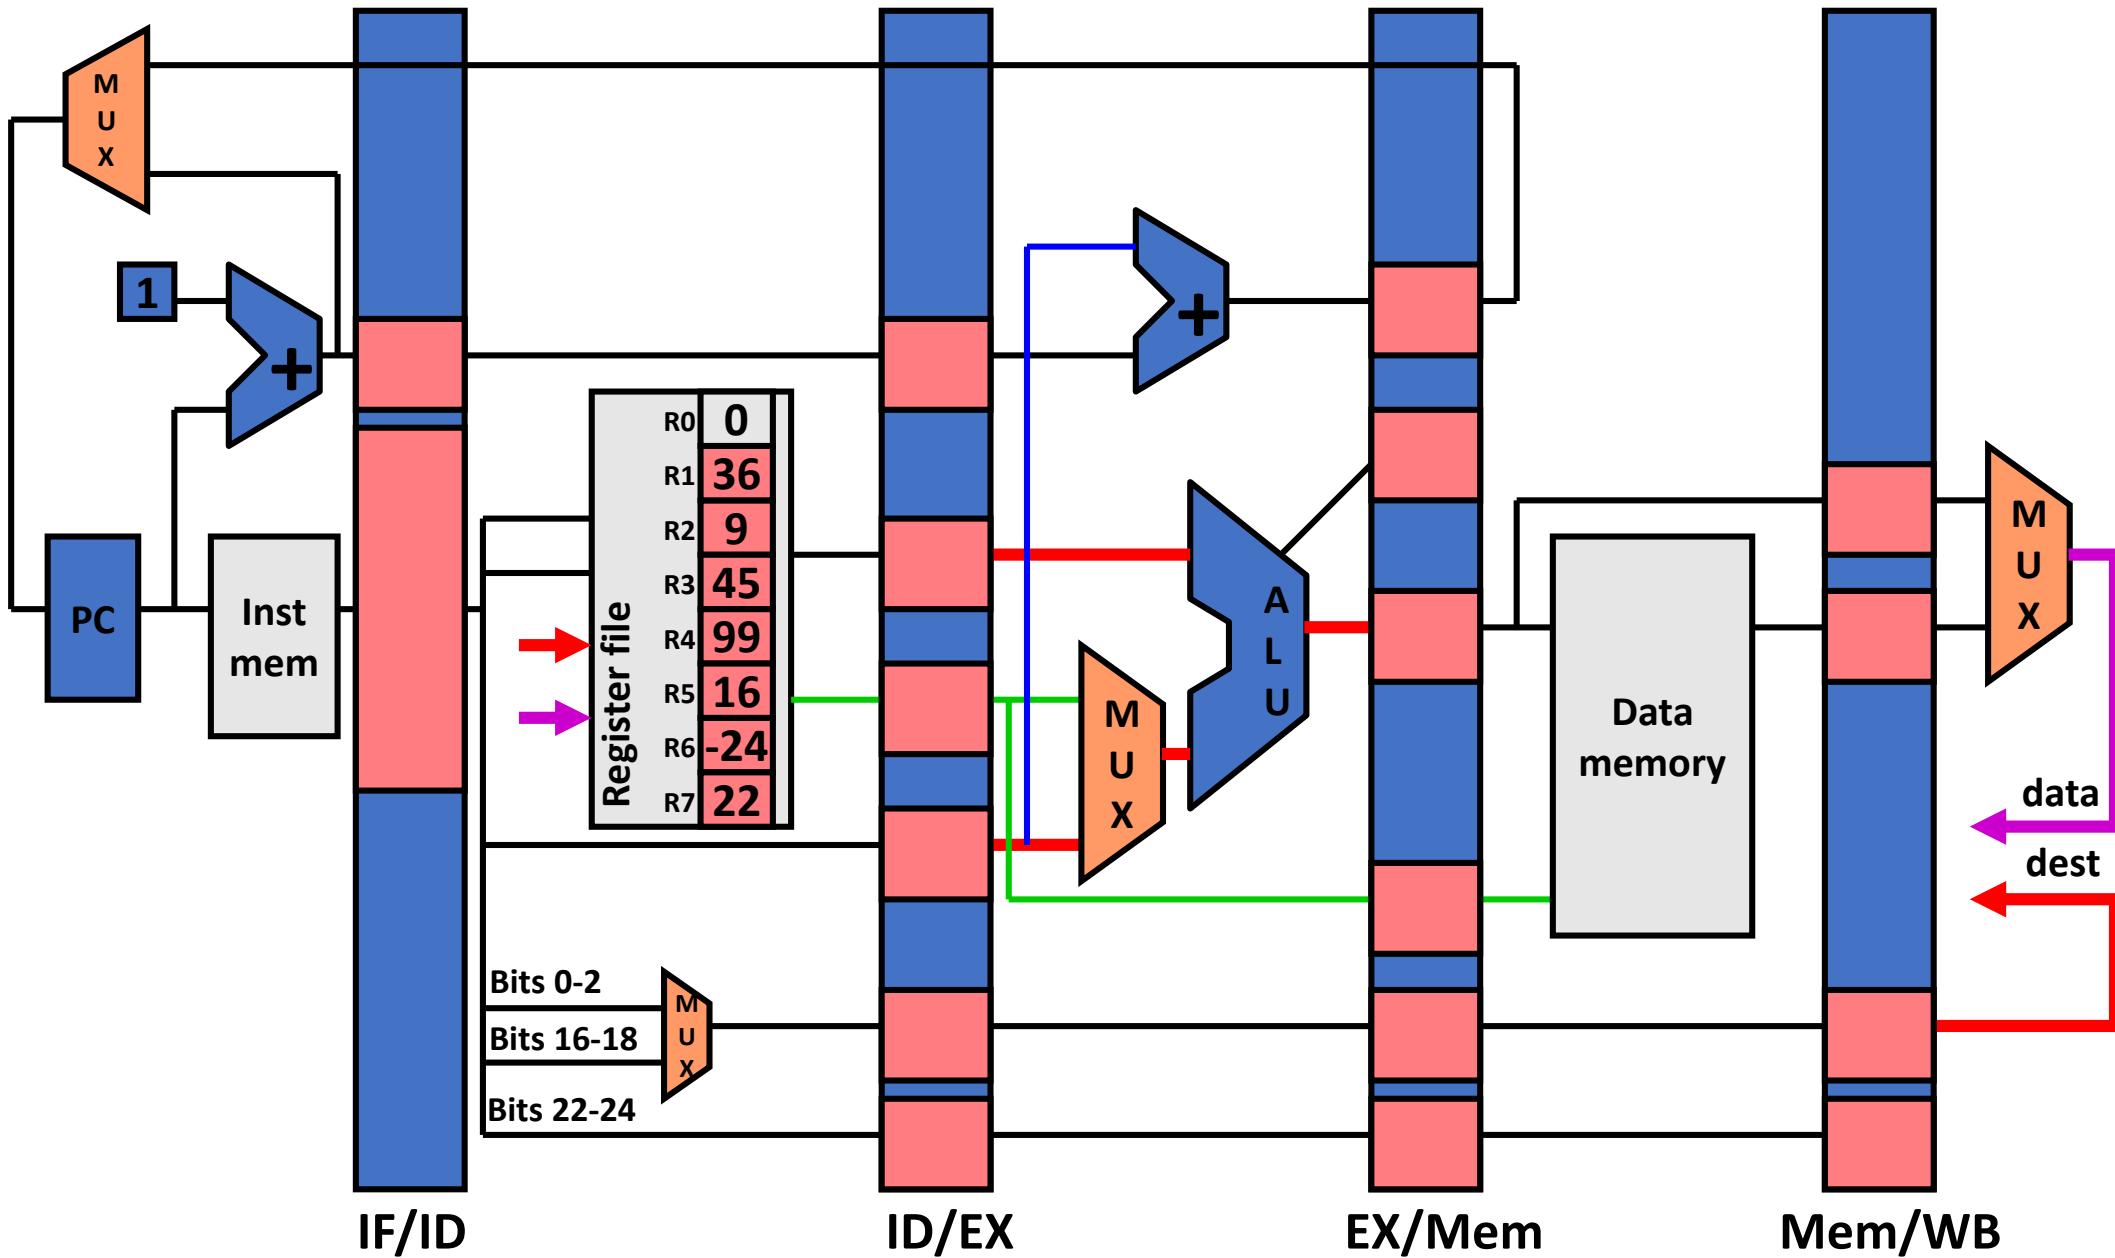
What all happens on the next clock cycle?



# Time 8 – no more instructions



# Time 9 – no more instructions



SW

# Pipelining - What can go wrong?

- **Data hazards**: since register reads occur in stage 2 and register writes occur in stage 5 it is possible to read an old / stale value before the correct value is written back.
- **Control hazards**: A branch instruction may change the PC, but not until stage 4. What do we fetch before that?
- **Exceptions**: Sometimes we need to pause execution, switch to another task (maybe the OS), and then resume execution... how do we make sure we resume at the right spot
- **Now - Data hazards**
  - What are they?
  - How do you detect them?
  - How do you deal with them?



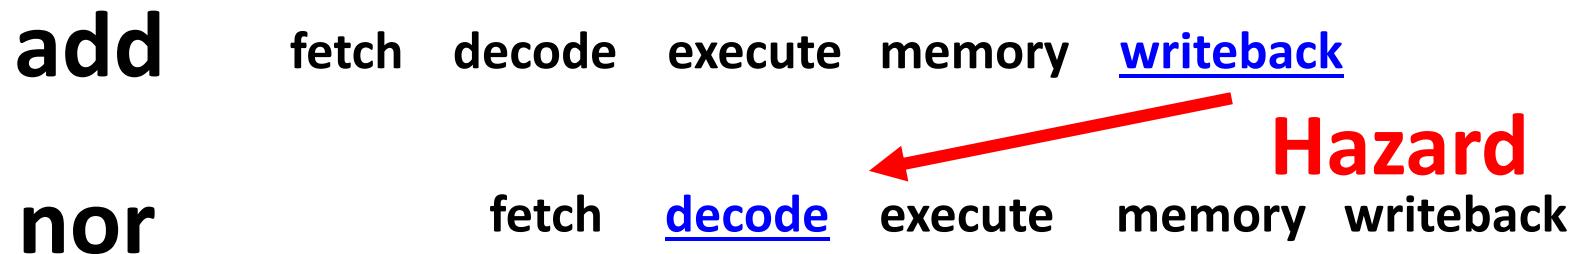
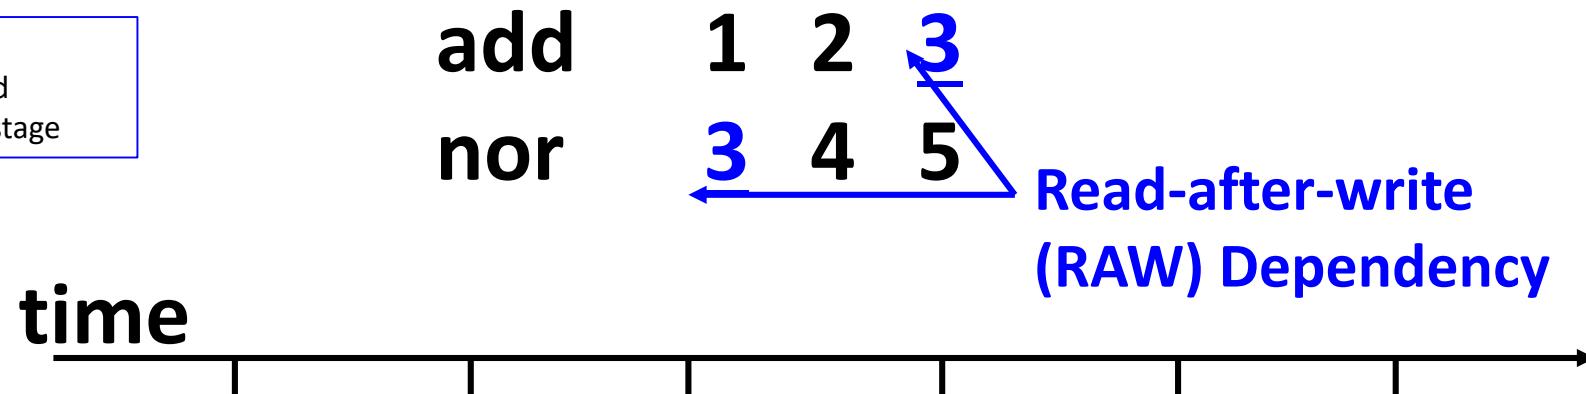
# Pipeline function for ADD

- Fetch: read instruction from memory
- Decode: **read source operands from reg**
- Execute: calculate sum
- Memory: pass results to next stage
- Writeback: **write sum into register file**



# Data Hazards

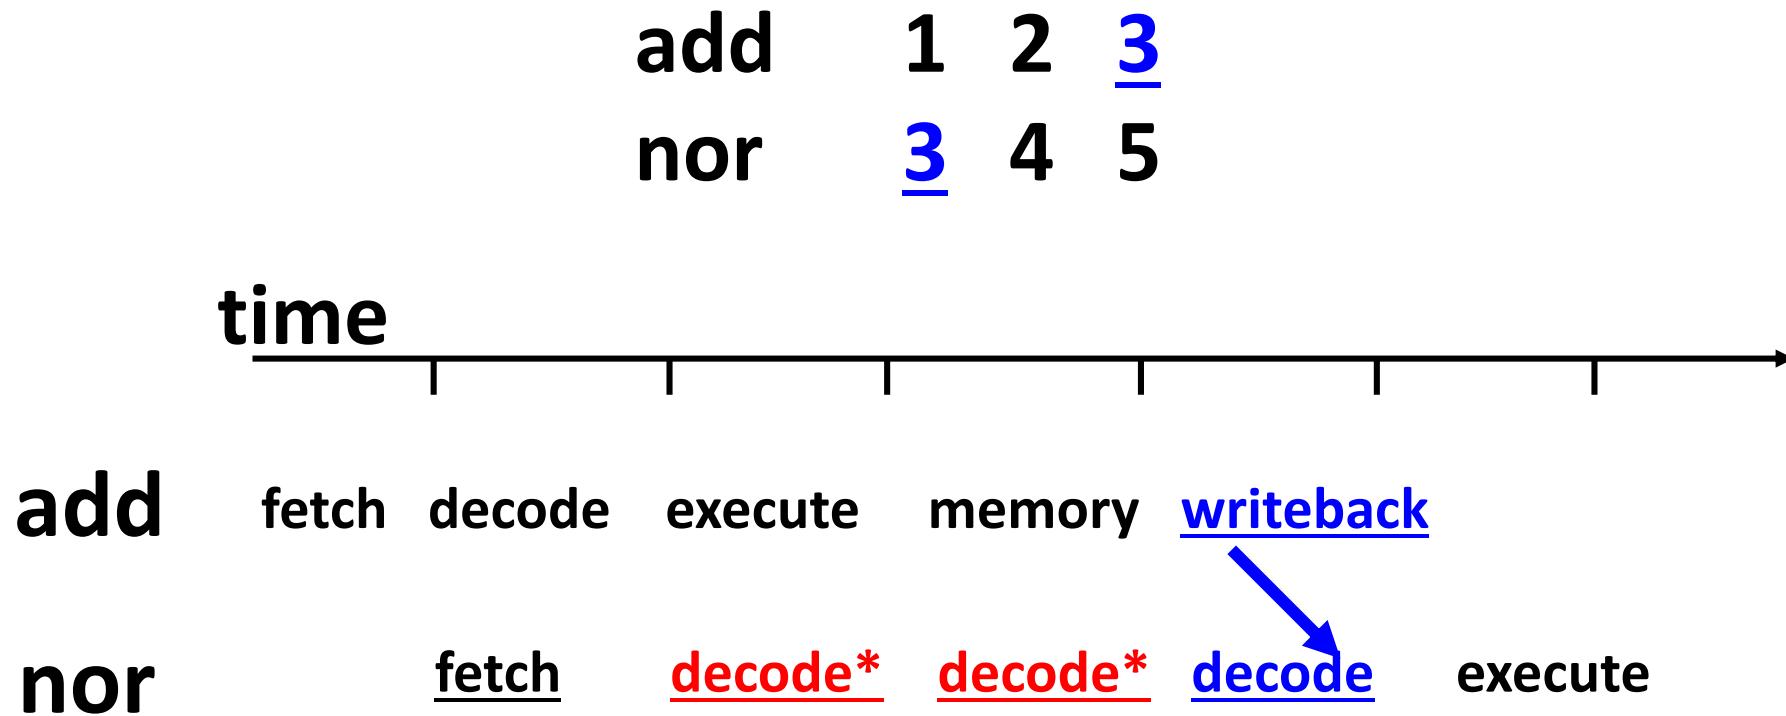
Recall: registers  
are read /sourced  
In the “decode” stage



If not careful, nor will read a stale value of register 3



# Data Hazards



Assume Register File gives the right value of **register 3** when read/written during same cycle. This is consistent with most processors (ARM/x86), but not Project 3.



# Definitions

- Data Dependency: *one instruction uses the result of a previous one*
  - Doesn't necessarily cause a problem
- Data Hazard: *one instruction has a data dependency that will cause a problem if we don't "deal with it"*



# Class Problem 1

Which of these instructions has a data dependency on an earlier one? Which of those are data hazards in our 5-stage pipeline?

1. add 1 2 3
2. nor 3 4 5
3. add 6 3 7
4. lw 3 6 10
5. sw 6 2 12

What about here?

1. add 1 2 3
2. beq 3 4 1
3. add 3 5 6
4. add 3 6 7



# Class Problem 1

Which read-after-write (RAW) dependences do you see?

Which of those are data hazards?

1. add 1 2 3

2. nor 3 4 5

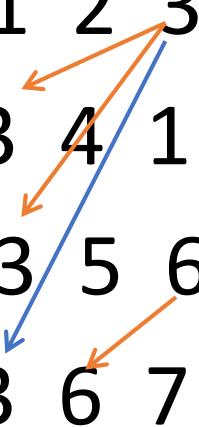
3. add 6 3 7

4. lw 3 6 10

5. sw 6 2 12

What about here?

1. add 1 2 3
2. beq 3 4 1
3. add 3 5 6
4. add 3 6 7

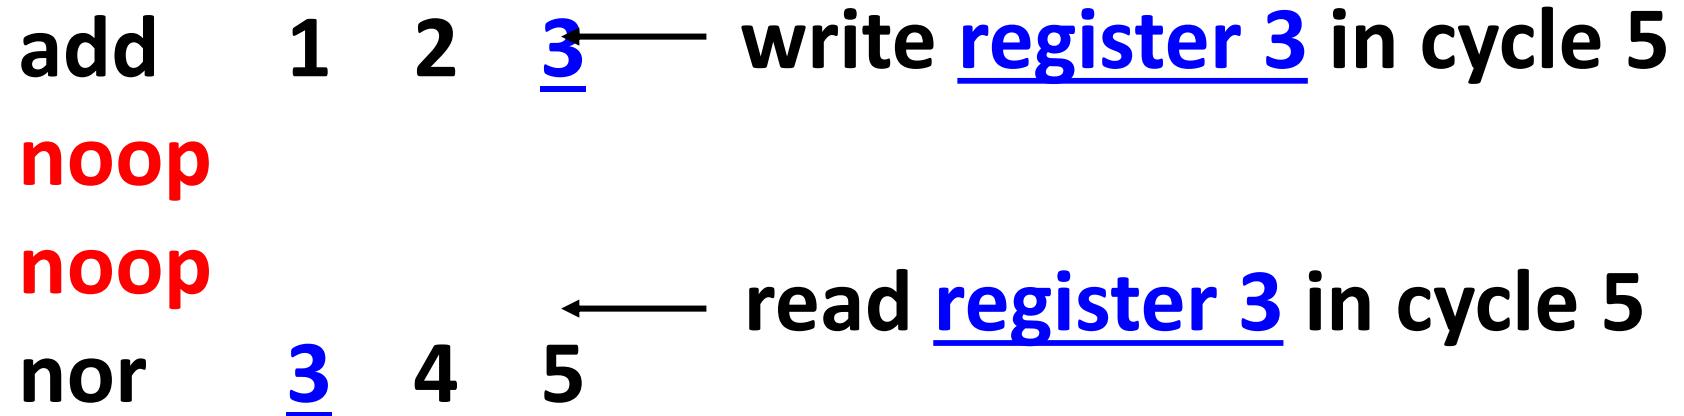


# Three approaches to handling data hazards

- Avoid
  - Make sure there are no hazards in the code
- Detect and Stall
  - If hazards exist, stall the processor until they go away.
- Detect and Forward
  - If hazards exist, fix up the pipeline to get the correct value (if possible)

# Handling data hazards I: Avoid all hazards

- Assume the programmer (or the compiler) knows about the processor implementation.
  - Make sure no hazards exist.
    - Put noops between any dependent instructions.



# Problems with this solution

- Old programs (legacy code) may not run correctly on new implementations
  - Longer pipelines need more noops
- Programs get larger as noops are included
  - Especially a problem for machines that try to execute more than one instruction every cycle
  - Intel EPIC: Often 25% - 40% of instructions are noops
- Program execution is slower
  - CPI is 1, but some instructions are noops

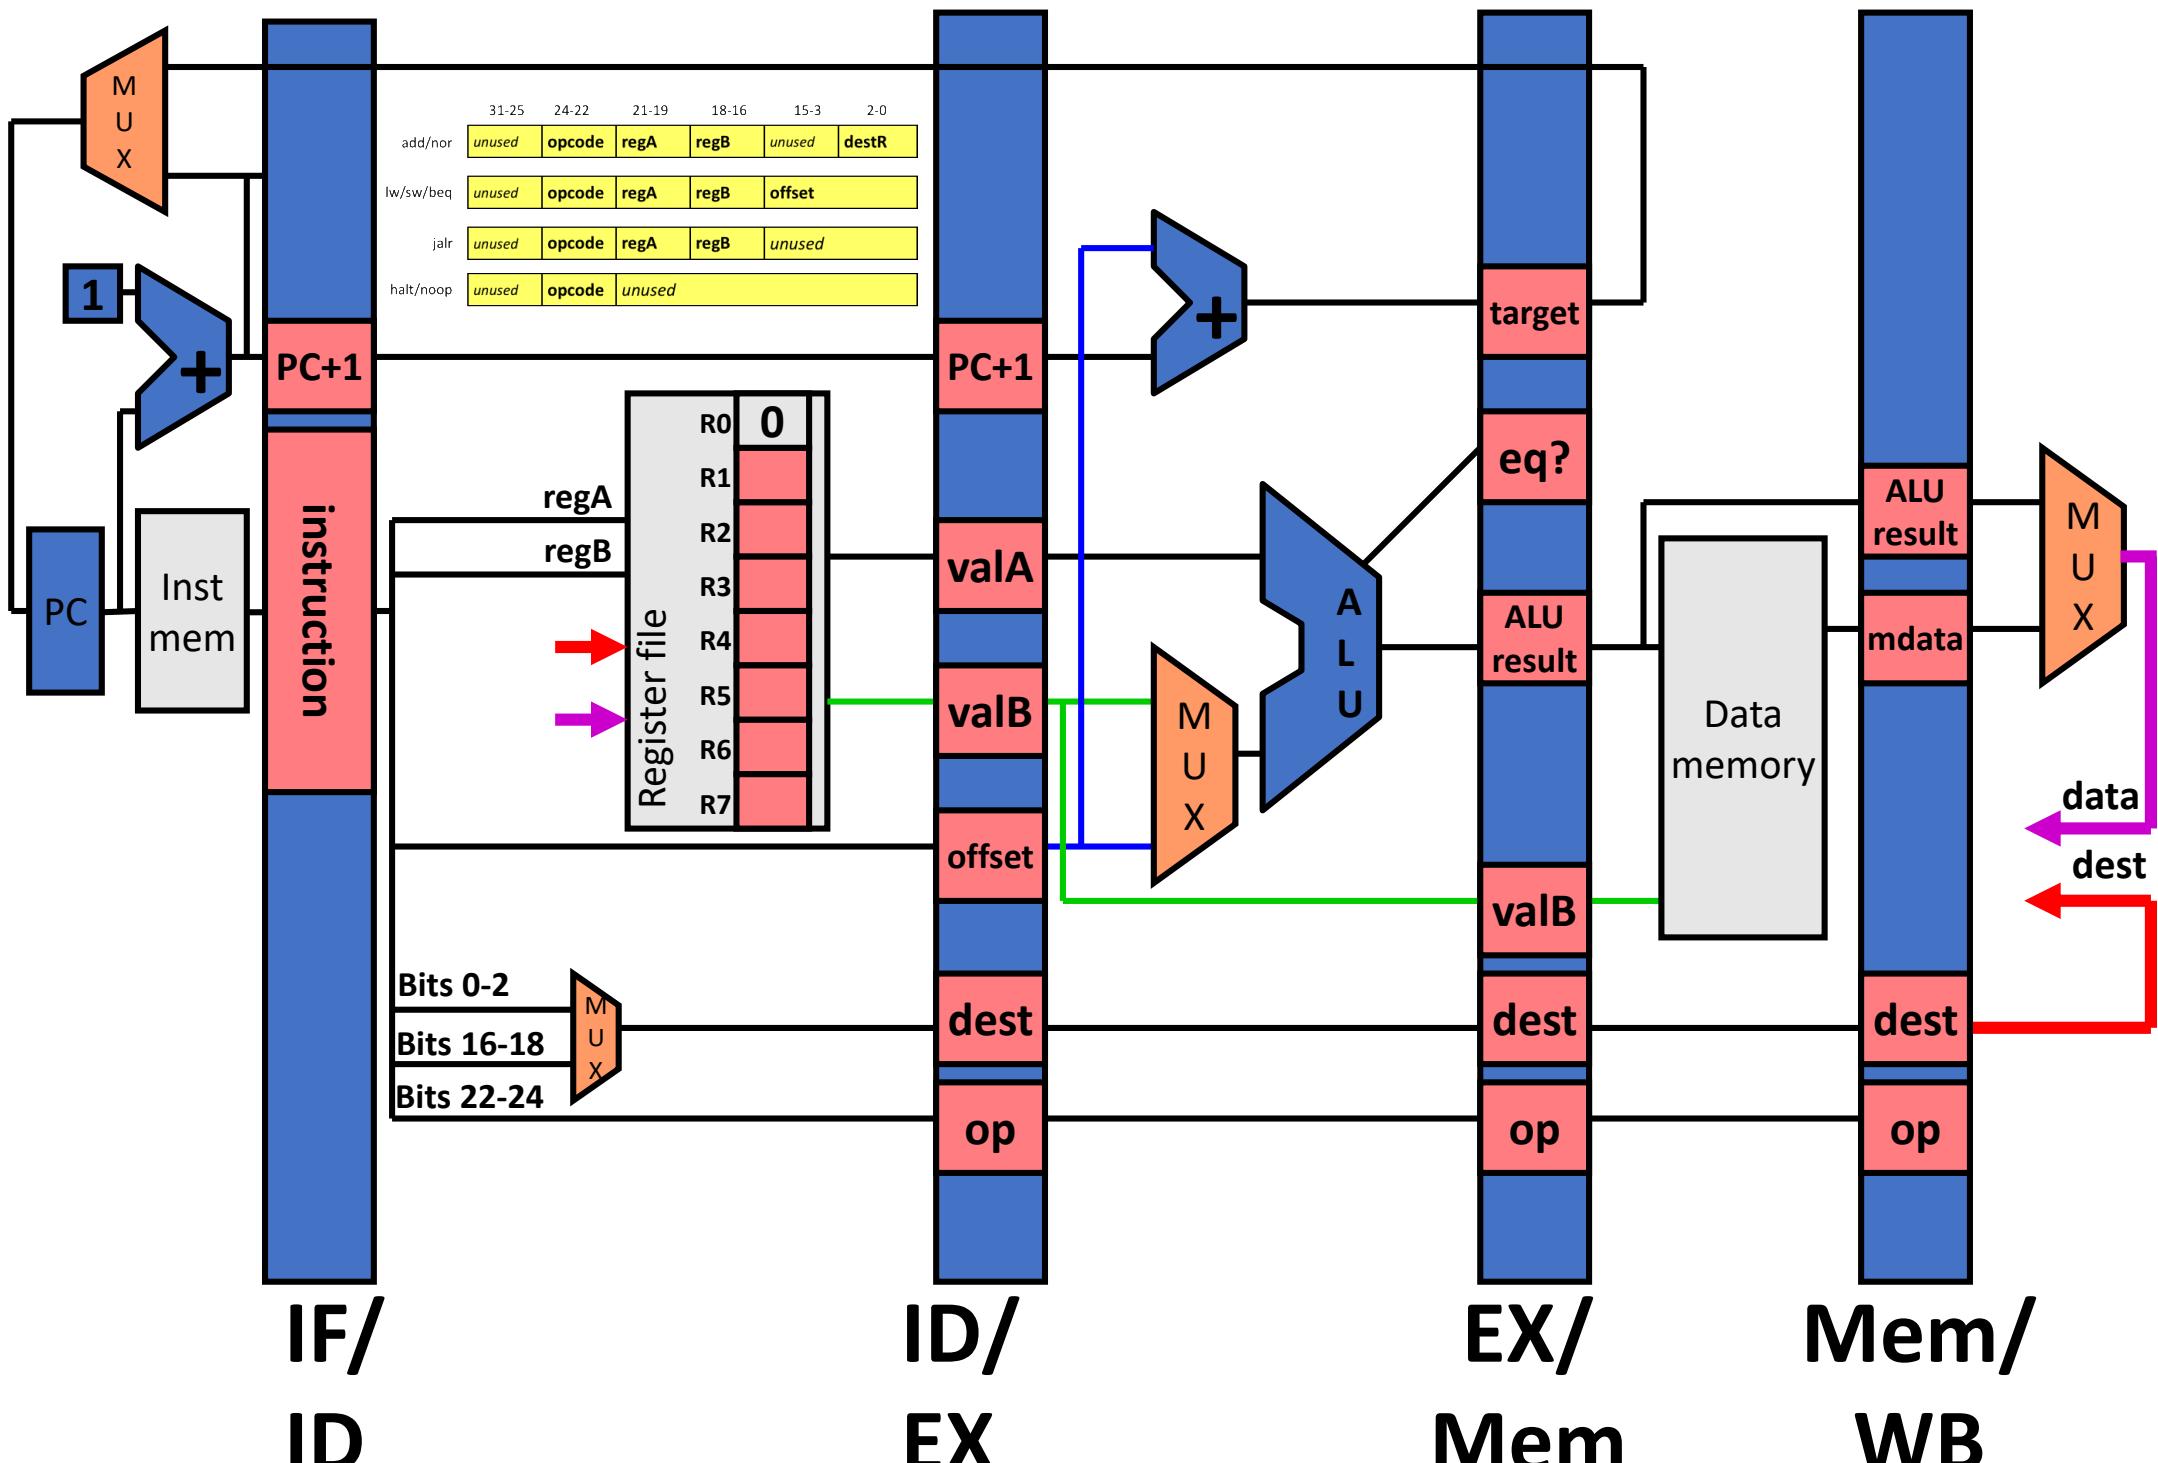


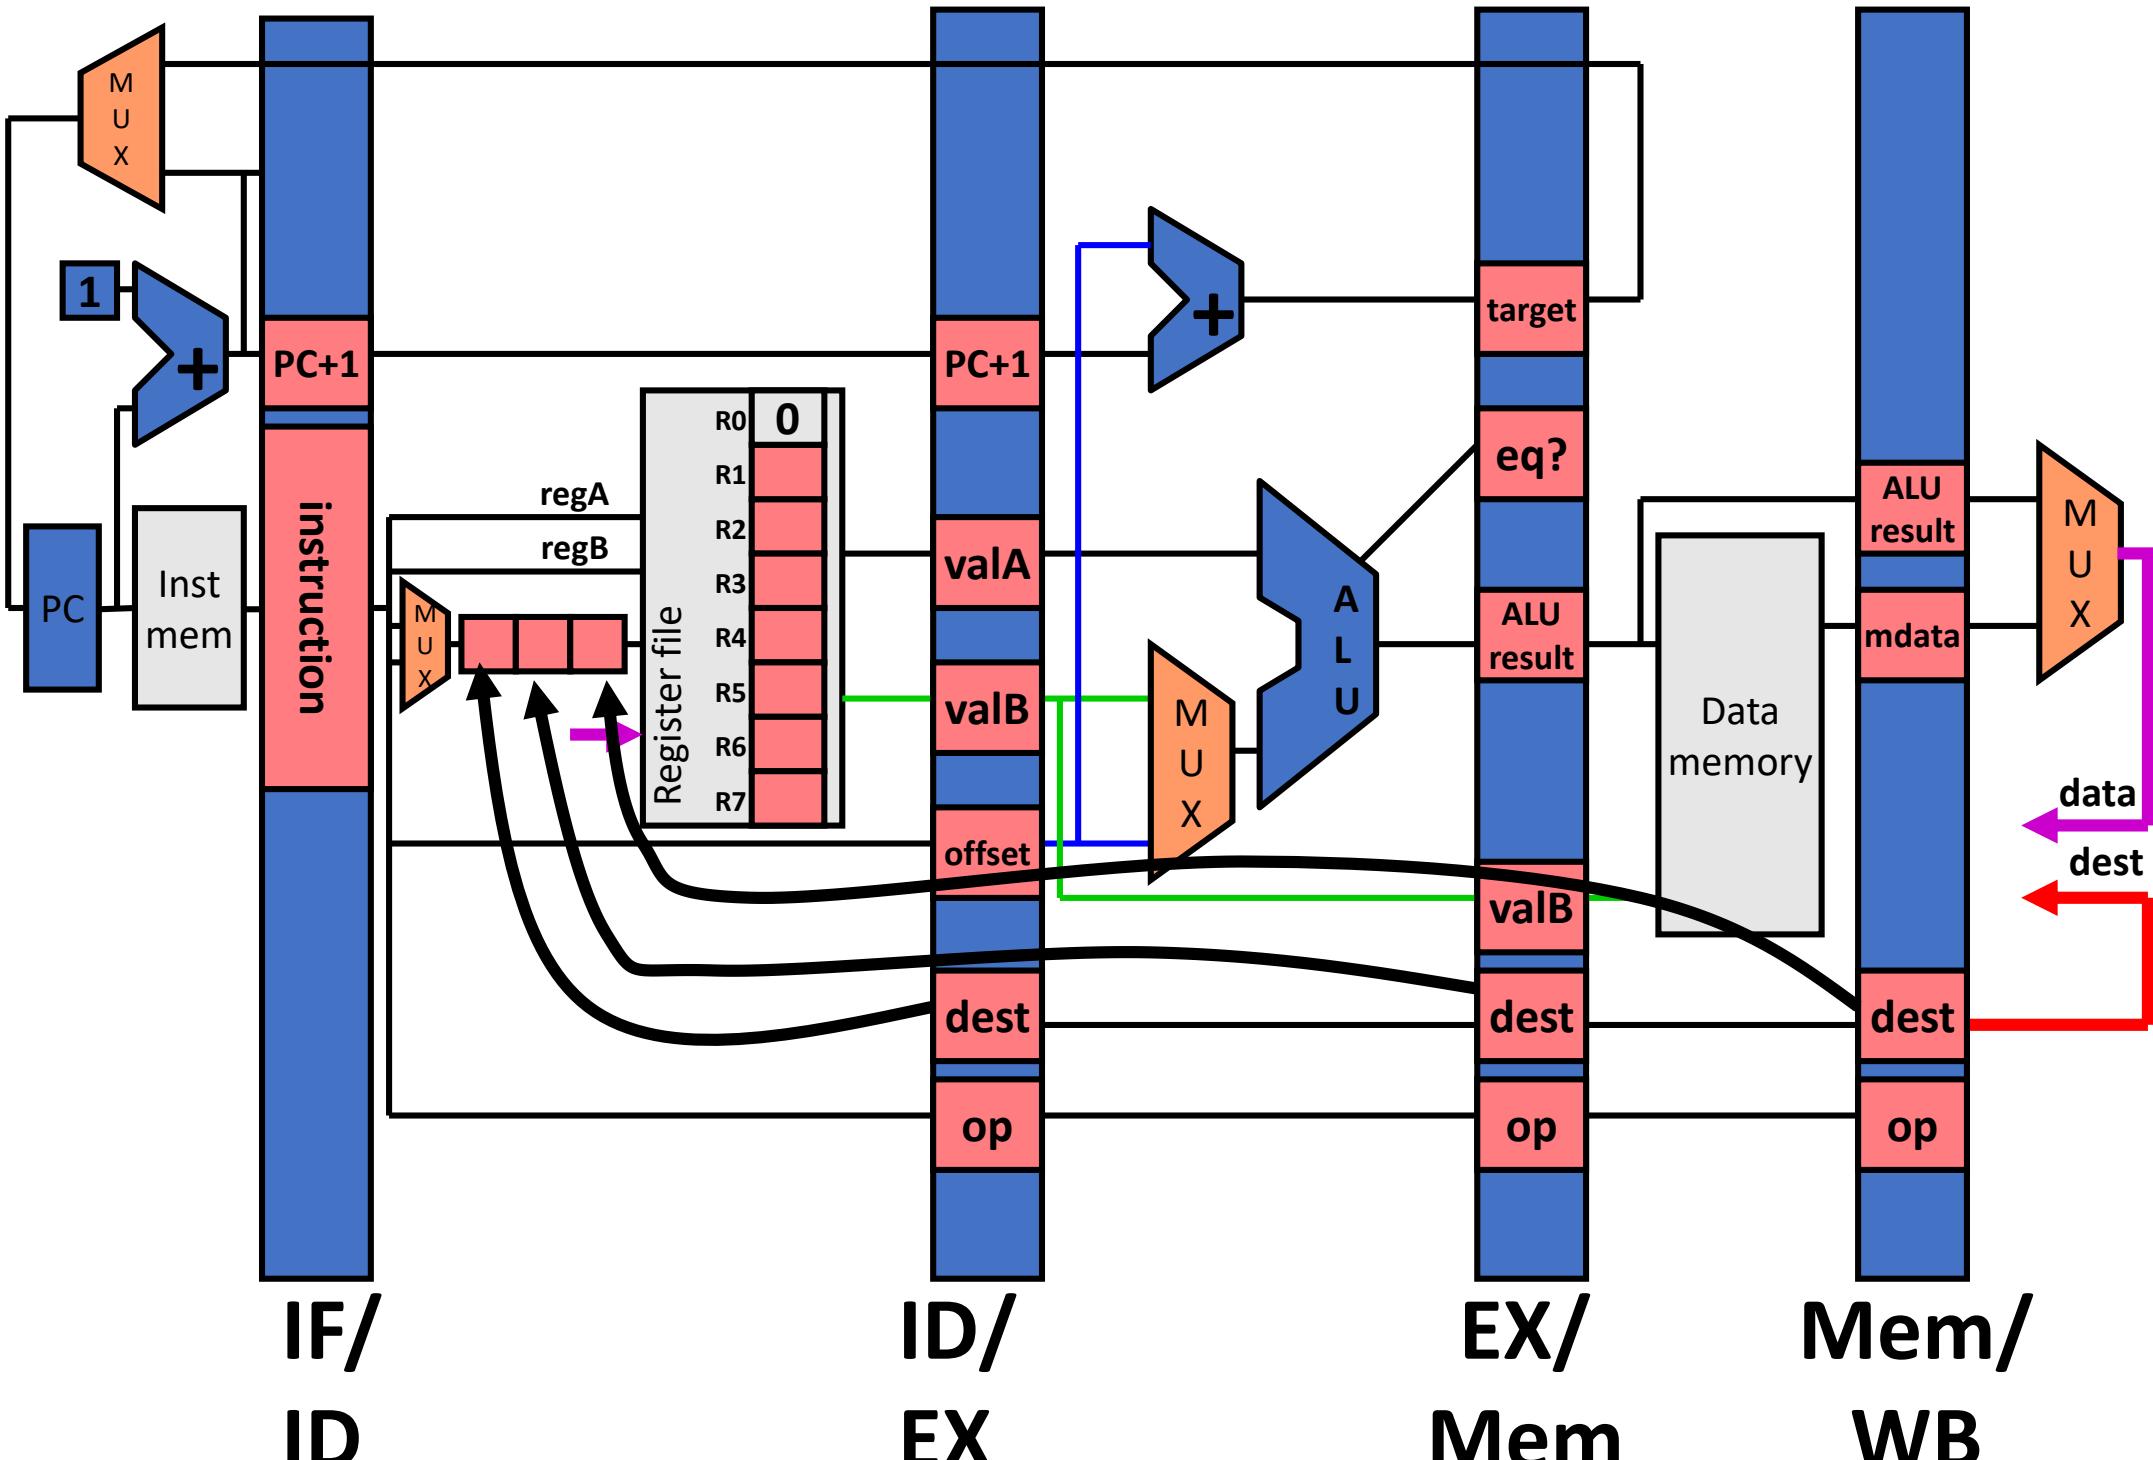
# Handling data hazards II: Detect and stall until ready

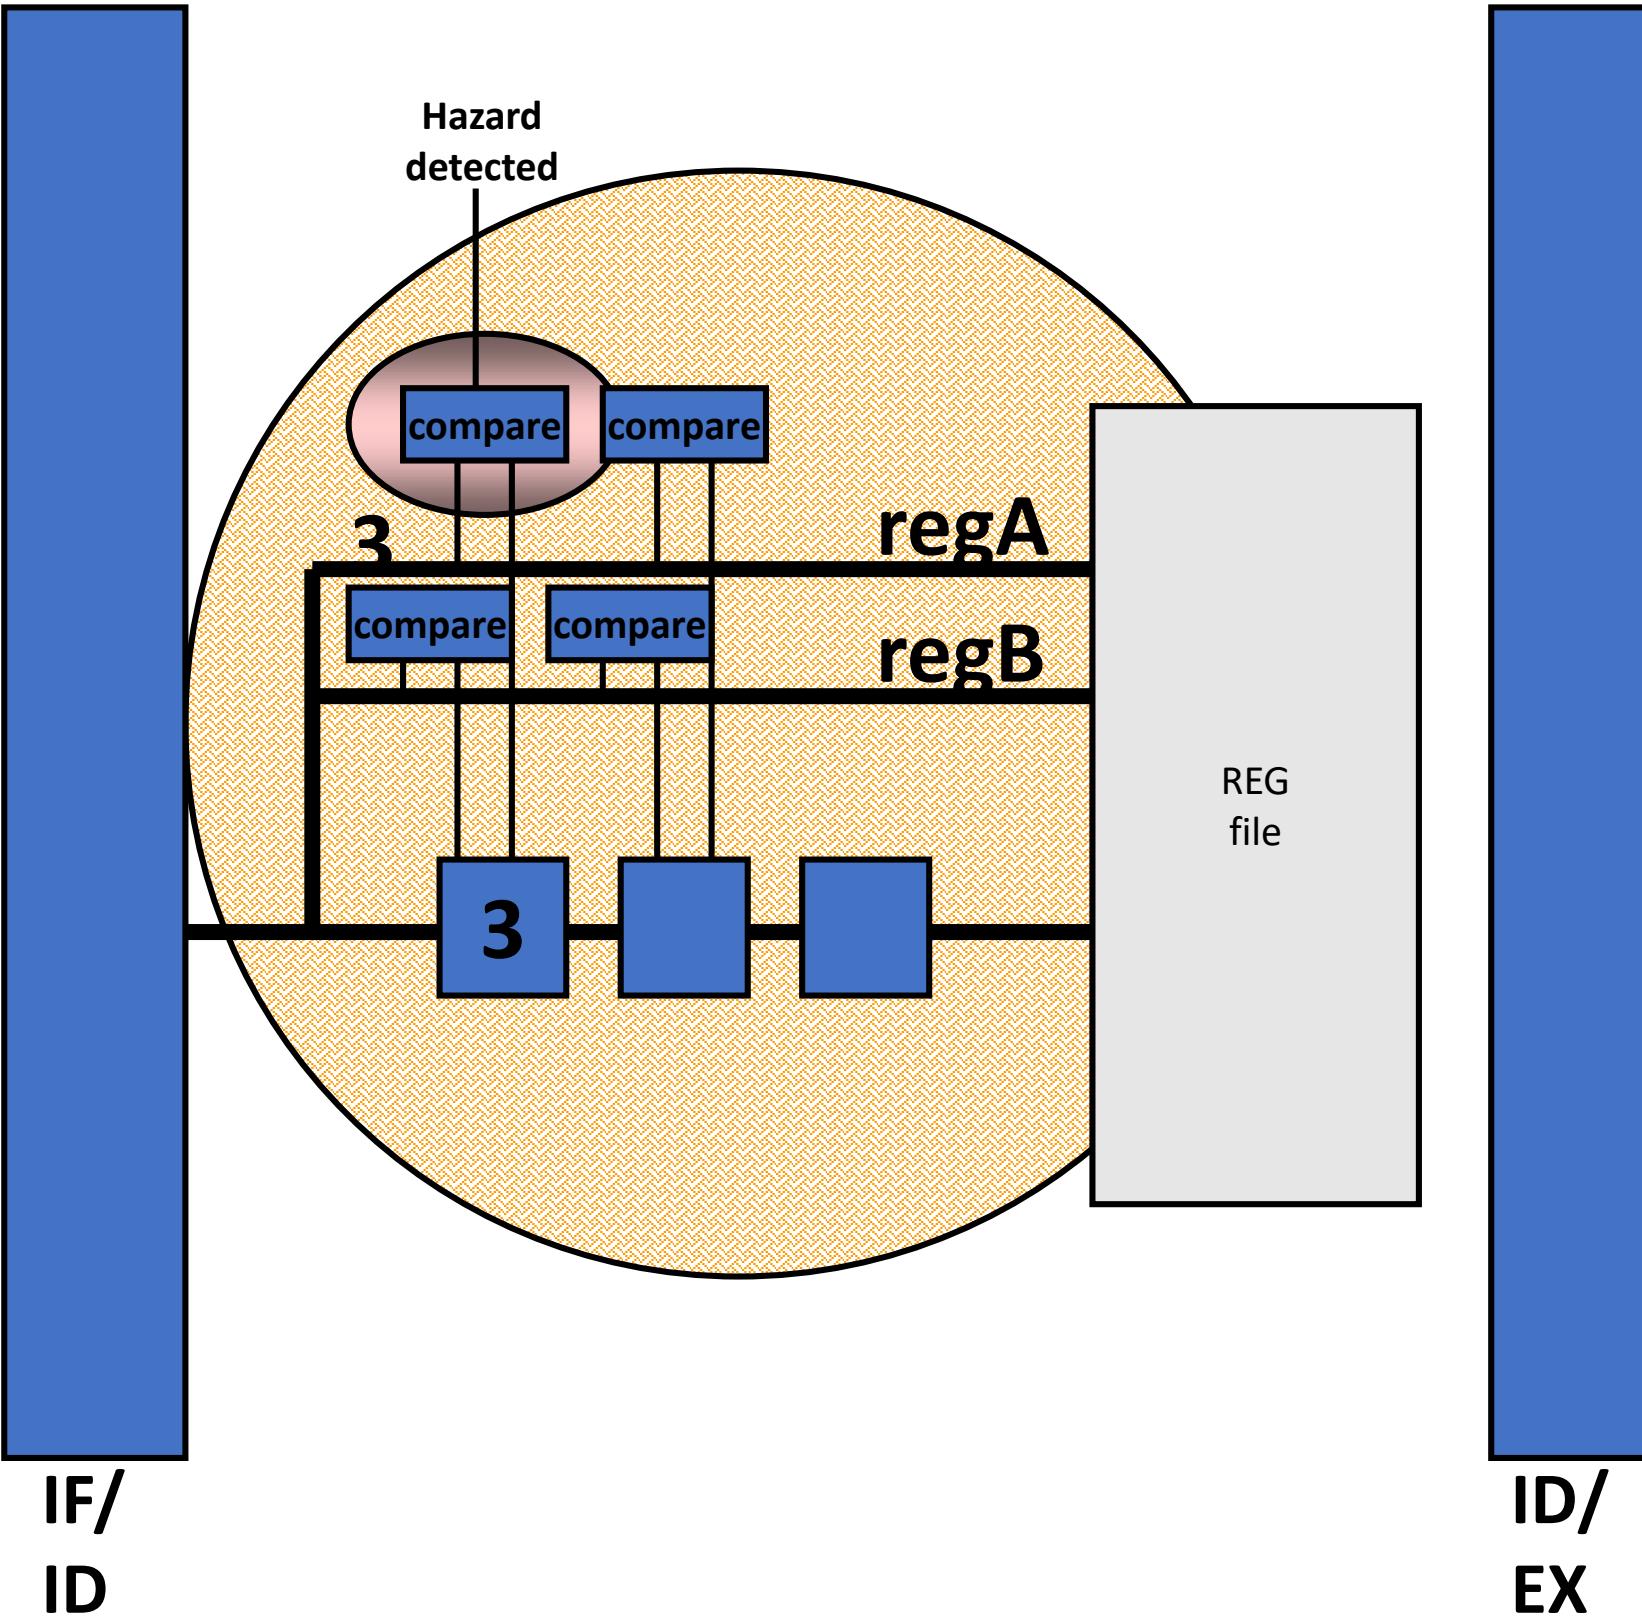
- Detect:
  - Compare regA with previous DestRegs
    - 3 bit operand fields
  - Compare regB with previous DestRegs
    - 3 bit operand fields
- Stall:
  - Keep current instructions in fetch and decode
  - Pass a noop to execute
- How do we modify the pipeline to do this?

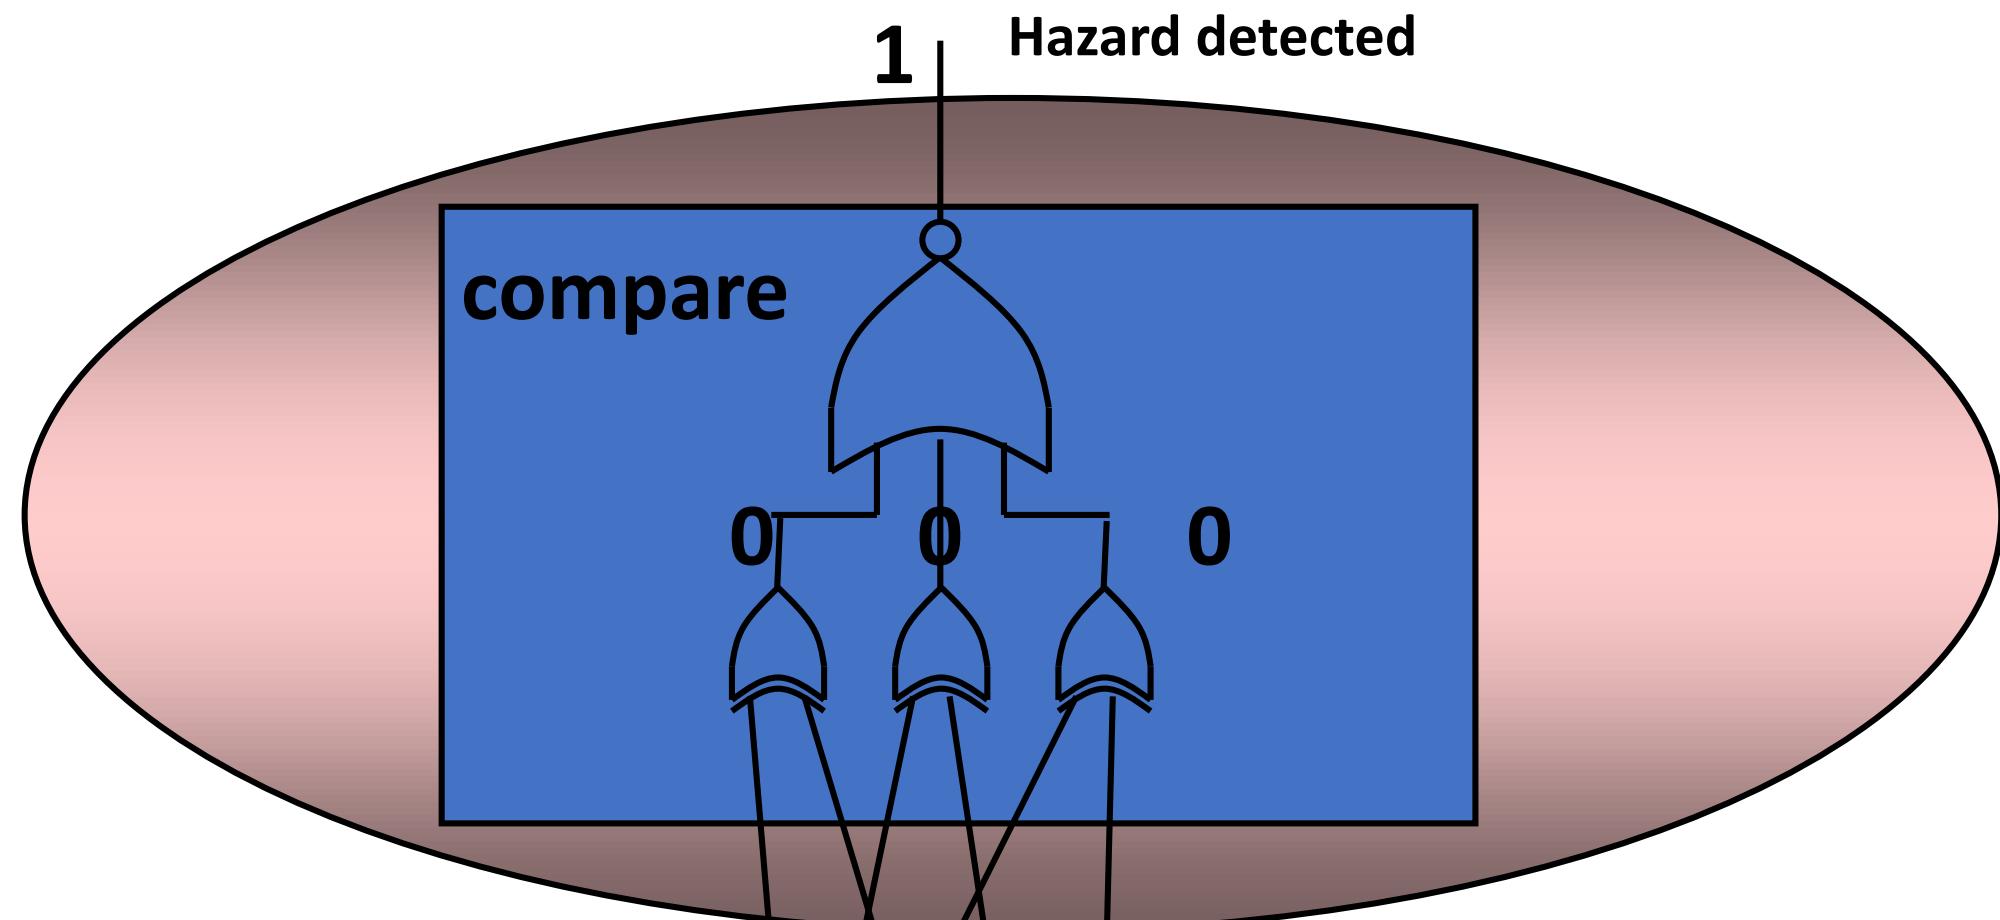


Our pipeline currently does not handle hazards—let's fix it









regA

regB

0 1 1

3

# Example

- Let's run this program with a data hazard through our 5-stage pipeline

**add 1 2 3**

**nor 3 4 5**

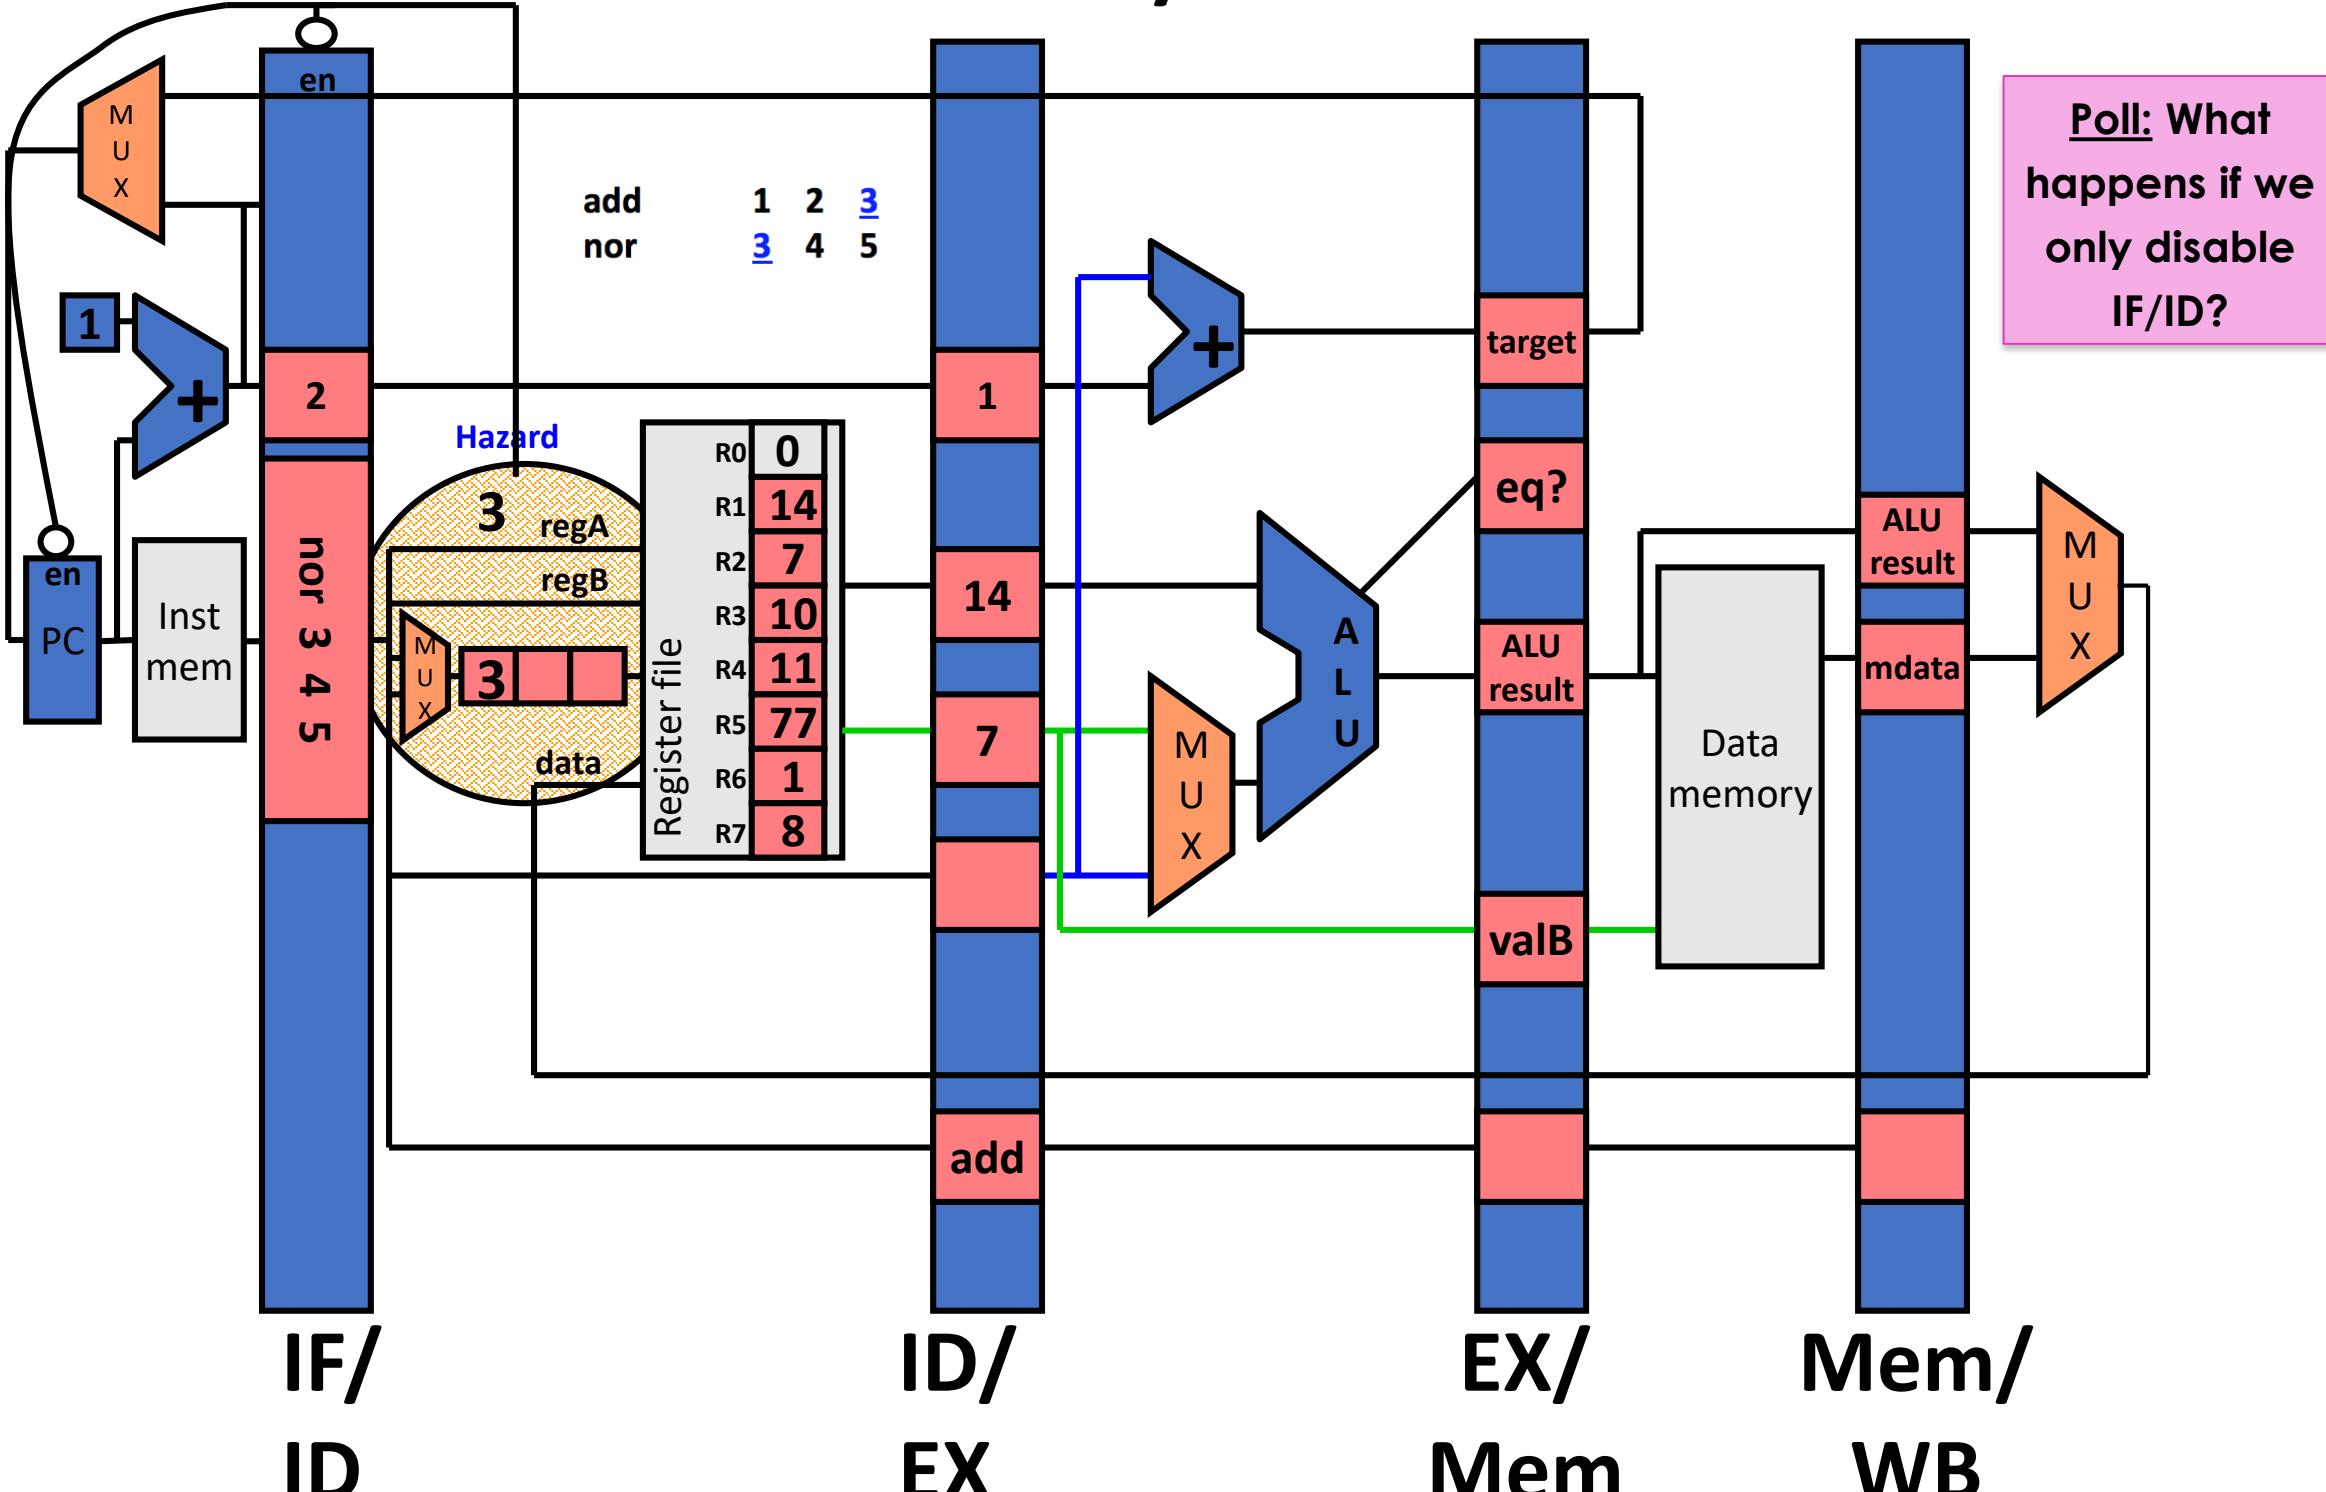
- We will start at the beginning of cycle 3, where add is in the EX stage, and nor is in the ID stage, about to read a register value

Time:	1	2	3
add 1 2 3	IF	ID	EX
nor 3 4 5		IF	ID

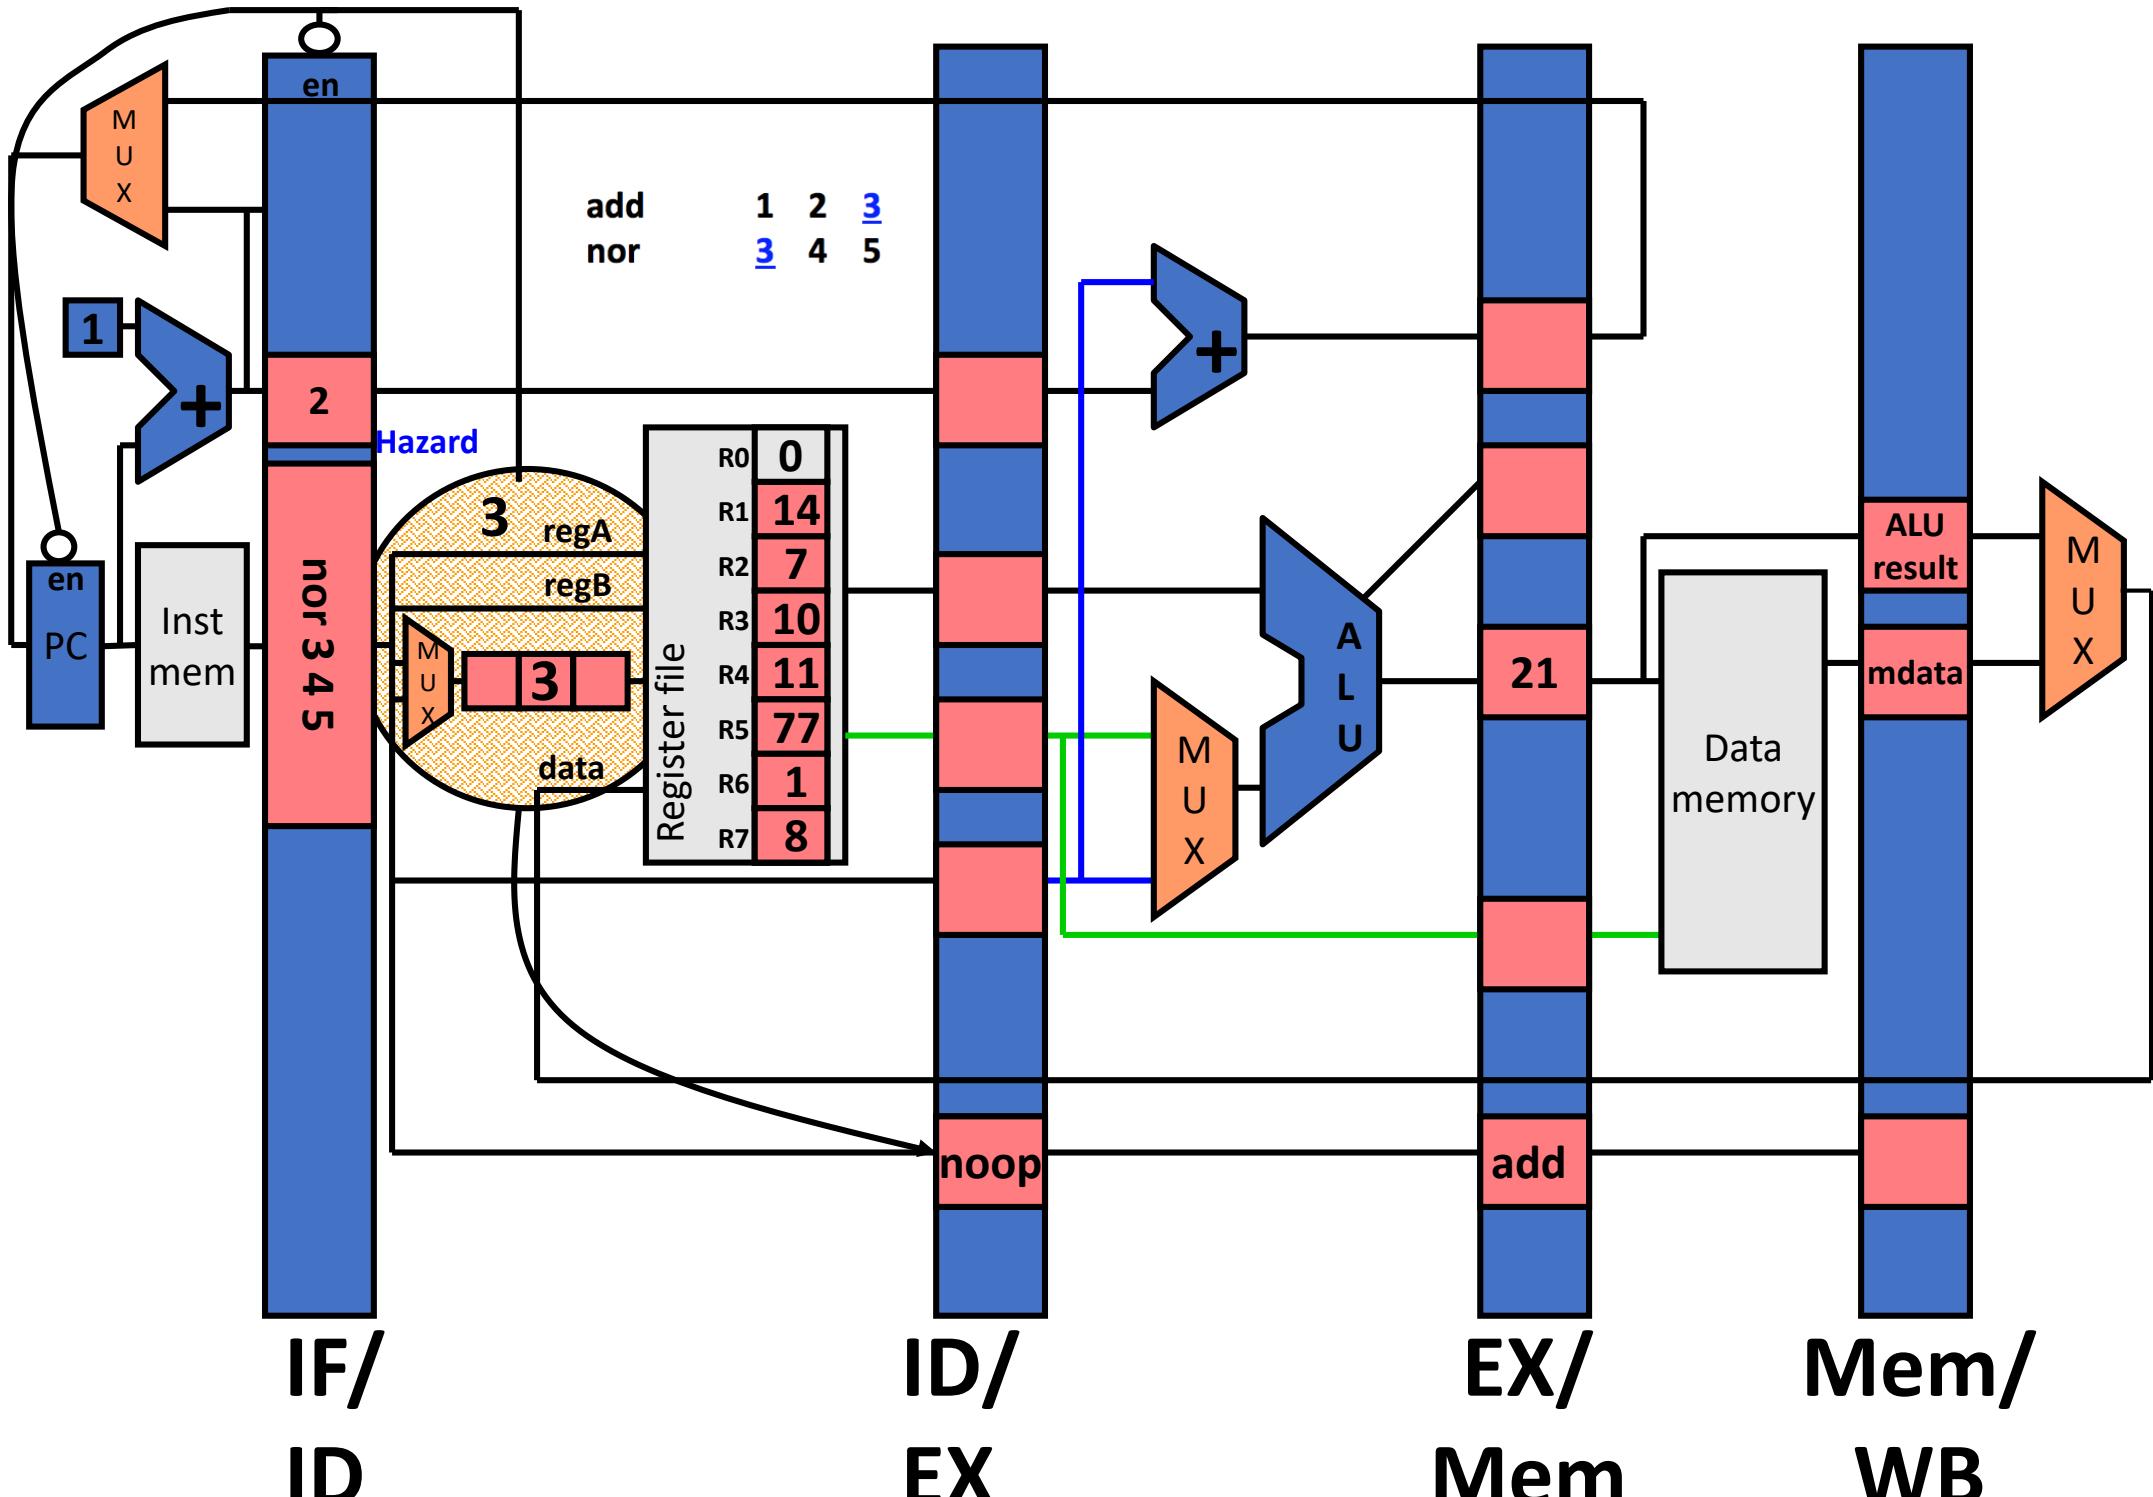
Hazard!



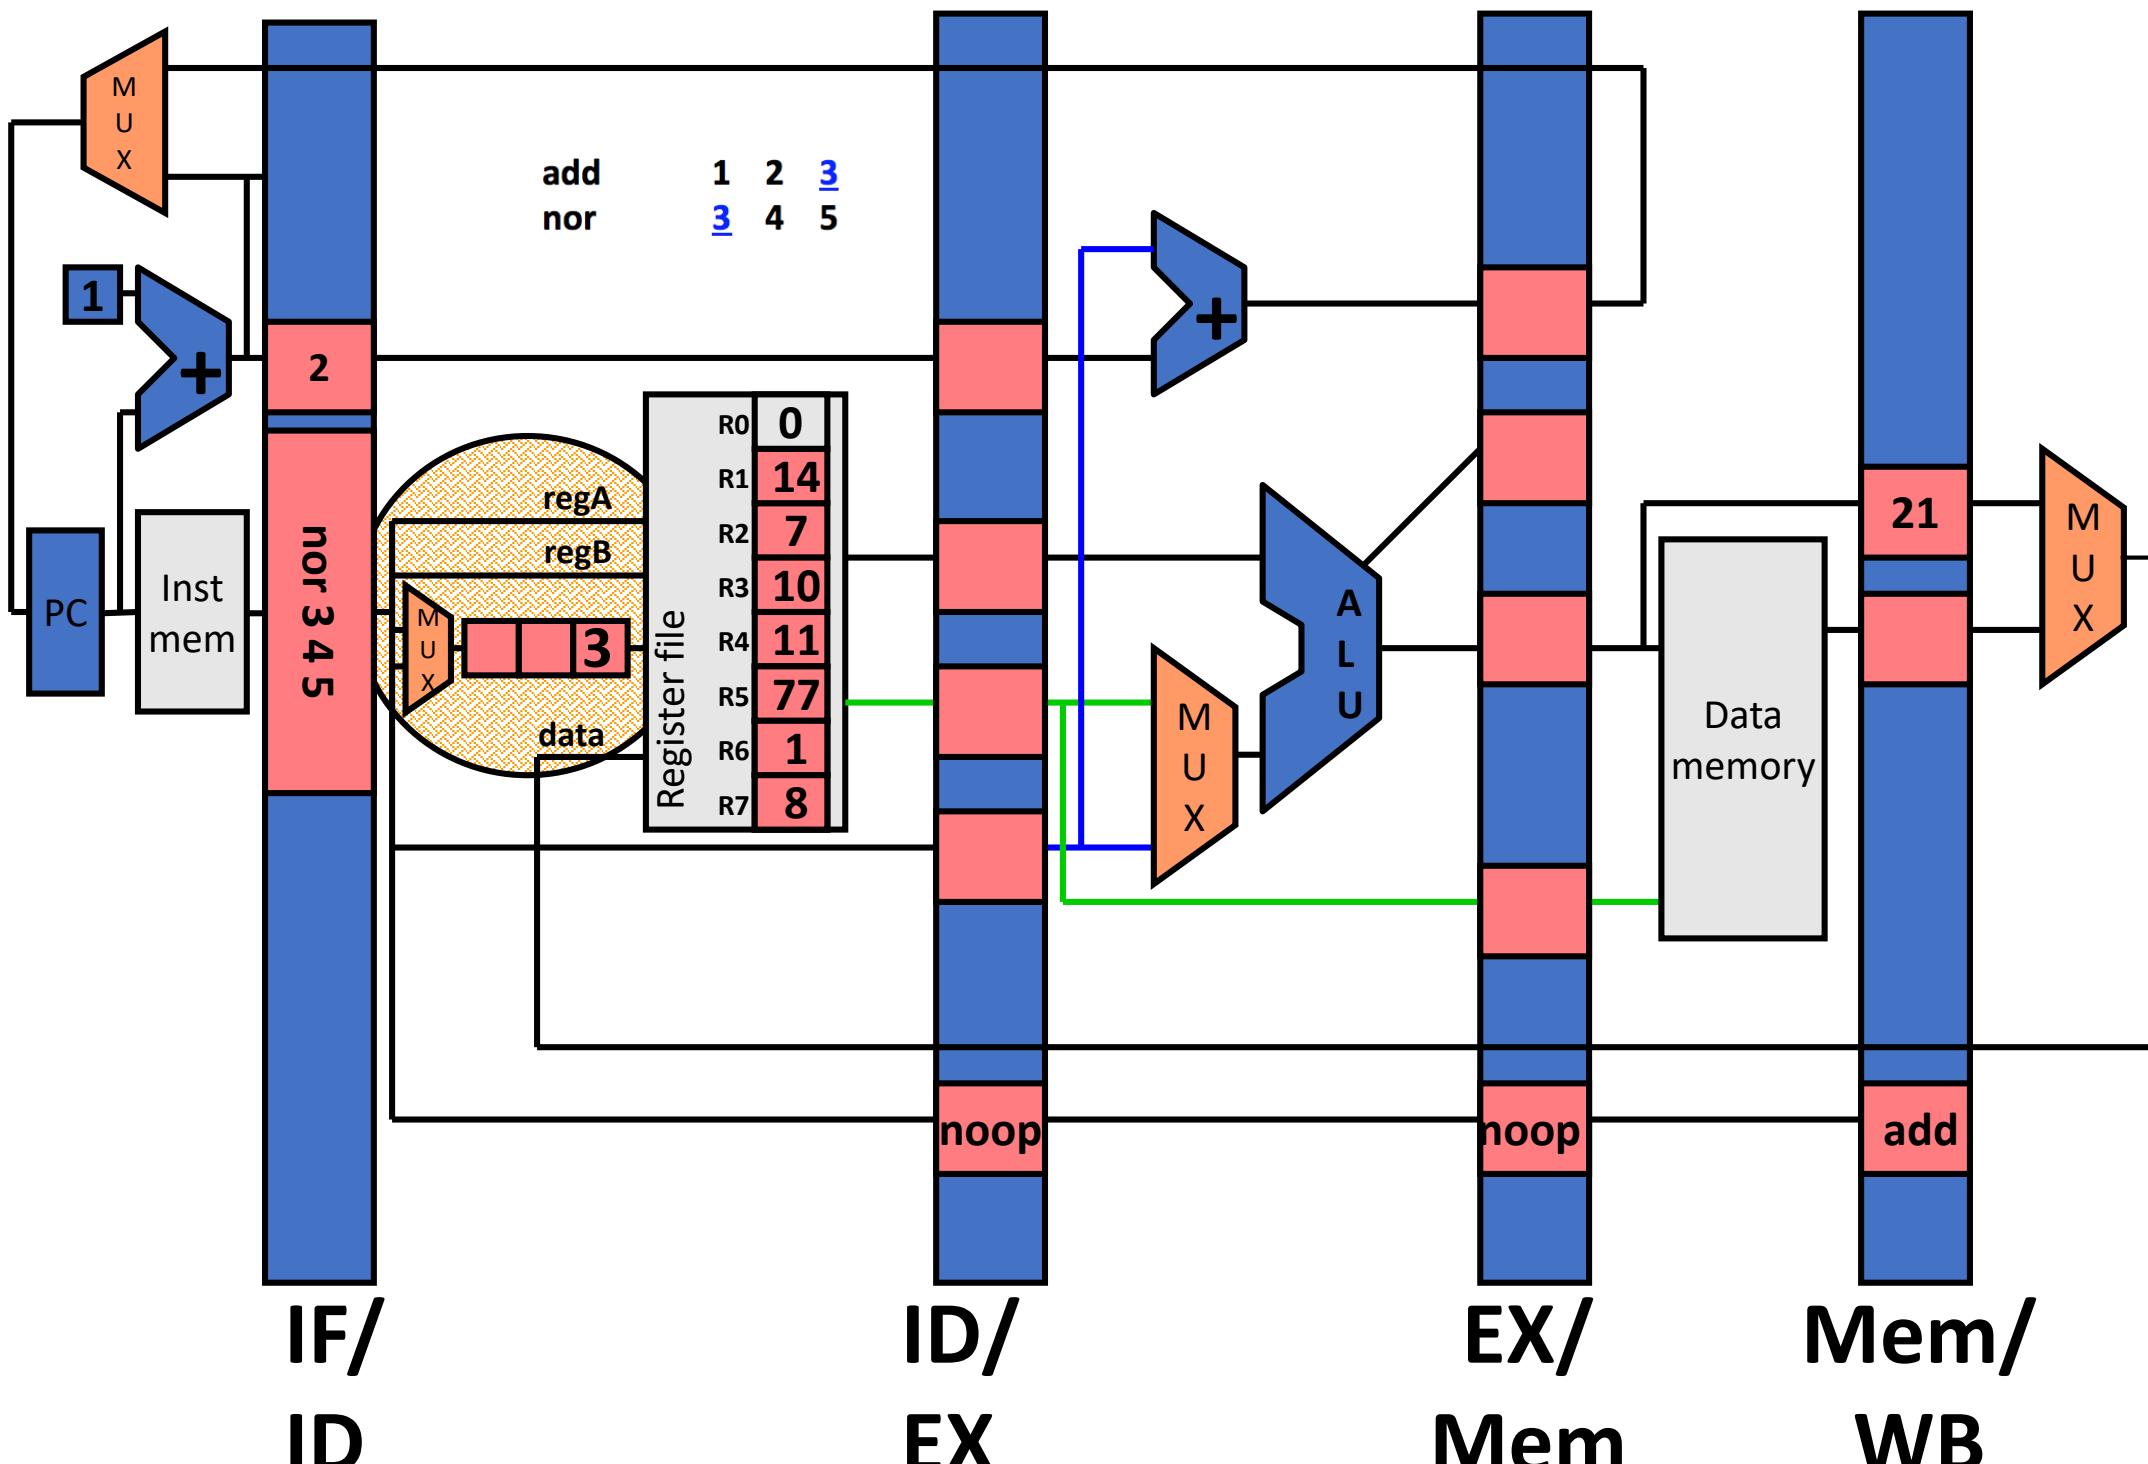
# First half of cycle 3



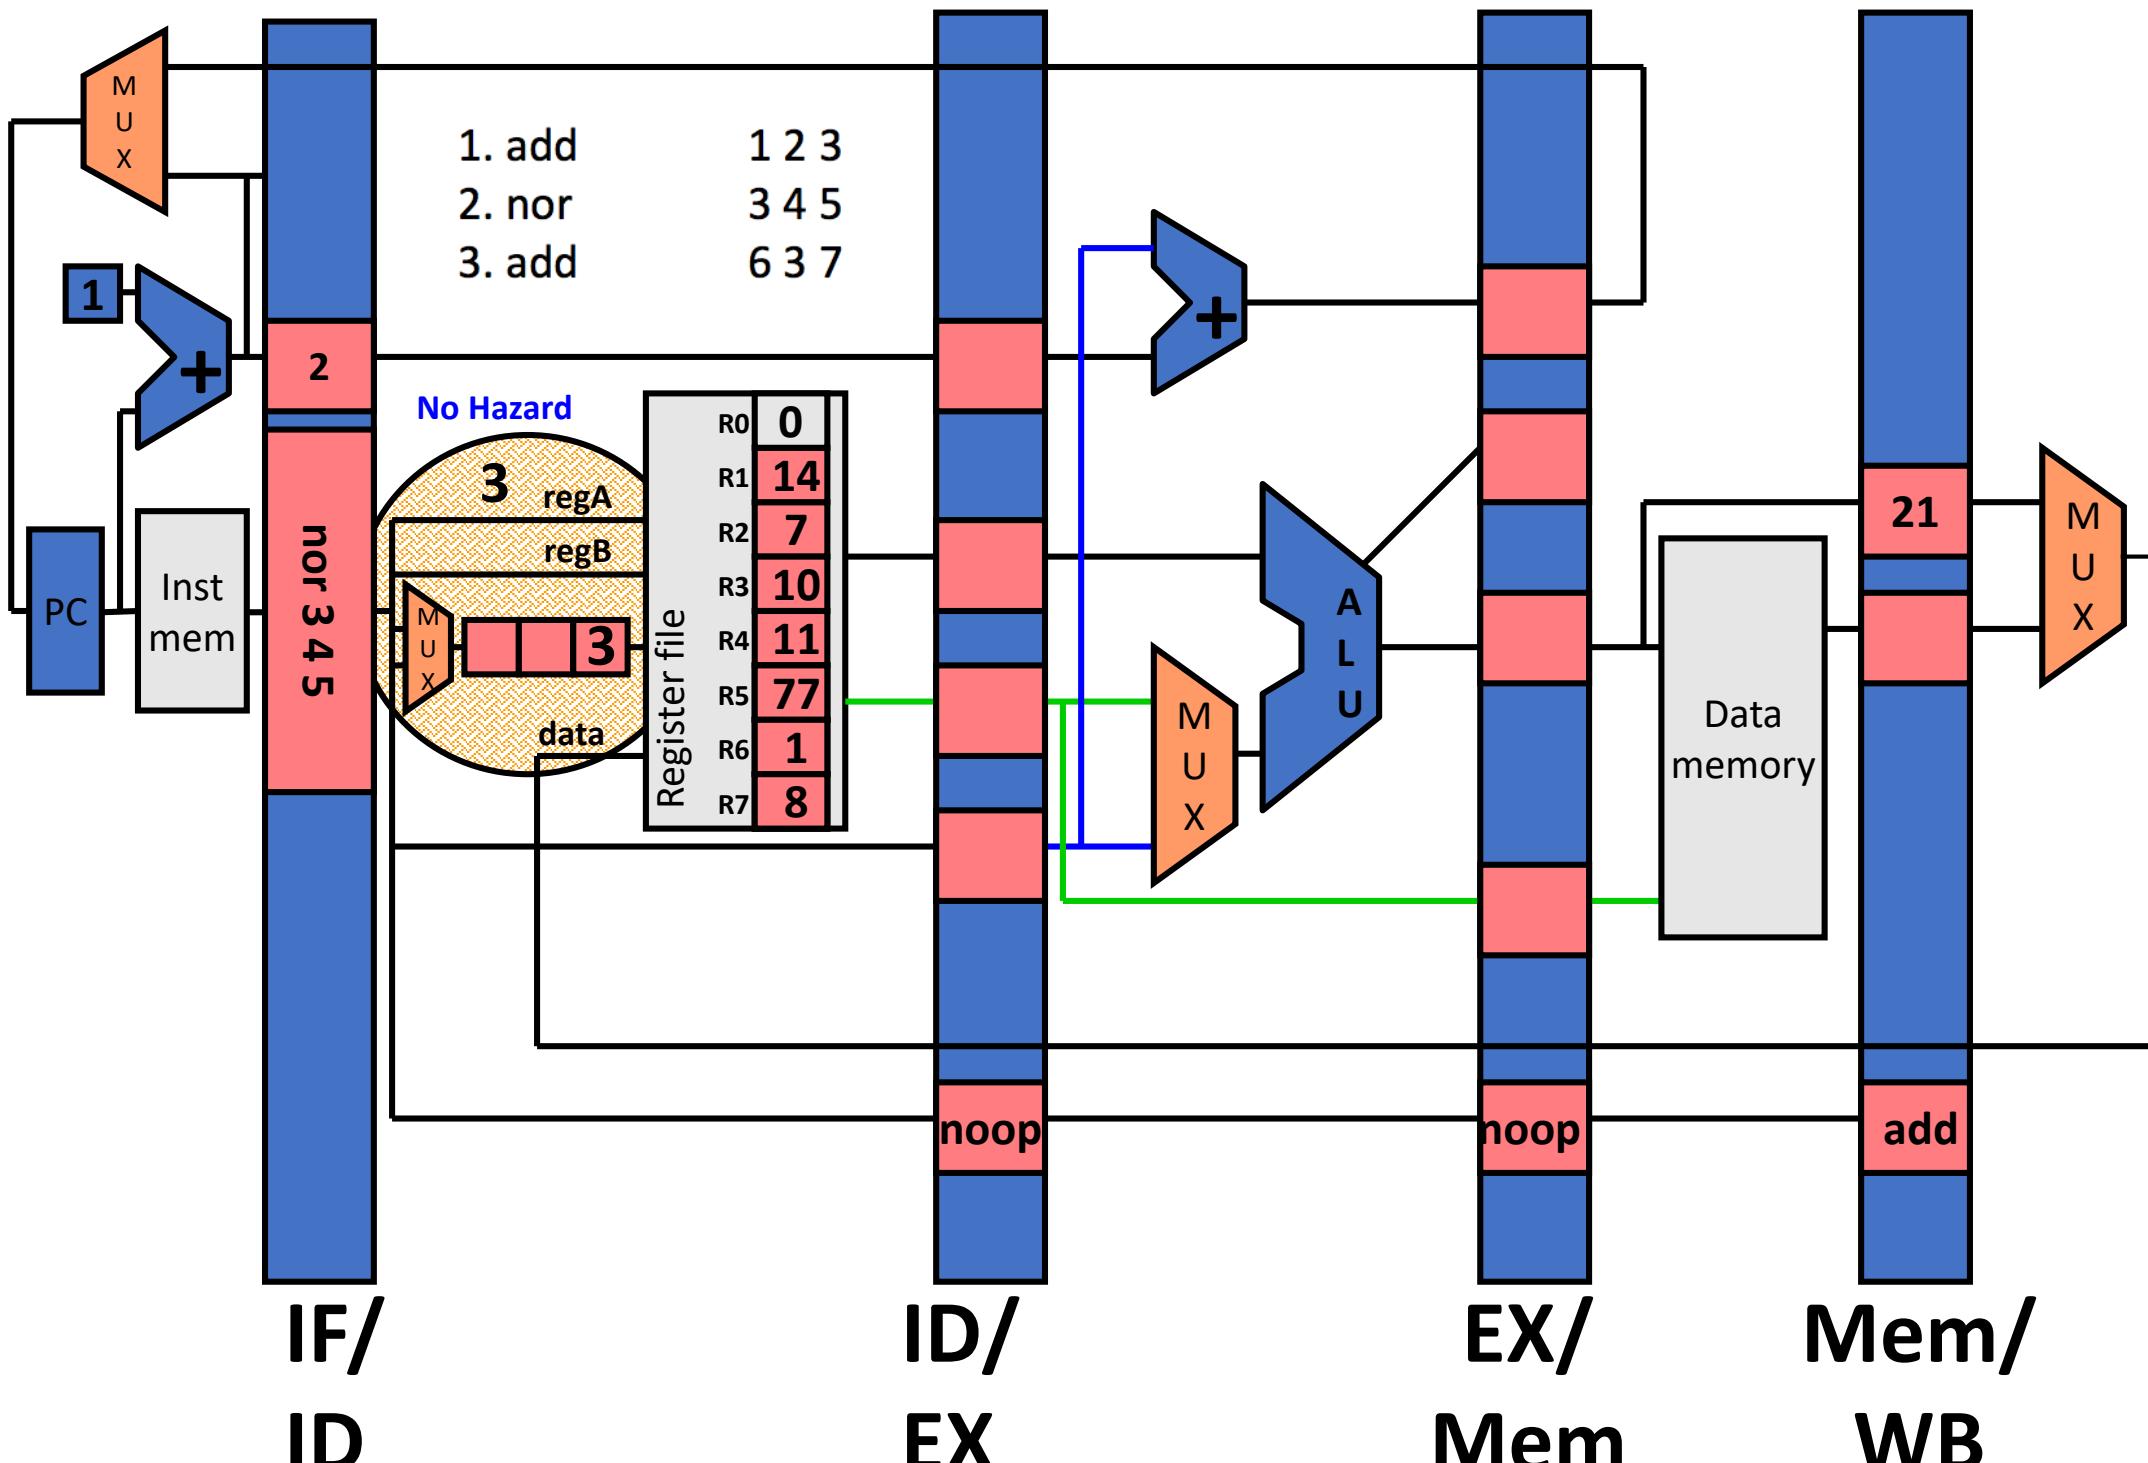
# First half of cycle 4



# End of cycle 4

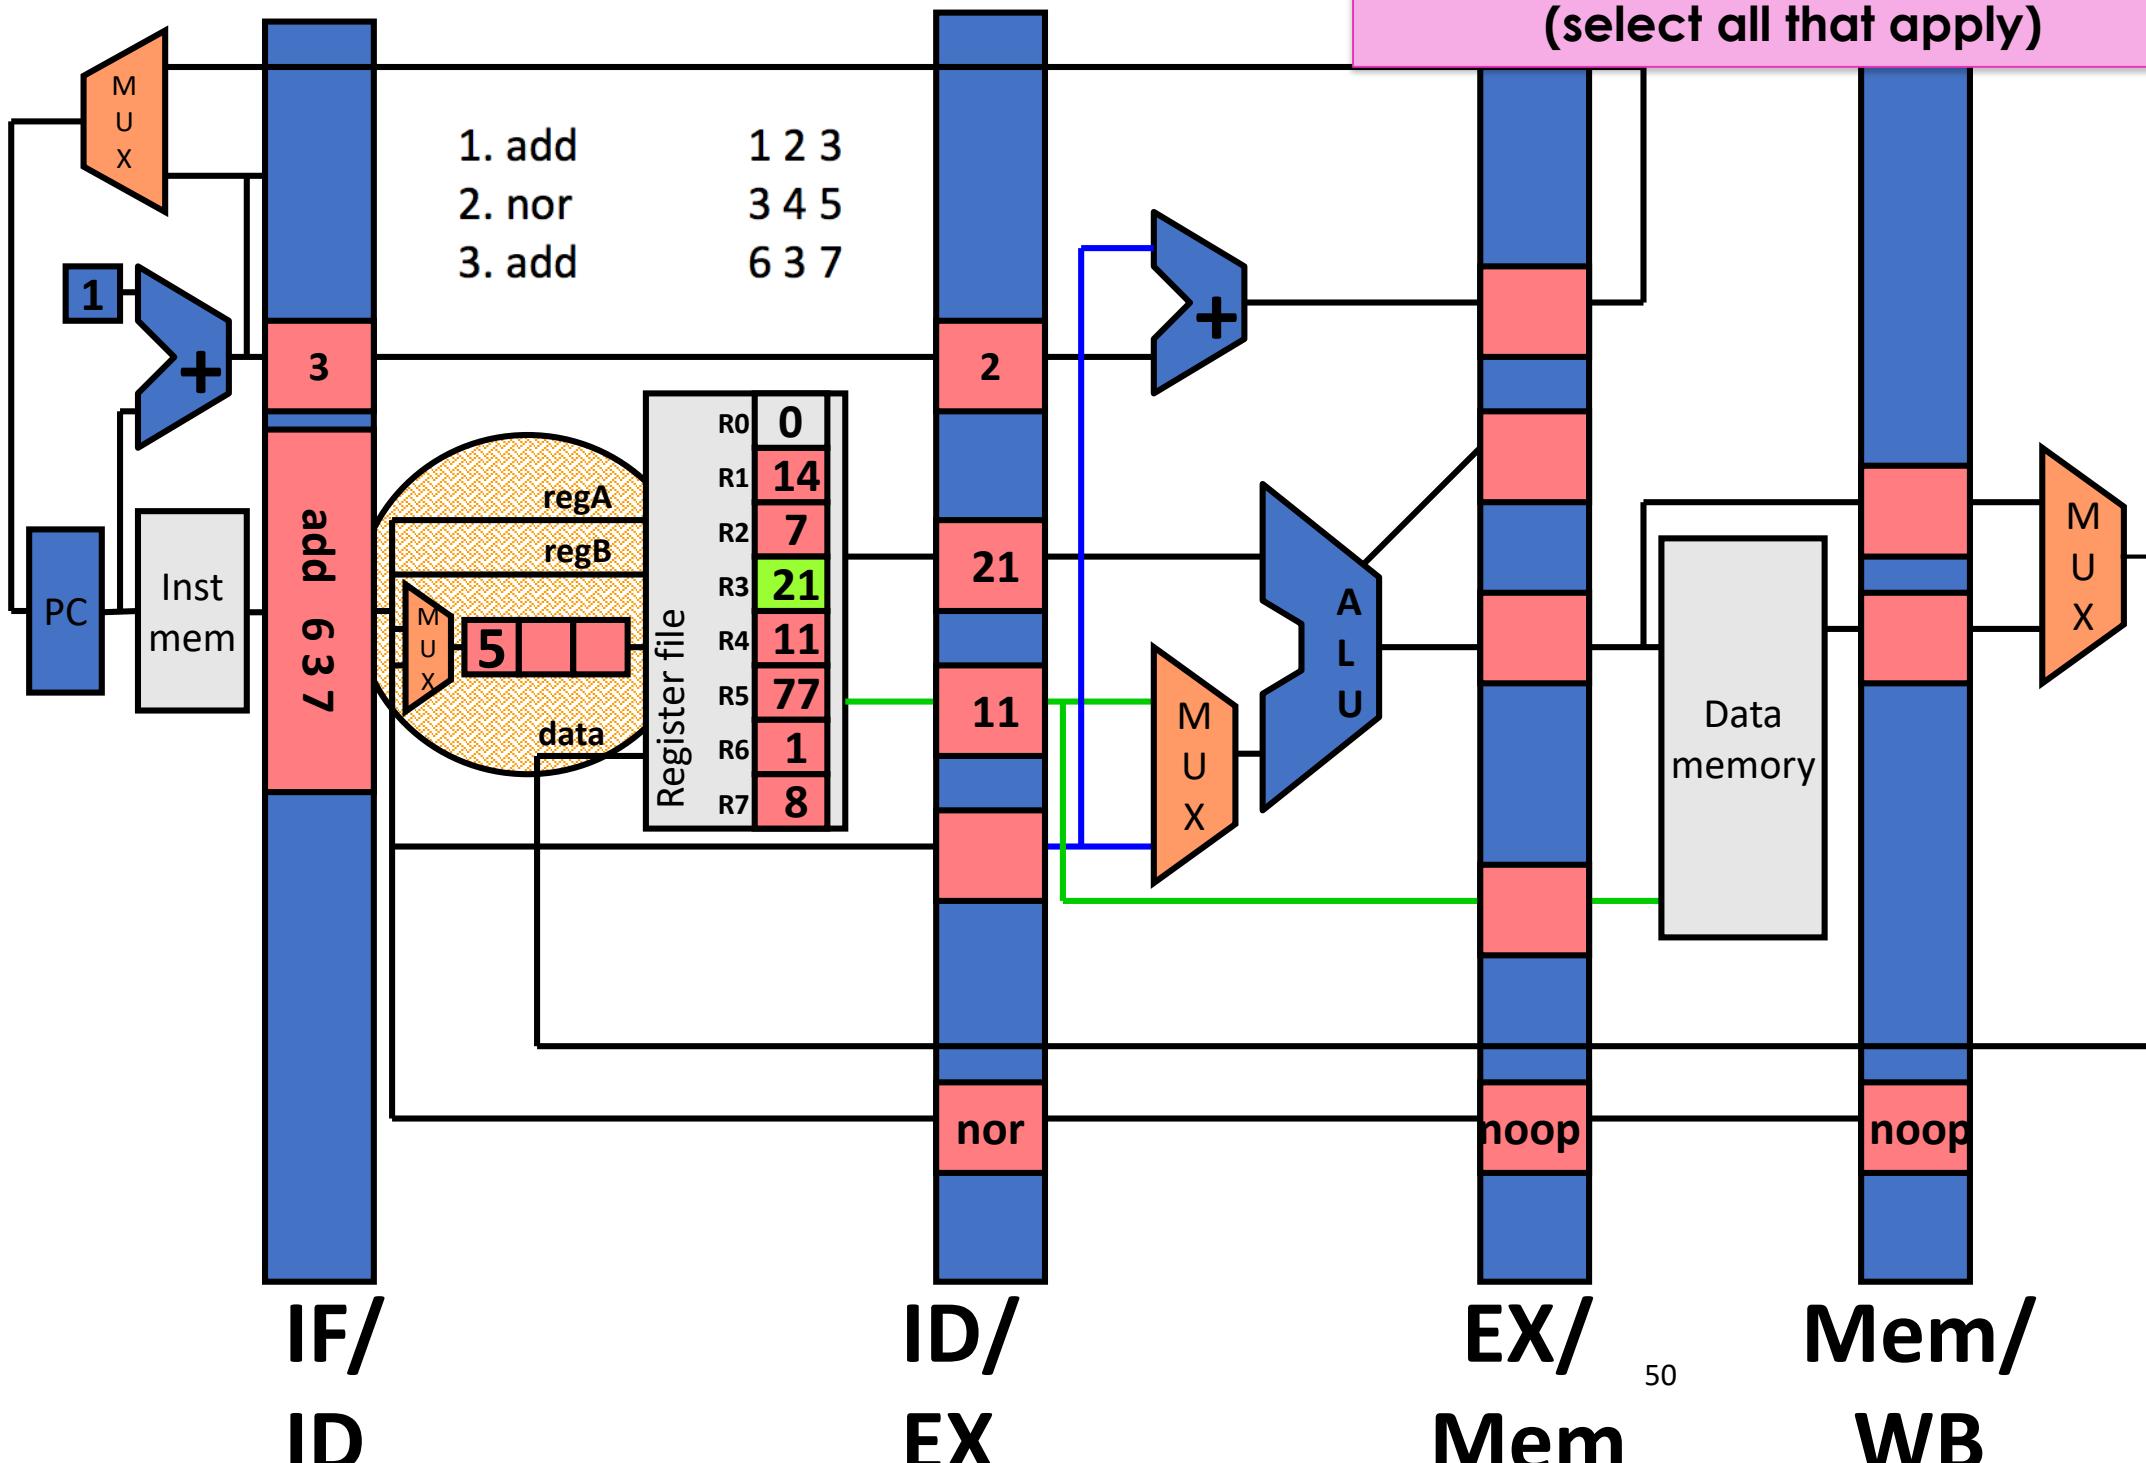


# First half of cycle 5



# End of cycle 5

How many cycles of delay can each hazard introduce? 3, 2, 1, 0  
(select all that apply)



# Time Graph

Time:	1	2	3	4	5	6	7	8	9	10	11	12	13
add 1 2 3	IF												
nor 3 4 5													

# Time Graph

Time:	1	2	3	4	5	6	7	8	9	10	11	12	13
add 1 2 3	IF	ID	EX	ME	WB								
nor 3 4 5		IF	<b>ID*</b>	<b>ID*</b>	ID	EX	ME	WB					
add 6 3 7													
lw 3 6 10													
sw 6 2 12													

# Solution

Which problems does "detect and stall" fix over  
"avoid hazards"? (select all)

1. Breaking backwards compatibility
2. Larger programs
3. Slower programs

Time:	1	2	3	4	5	6	7	8	9	10	11	12	13
add 1 2 3	IF	ID	EX	ME	WB								
nor 3 4 5		IF	<b>ID*</b>	<b>ID*</b>	ID	EX	ME	WB					
add 6 3 7					IF	ID	EX	ME	WB				
lw 3 6 10						IF	ID	EX	ME	WB			
sw 6 2 12							IF	<b>ID*</b>	<b>ID*</b>	ID	EX	ME	WB

# Next time

- Resolving Hazards